

OS Lab – 2

Madhav Bhardwaj (CS23B035)

August 2025

Problem 1

Files modified and changes made

- First of all, the trace is a system call and we are also testing it by calling this system call from trace.c file. It is not a user program by itself but a system call.
- Since we are adding trace.c for testing, we need to add it in UPROGS to build it. Therefore after building it, we can just call trace with suitable command line arguments.
- We have defined int trace(int) in the user/user.h. This trace will call the syscall trace internally. This is just telling the compiler there is an implementation of trace to link with somewhere and can let other user programs call this trace(mask).

```
int trace(int);
```

- Entry in user/usys.pl to generate the usys.S file, inside the user/usys.pl we have

```
entry("trace");
```

It expands to a small assembly code in usys.S file which has the system call stub (performs the ecall). When this ecall is called the trap handler is invoked which is in kernel/trap.c. In that based on the type of trap it is (here it is a system call), control is transferred to kernel/syscall.c. There, based on the syscall number, suitable function is called (in case of trace, it will be sys_trace).

- Every system call has a number so assign a system call number for trace in kernel/syscall.h.
- Now for the implementation part of the system call, it is in kernel/sysproc.c. we write:

```
uint64
sys_trace(void)
{
    int mask;
    if(argint(0, &mask) < 0)
        return -1;
    myproc()->trace_mask = mask;
```

```
    return 0;
}
```

therefore when sys_trace is called, the mask is set in the process information struct proc.

- Modify fork() in kernel/proc.c here we need to copy the trace mask to forked processes so that they are also traced.
- Finally inside the syscall function in syscall.c the following changes are made just to check if the mask for the current system call is set and if yes print a msg.

```
if(p->trace_mask & (1 << p->trapframe->a7)) {
    printf("%d: syscall %s -> %lu\n",
           p->pid,
           syscalls[p->trapframe->a7].name,
           p->trapframe->a0);
}
```

- Then the trace(mask) returns to the main in trace.c and then the actual program based on the commandline arguments is executed where the systemcalls will be tracked.

Problem 2

- Modified kernel/riscv.h:

- Added a static inline function r_fp() which uses inline assembly to read the frame pointer (the s0 register) and return it as a uint64. it is placed in kernel/riscv.h because it contains all the architecture level helper functions that can read the registers directly.
- Code:

```
static inline uint64
r_fp()
{
    uint64 x;
    asm volatile("mv %0, s0" : "=r" (x));
    return x;
}
```

- Modified kernel/printf.c:

- Implemented the backtrace() function.
- The function begins by reading the current frame pointer (fp = r_fp()).
- The base of the stack page is computed as fp_base = PGROUNDDOWN(fp).
- For each frame:
 - * The return address is stored at fp - 8 and is printed using printf.
 - * The previous frame pointer is stored at fp - 16.

-
- * If the previous frame pointer goes outside the current stack page (fp_base to fp_base + PGSIZE), the loop terminates.
 - * Otherwise, fp is updated to the previous frame pointer, continuing the traversal.
- Code:

```
void
backtrace(void)
{
    uint64 fp;
    fp = r_fp();
    uint64 fp_base = PGROUNDDOWN(fp);
    while(fp >= fp_base && fp >= 0 && PGROUNDDOWN(fp) == fp_base)
    {
        printf("%p\n", (void*)(uint64)(fp - 8));
        if (*(uint64*)(fp - 16) >= fp_base + PGSIZE ||
            (*(uint64*)(fp - 16) < fp_base)) break;
        fp = *(uint64*)(fp - 16);
    }
}
```

- Modified kernel/defs.h:
 - Added a prototype for the backtrace() function so it can be called from other files.
- Usage of backtrace():
 - Called inside sys_sleep() to test with bttest.
 - Called inside panic() in kernel/printf.c to print stack traces when the kernel panics.
- Internal Working:
 - The function works by traversing the list of stack frames maintained using the s0 register (frame pointer).
 - Each frame stores the caller's frame pointer and return address.
 - By following the chain of previous frame pointers, the function can unwind the stack and print the sequence of return addresses that led to the current function.