

# Synchronization

Chester Rebeiro  
IIT Madras

# Motivating Scenario

program 0

```
{  
    *  
    *  
    counter++  
    *  
}
```

shared variable

```
int counter=5;
```

program 1

```
{  
    *  
    *  
    counter--  
    *  
}
```

- Single core
  - Program 1 and program 2 are executing at the same time but sharing a single core



→CPU usage wrt time

# Motivating Scenario

program 0

```
{  
    *  
    *  
    counter++  
    *  
}
```

Shared variable

```
int counter=5;
```

program 1

```
{  
    *  
    *  
    counter--  
    *  
}
```

- What is the value of counter?
  - expected to be 5
  - but could also be 4 and 6

# Motivating Scenario

program 0

```
{  
    *  
    *  
    counter++  
    *  
}
```

Shared variable

```
int counter=5;
```

program 1

```
{  
    *  
    *  
    counter--  
    *  
}
```

$R1 \leftarrow \text{counter}$   
 $R1 \leftarrow R1 + 1$   
 $\text{counter} \leftarrow R1$   
 $R2 \leftarrow \text{counter}$   
 $R2 \leftarrow R2 - 1$   
 $\text{counter} \leftarrow R2$

context  
switch

counter = 5

$R1 \leftarrow \text{counter}$   
 $R2 \leftarrow \text{counter}$   
 $R2 \leftarrow R2 - 1$   
 $\text{counter} \leftarrow R2$   
 $R1 \leftarrow R1 + 1$   
 $\text{counter} \leftarrow R1$

counter = 6

$R2 \leftarrow \text{counter}$   
 $R2 \leftarrow \text{counter}$   
 $R2 \leftarrow R2 + 1$   
 $\text{counter} \leftarrow R2$   
 $R2 \leftarrow R2 - 1$   
 $\text{counter} \leftarrow R2$

counter = 4

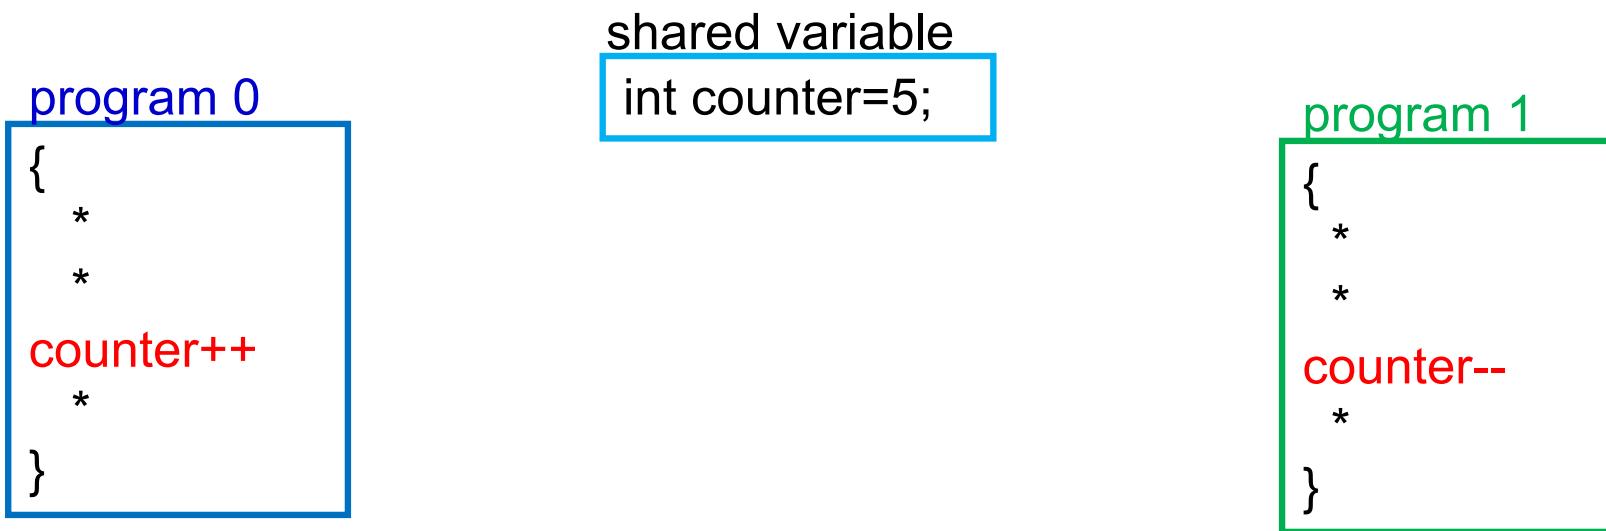
# Race Conditions

- Race conditions
  - A situation where several processes access and manipulate the same data (*critical section*)
  - The outcome depends on the order in which the access take place
  - Prevent race conditions by synchronization
    - Ensure only one process at a time manipulates the critical data

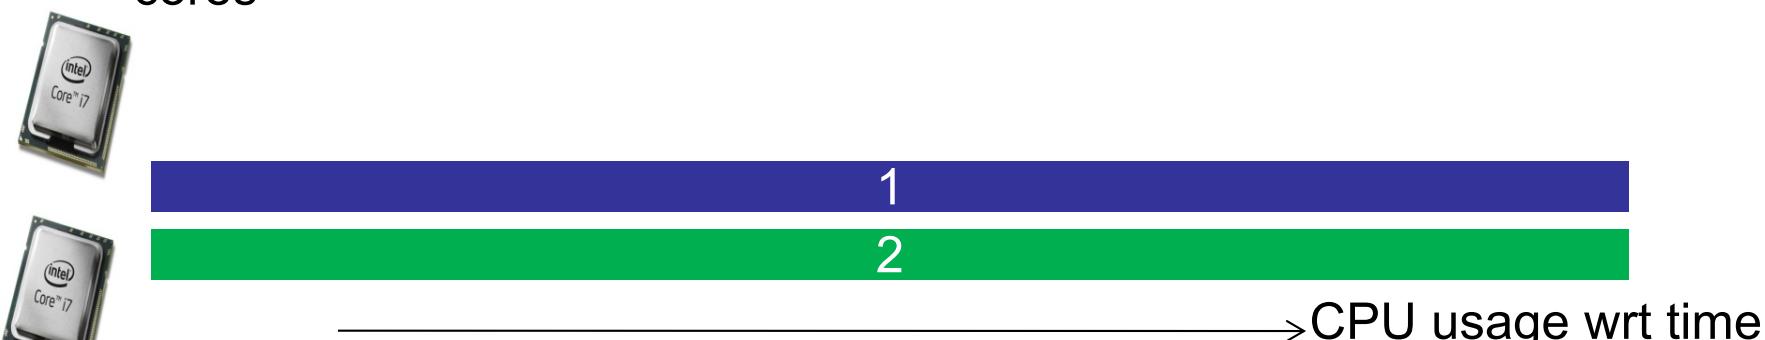


*No more than one  
process should execute in  
critical section at a time*

# Race Conditions in Multicore



- Multi core
  - Program 1 and program 2 are executing at the same time on different cores



# Critical Section

- Any solution should satisfy the following requirements
  - **Mutual Exclusion** : No more than one process in critical section at a given time
  - **Progress** : When no process is in the critical section, any process that requests entry into the critical section must be permitted without any delay
  - **No starvation (bounded wait)**: There is an upper bound on the number of times a process enters the critical section, while another is waiting.

# Locks and Unlocks

program 0

```
{  
    *  
    *  
    lock(L)  
    counter++  
    unlock(L)  
    *  
}
```

shared variable

```
int counter=5;  
lock_t L;
```

program 1

```
{  
    *  
    *  
    lock(L)  
    counter--  
    unlock(L)  
    *  
}
```

- **lock(L)** : acquire lock L exclusively
  - Only the process with L can access the critical section
- **unlock(L)** : release exclusive access to lock L
  - Permitting other processes to access the critical section

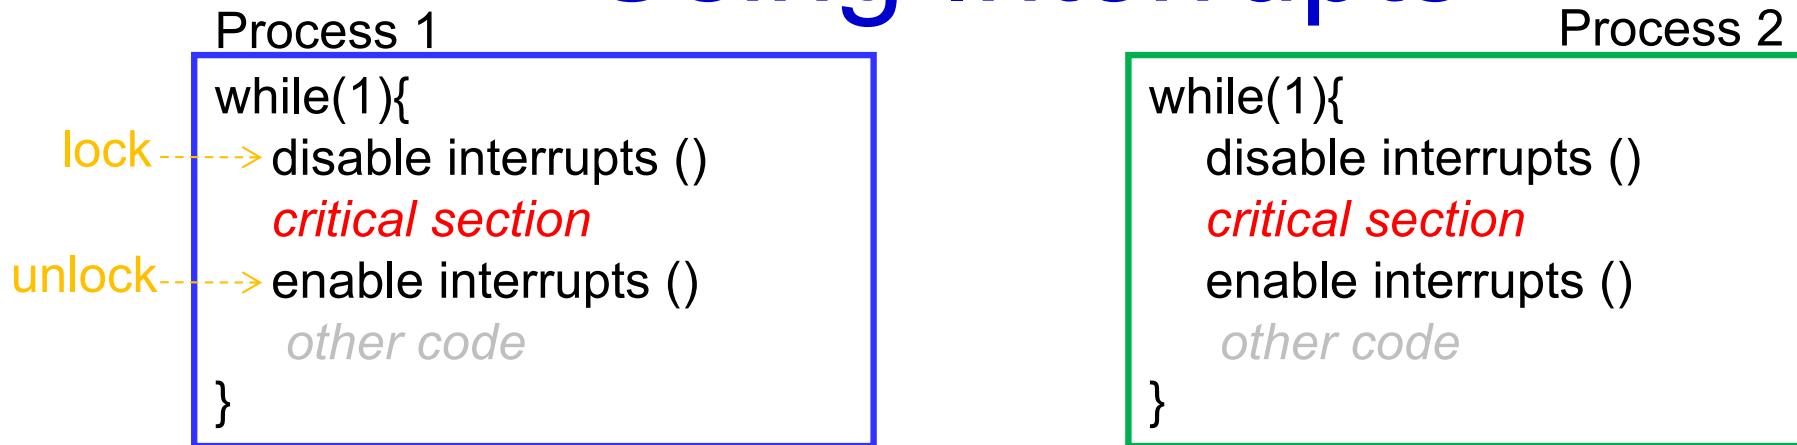
# When to have Locking?

- Single instructions by themselves are atomic
  - eg. add %eax, %ebx
- Multiple instructions need to be explicitly made atomic
  - Each piece of code in the OS must be checked if they need to be atomic

# **How to Implement Locking (Software Solutions)**

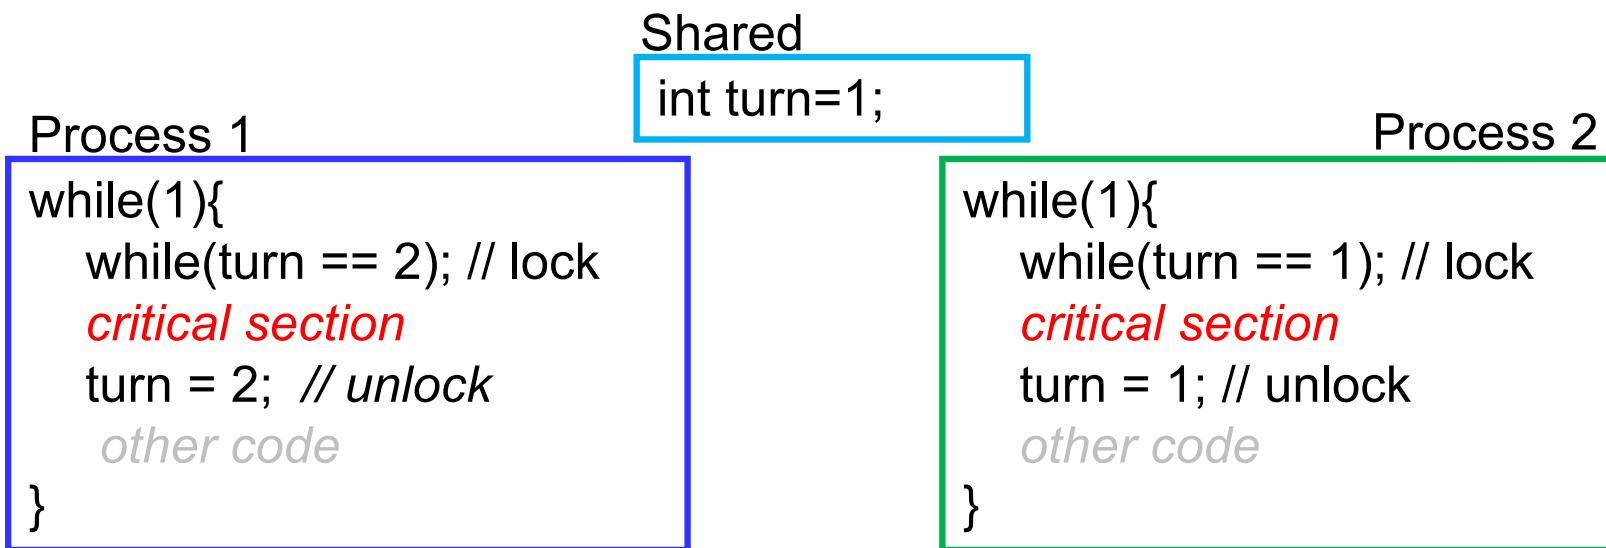
Chester Rebeiro  
IIT Madras

# Using Interrupts



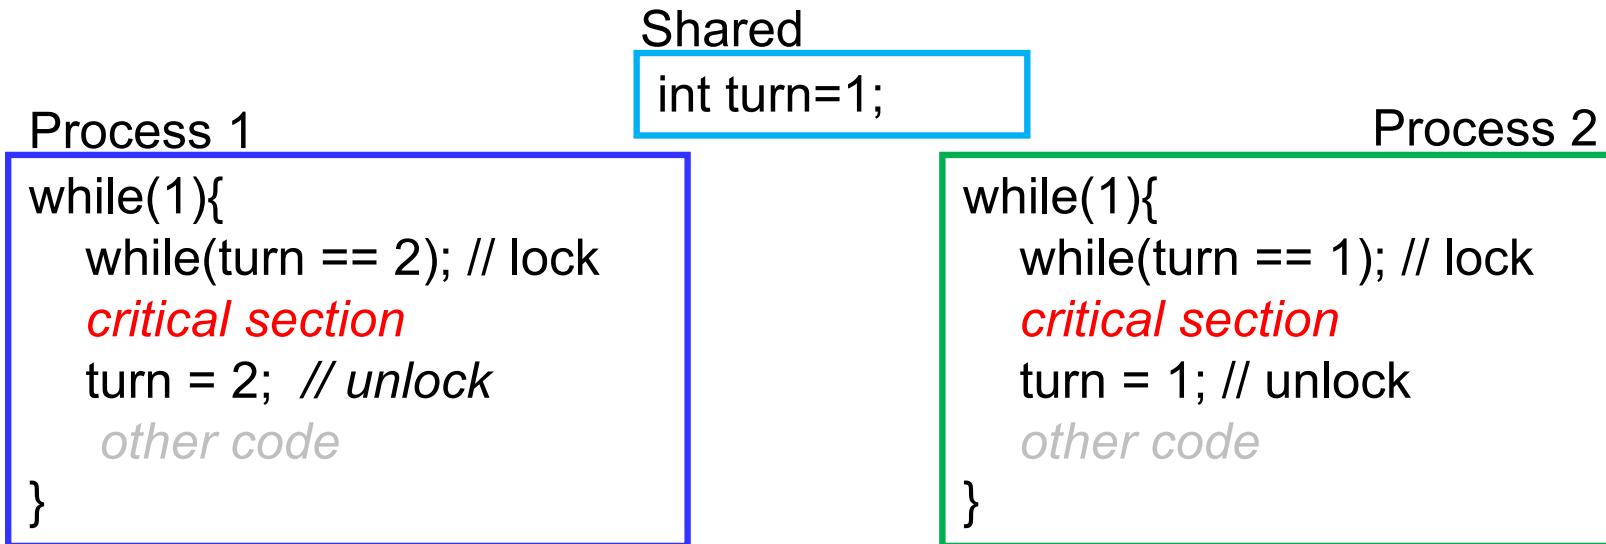
- Simple
  - When interrupts are disabled, context switches won't happen
- Requires privileges
  - User processes generally cannot disable interrupts
- Not suited for multicore systems

# Software Solution (Attempt 1)



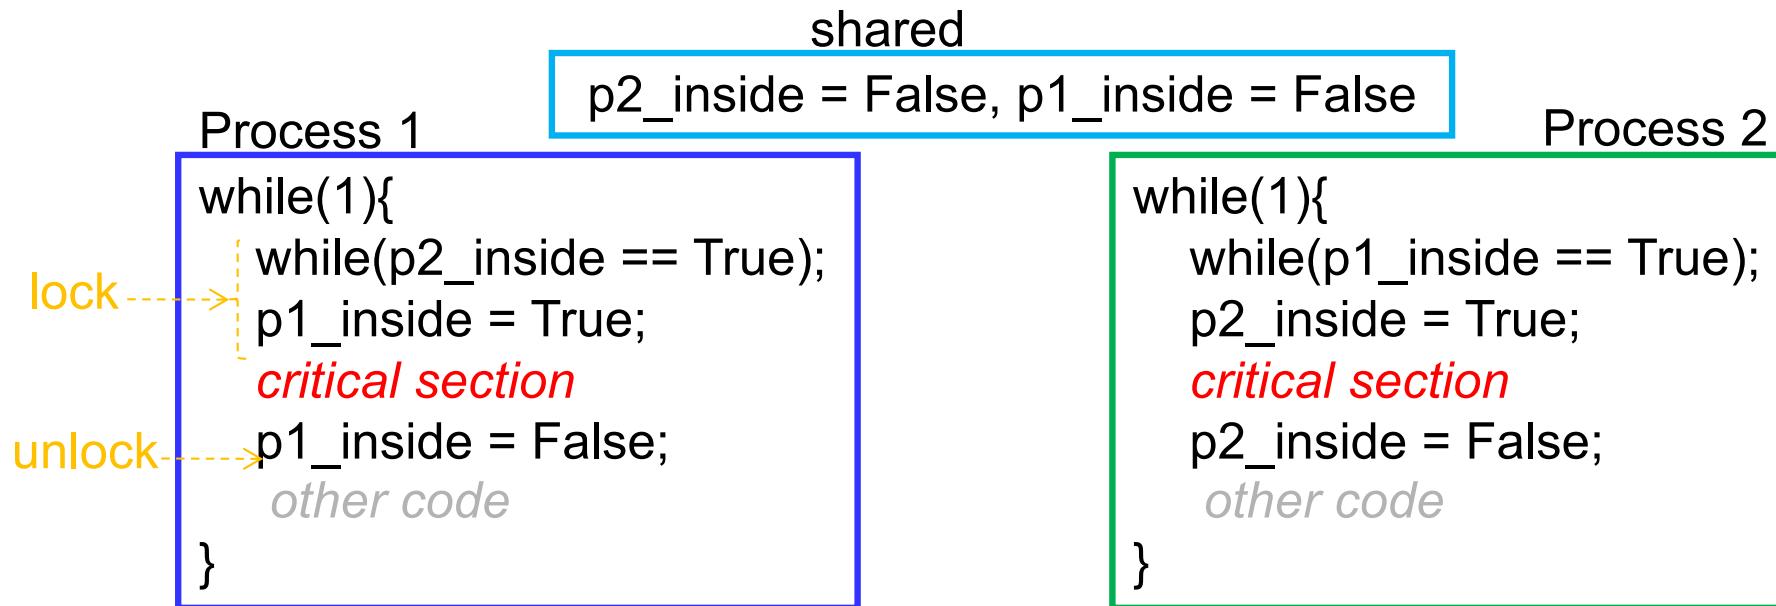
- Achieves mutual exclusion
- Busy waiting – waste of power and time
- Needs to alternate execution in critical section
  - process1 → process2 → process1 → process2*

# Problem with Attempt 1



- Had a common turn flag that was modified by both processes
- This required processes to alternate.
- Possible Solution : Have two flags – one for each process

# Software Solution (Attempt 2)



- Need not alternate execution in critical section
- Does not guarantee mutual exclusion

# Attempt 2: No mutual exclusion



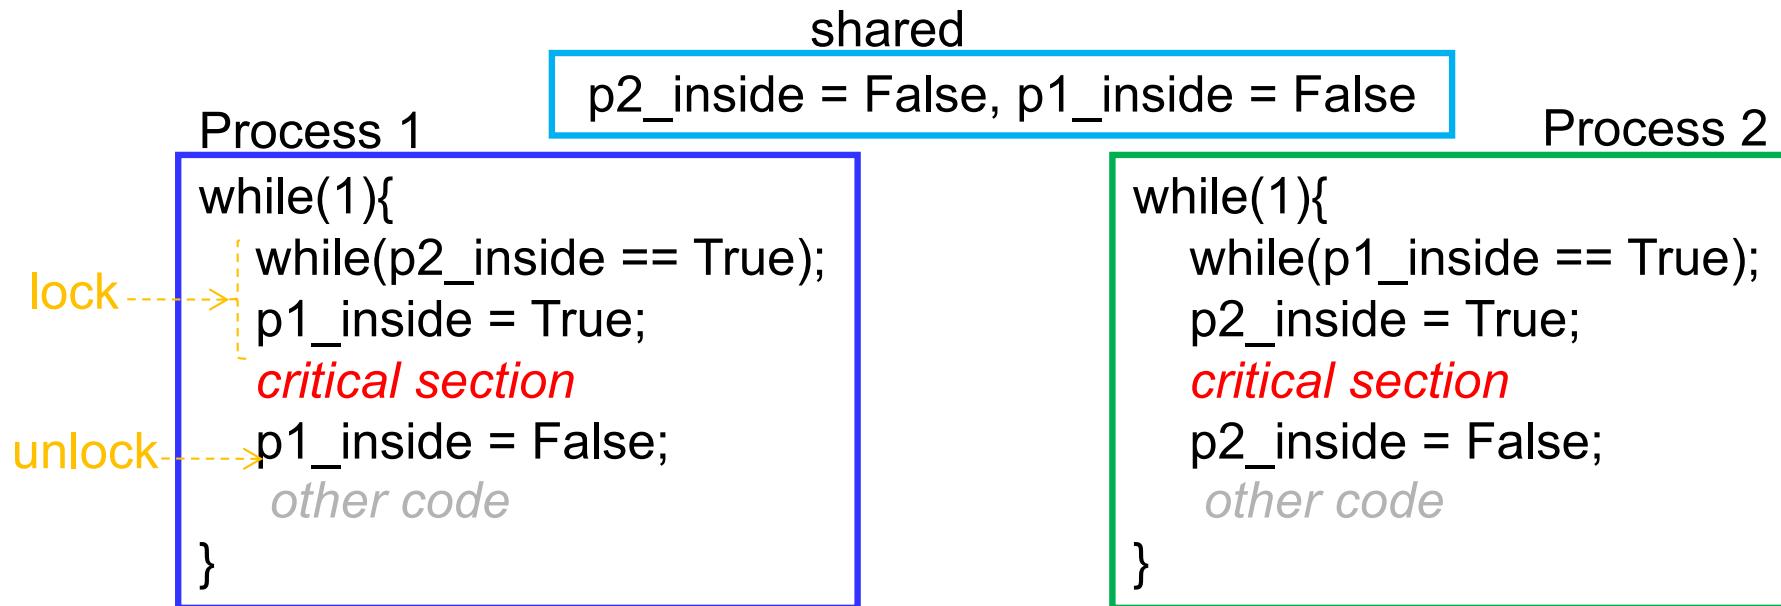
CPU	p1_inside	p2_inside
while(p2_inside == True);	False	False
context switch		
while(p1_inside == True);	False	False
p2_inside = True;	False	True
context switch		
p1_inside = True;	True	True

Both p1 and p2 can enter into the critical section at the same time

```
while(1){  
    while(p2_inside == True);  
    p1_inside = True;  
    critical section  
    p1_inside = False;  
    other code  
}
```

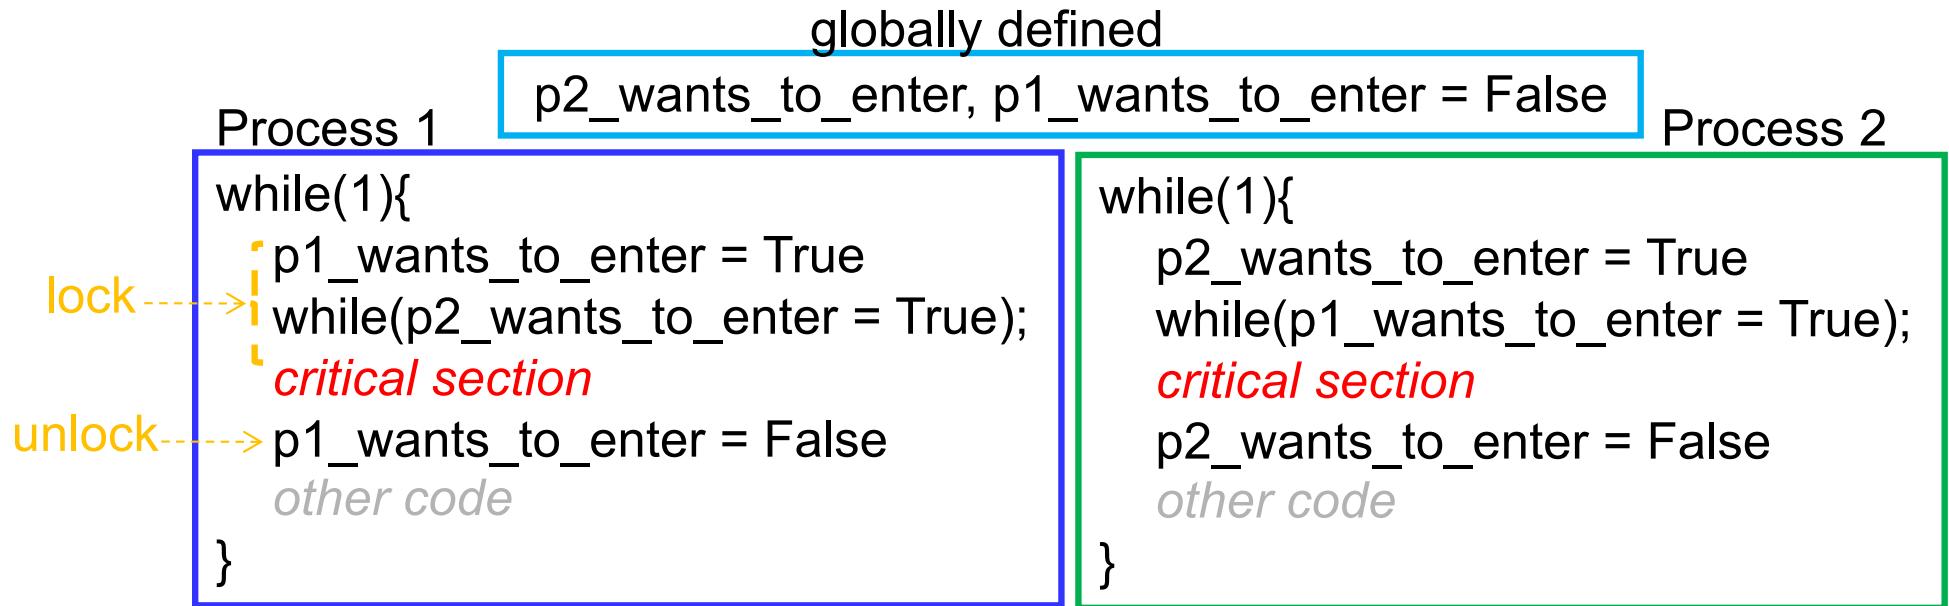
```
while(1){  
    while(p1_inside == True);  
    p2_inside = True;  
    critical section  
    p2_inside = False;  
    other code  
}
```

# Problem with Attempt 2



- The flag (p1\_inside, p2\_inside), is set after we break from the while loop.

# Software Solution (Attempt 3)



- Achieves mutual exclusion
- Does not achieve progress (could deadlock)

# Attempt 3: No Progress

CPU	p1_inside	p2_inside
p1_wants_to_enter = True	False	False
context switch		
p2_wants_to_enter = True	False	False

time ↓

There is a tie!!!

Both p1 and p2 will loop infinitely

Progress not achieved  
Each process is waiting for the other  
this is a deadlock

```
while(1){  
    p2_wants_to_enter = True  
    while(p1_wants_to_enter = True);  
    critical section  
    p2_wants_to_enter = False  
    other code  
}
```

# Deadlock

↓ time

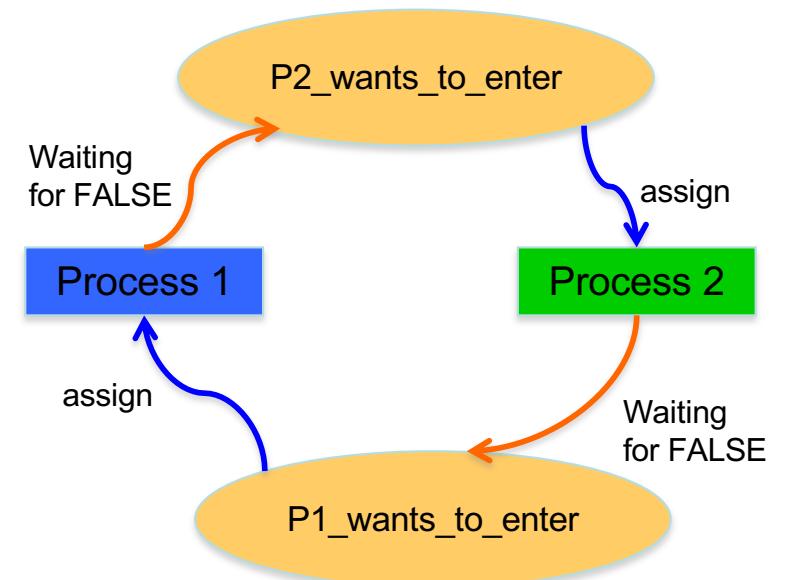
CPU	p1_inside	p2_inside
p1_wants_to_enter = True	False	False
context switch		
p2_wants_to_enter = True	False	False

There is a tie!!!

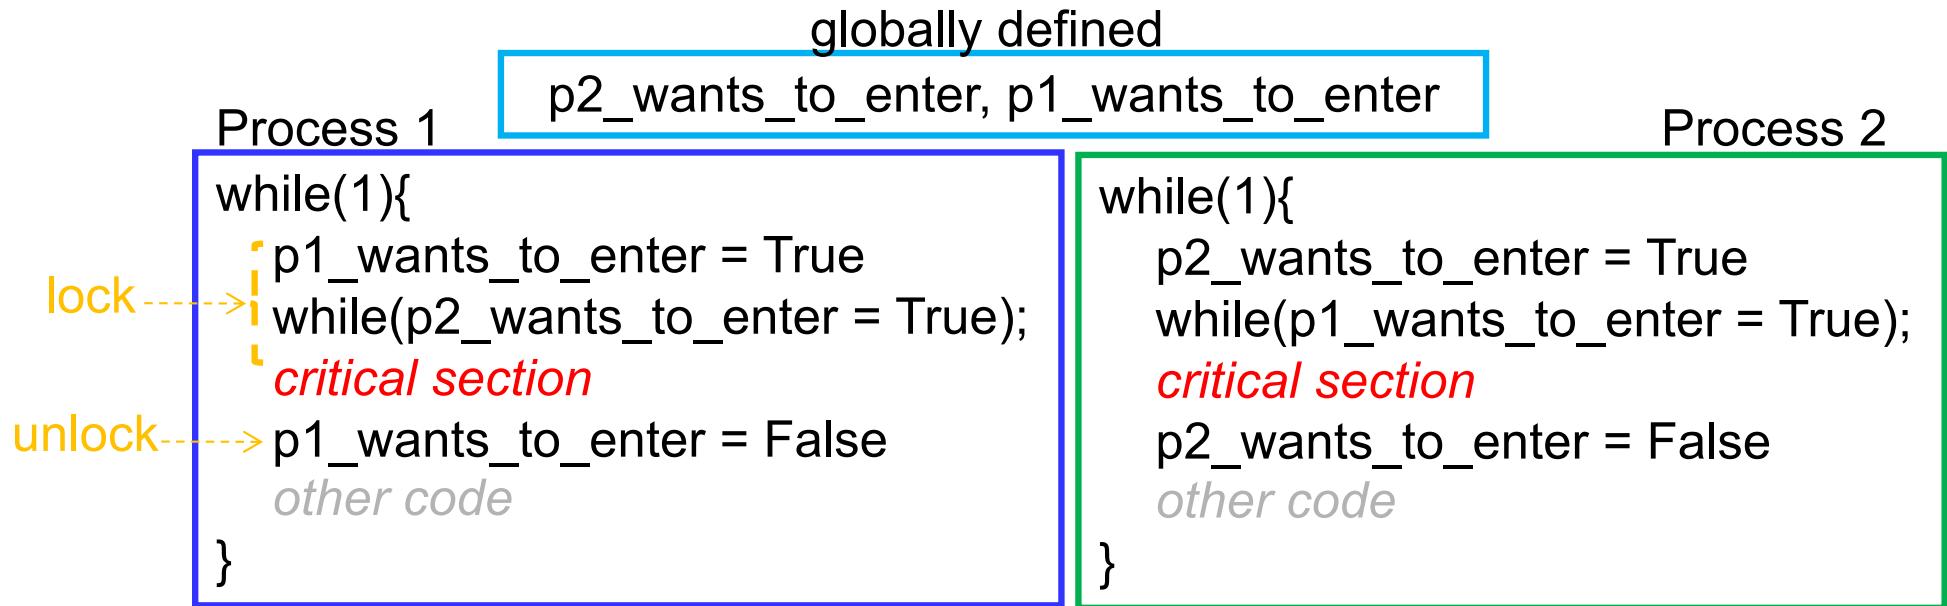
Both p1 and p2 will loop infinitely

Progress not achieved

Each process is waiting for the other  
this is a deadlock

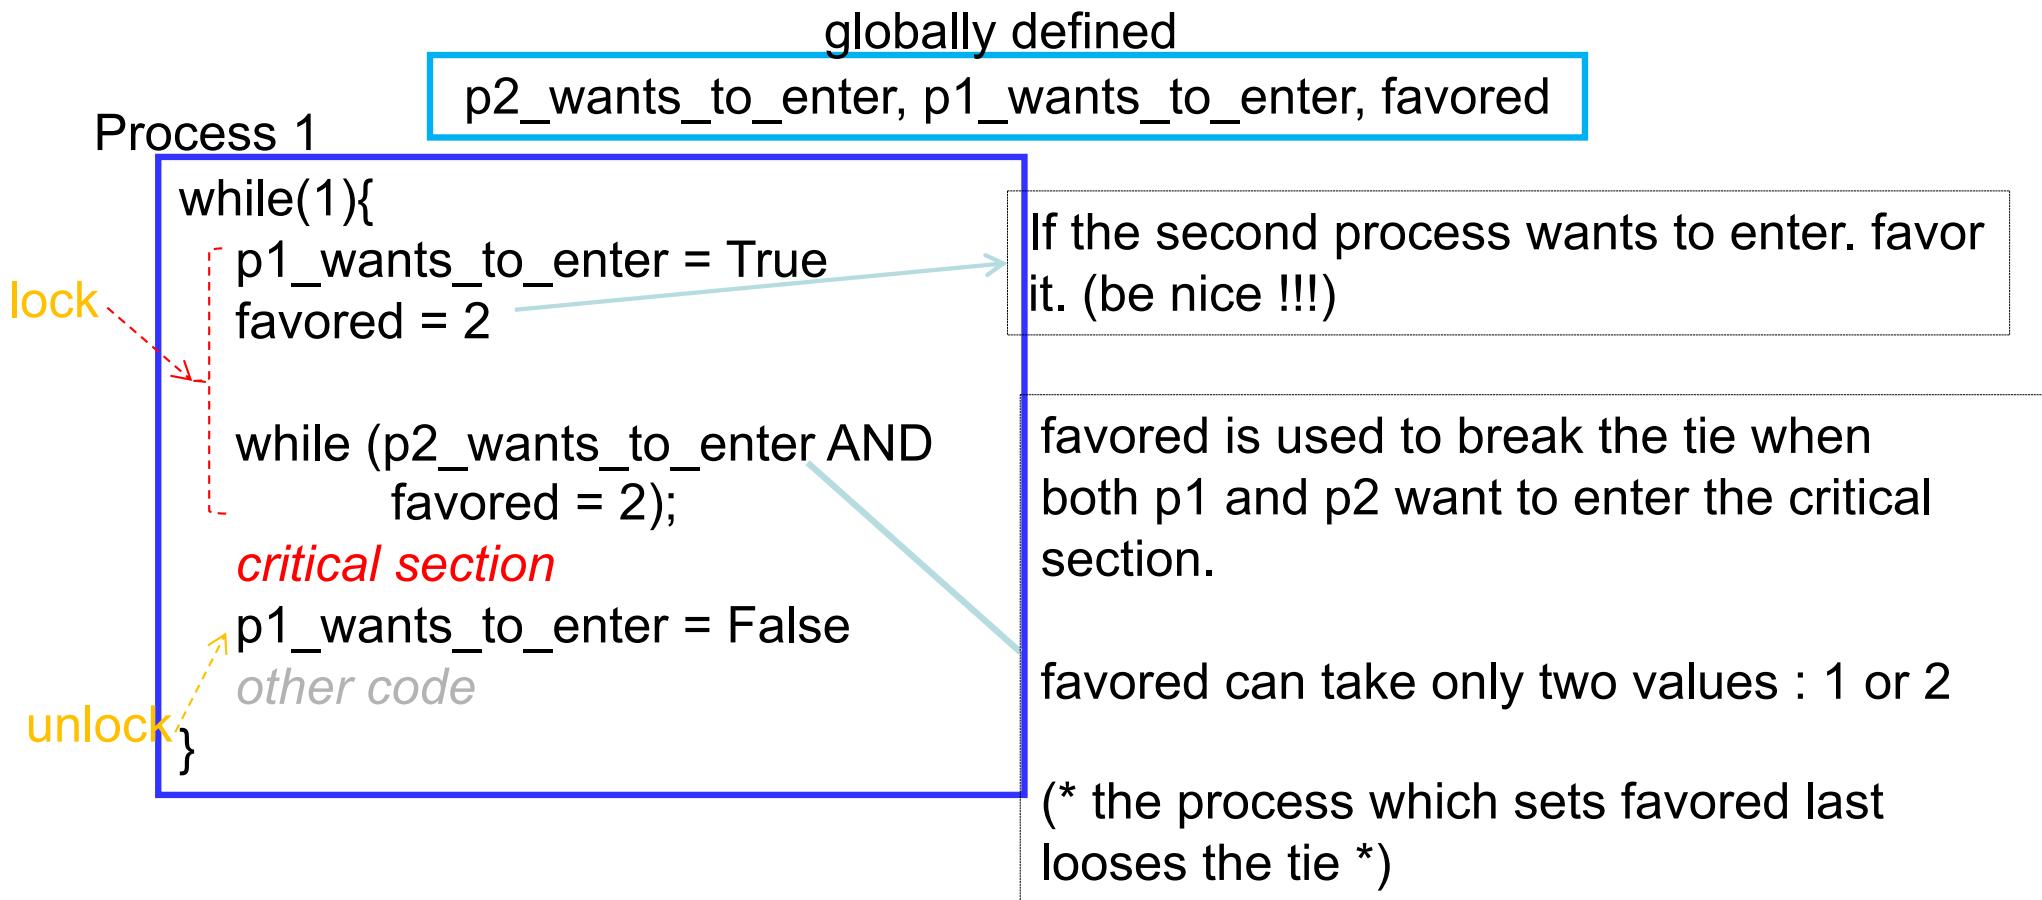


# Problem with Attempt 3



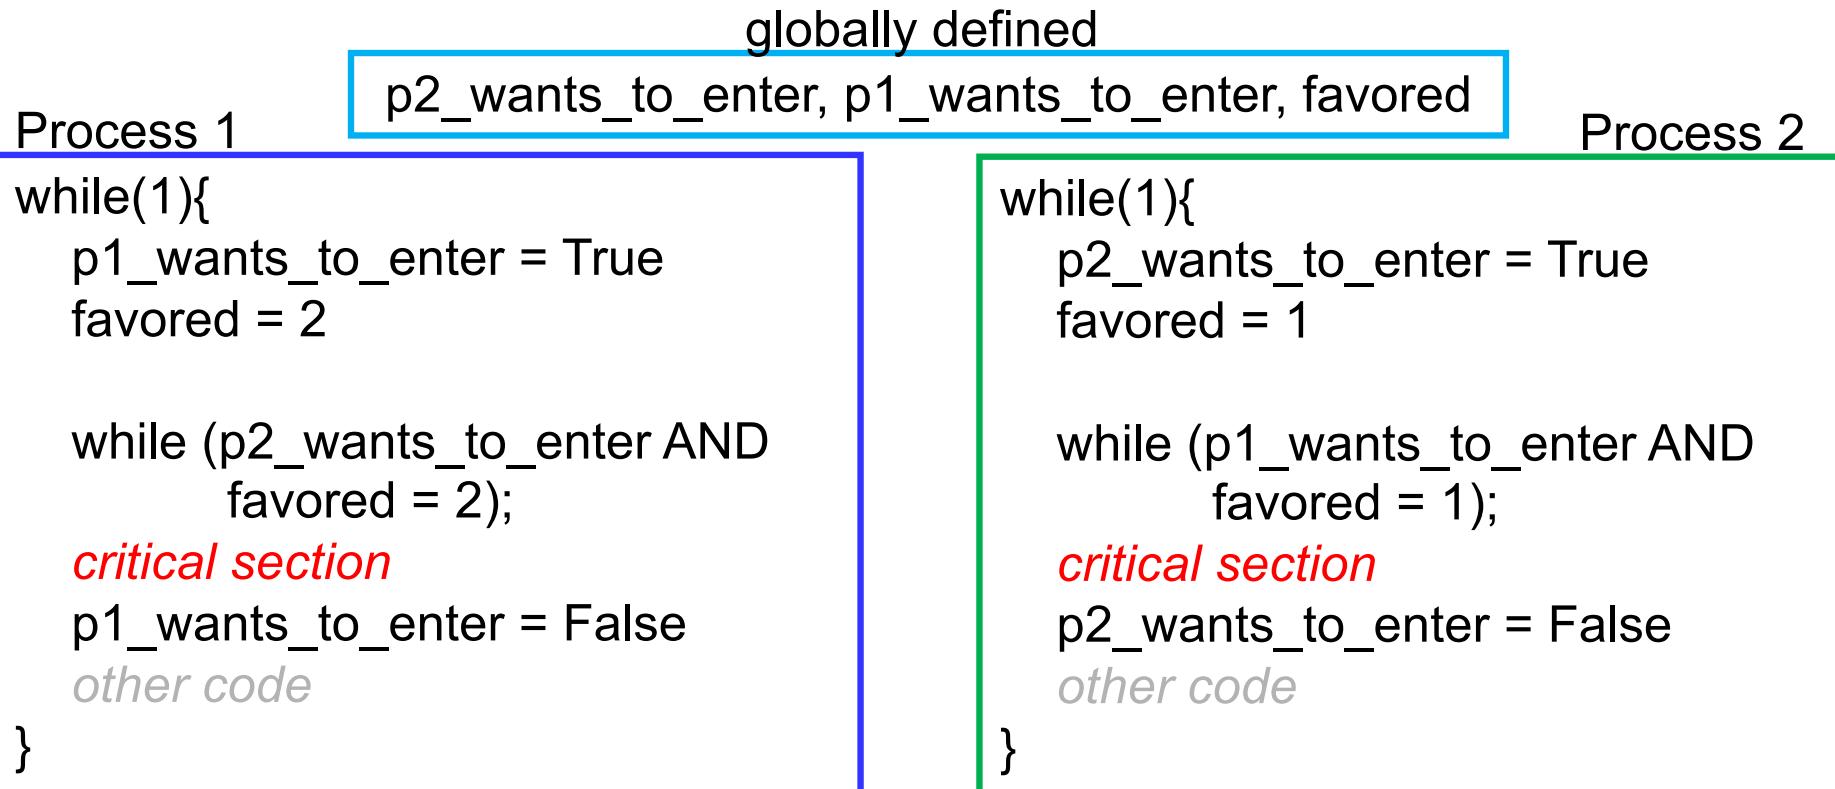
- Deadlock  
Have a way to break the deadlock

# Peterson's Solution



Break the deadlock with a ‘favored’ process

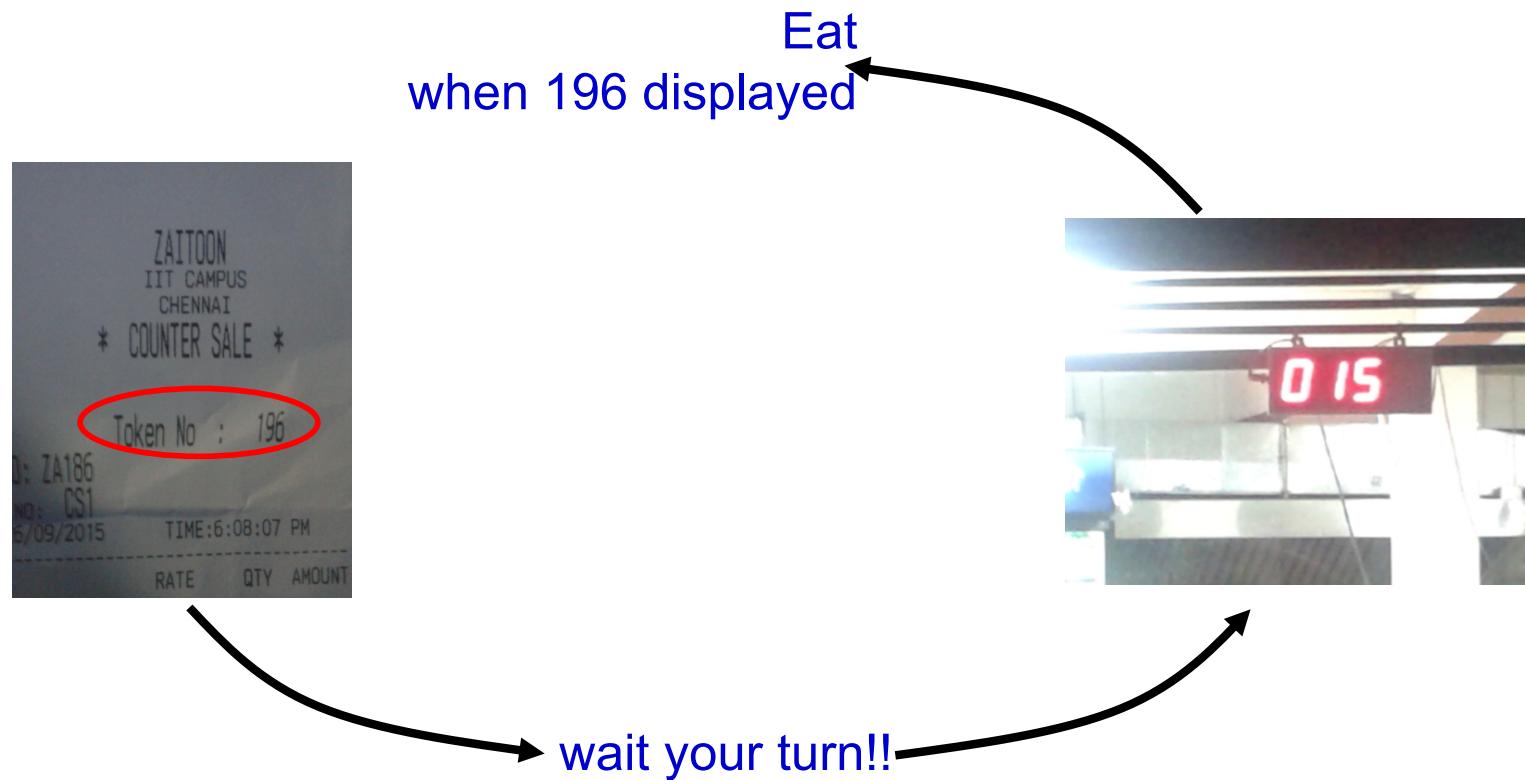
# Peterson's Solution



- Deadlock broken because favored can only be 1 or 2.
  - Therefore, tie is broken. Only one process will enter the critical section
- Solves Critical Section problem for two processes

# Bakery Algorithm

- Synchronization between  $N > 2$  processes
- By Leslie Lamport



# Simplified Bakery Algorithm

- Processes numbered 0 to N-1
- num is an array N integers (initially 0).
  - Each entry corresponds to a process

```
lock(i){  
    num[i] = MAX(num[0], num[1], ..., num[N-1]) + 1  
    for(p = 0; p < N; ++p){  
        while (num[p] != 0 and num[p] < num[i]);  
    }  
}
```

critical section

```
unlock(i){  
    num[i] = 0;  
}
```

This is at the doorway!!!  
It has to be atomic  
to ensure two processes  
do not get the same token

# Simplified Bakery Algorithm (example)

- Processes numbered 0 to N-1
- num is an array N integers (initially 0).
  - Each entry corresponds to a process

```
lock(i){  
    num[i] = MAX(num[0], num[1], ..., num[N-1]) + 1  
    for(p = 0; p < N; ++p){  
        while (num[p] != 0 and num[p] < num[i]);  
    }  
}
```

critical section

	P1	P2	P3	P4	P5
0	0	0	0	2	3

```
unlock(i){  
    num[i] = 0;  
}
```

# Simplified Bakery Algorithm (example)

Processes numbered 0 to N-1  
num is an array N integers (initially 0).  
Each entry corresponds to a process

```
lock(i){  
    num[i] = MAX(num[0], num[1], ..., num[N-1]) + 1  
    for(p = 0; p < N; ++p){  
        while (num[p] != 0 and num[p] < num[i]);  
    }  
}
```

critical section

```
unlock(i){  
    num[i] = 0;  
}
```

P1	P2	P3	P4	P5
0	0	0	0	0

# Simplified Bakery Algorithm (why atomic doorway?)

- Processes numbered 0 to N-1
- num is an array N integers (initially 0).
  - Each entry corresponds to a process

```
lock(i){  
    num[i] = MAX(num[0], num[1], ..., num[N-1]) + 1  
    for(p = 0; p < N; ++p){  
        while (num[p] != 0 and num[p] < num[i]);  
    }  
}
```

critical section

This is at the doorway!!!  
Assume it is not atomic

```
unlock(i){  
    num[i] = 0;  
}
```

P1	P2	P3	P4	P5
0	0	0	0	0

P4 and P5 can enter the critical section at the same time.

# Original Bakery Algorithm (making MAX atomic)

- Without atomic operation assumptions
- Introduce an array of N Booleans: *choosing*, initially all values False.

```
lock(i){  
    choosing[i] = True  
    num[i] = MAX(num[0], num[1], ...., num[N-1]) + 1  
    choosing[i] = False  
    for(p = 0; p < N; ++p){  
        while (choosing[p]);  
        while (num[p] != 0 and (num[p],p)<(num[i],i));  
    }  
}
```

doorway

critical section

```
unlock(i){  
    num[i] = 0;  
}
```

Choosing ensures that a process  
Is not at the doorway  
i.e., the process is not ‘choosing’  
a value for num

$(a, b) < (c, d)$  which is equivalent to:  $(a < c)$  or  $((a == c) \text{ and } (b < d))$

# Original Bakery Algorithm (making MAX atomic)

- Without atomic operation assumptions
- Introduce an array of N Booleans: *choosing*, initially all values False.

```
lock(i){  
    choosing[i] = True  
    num[i] = MAX(num[0], num[1], ...., num[N-1]) + 1  
    choosing[i] = False  
    for(p = 0; p < N; ++p){  
        while (choosing[p]);  
        while (num[p] != 0 and (num[p],p)<(num[i],i));  
    }  
}
```

doorway

critical section

```
unlock(i){  
    num[i] = 0;  
}
```

Favor one process when there is a conflict.

If there are two processes, with the same num value, favor the process with the smaller id (i)

# Original Bakery Algorithm (example)

- Without atomic operation assumptions
- Introduce an array of N Booleans: *choosing*, initially all values False.

```
lock(i){  
    choosing[i] = True  
    num[i] = MAX(num[0], num[1], ...., num[N-1]) + 1  
    choosing[i] = False  
    for(p = 0; p < N; ++p){  
        while (choosing[p]);  
        while (num[p] != 0 and (num[p],p)<(num[i],i));  
    }  
}
```

doorway

critical section

```
unlock(i){  
    num[i] = 0;  
}
```

P1	P2	P3	P4	P5
0	0	0	0	0

$(a, b) < (c, d)$  which is equivalent to:  $(a < c)$  or  $((a == c) \text{ and } (b < d))$

# Original Bakery Algorithm (example)

- Without atomic operation assumptions
- Introduce an array of N Booleans: *choosing*, initially all values False.

```
lock(i){  
    choosing[i] = True  
    num[i] = MAX(num[0], num[1], ...., num[N-1]) + 1  
    choosing[i] = False  
    for(p = 0; p < N; ++p){  
        while (choosing[p]);  
        while (num[p] != 0 and (num[p],p)<(num[i],i));  
    }  
}
```

doorway

critical section

```
unlock(i){  
    num[i] = 0;  
}
```

P1	P2	P3	P4	P5
0	3	1	2	2

$(a, b) < (c, d)$  which is equivalent to:  $(a < c)$  or  $((a == c) \text{ and } (b < d))$

# How to Implement Locking (Hardware Solutions and Usage)

# Analyze this

- Does this scheme provide mutual exclusion?

Process 1

```
while(1){  
    while(lock != 0);  
    lock= 1; // lock  
    critical section  
    lock = 0; // unlock  
    other code  
}
```

lock=0

Process 2

```
while(1){  
    while(lock != 0);  
    lock = 1; // lock  
    critical section  
    lock = 0; // unlock  
    other code  
}
```

No

```
lock = 0  
P1: while(lock != 0);  
P2: while(lock != 0);  
P2: lock = 1;  
P1: lock = 1;  
.... Both processes in critical section
```

context switch

# If only...

- We could make this operation atomic

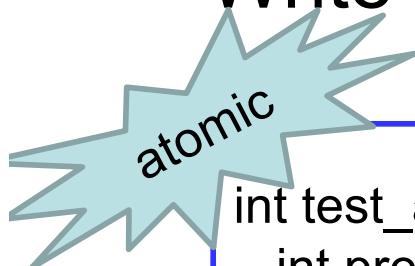
```
Process 1
while(1){
    while(lock != 0);
    lock= 1; // lock
    critical section
    lock = 0; // unlock
    other code
}
```

A blue dashed rectangle encloses the code block starting with 'critical section' and ending with 'lock = 0; // unlock'. A blue arrow points from this rectangle to the text 'Make atomic'.

Hardware to the rescue....

# Hardware Support (Test & Set Instruction)

- Write to a memory location, return its old value



```
int test_and_set(int *L){  
    int prev = *L;  
    *L = 1;  
    return prev;  
}
```

```
while(1){  
    while(test_and_set(&lock) == 1);  
    critical section  
    lock = 0; // unlock  
    other code  
}
```

equivalent software representation  
(the entire function is executed atomically)

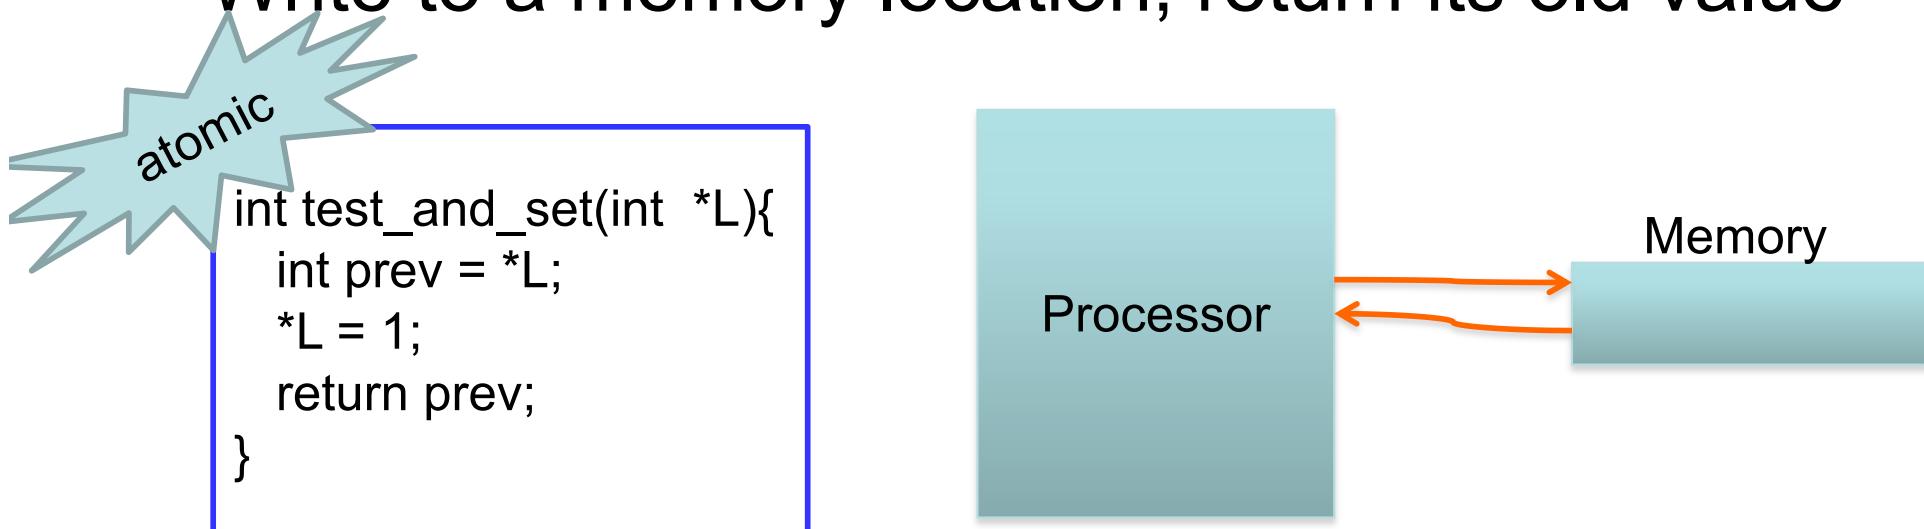
Usage for locking

**Why does this work?** If two CPUs execute `test_and_set` at the same time, the hardware ensures that one `test_and_set` does both its steps before the other one starts.

So the first invocation of `test_and_set` will read a 0 and set lock to 1 and return. The second `test_and_set` invocation will then see lock as 1, and will loop continuously until lock becomes 0

# Hardware Support (Test & Set Instruction)

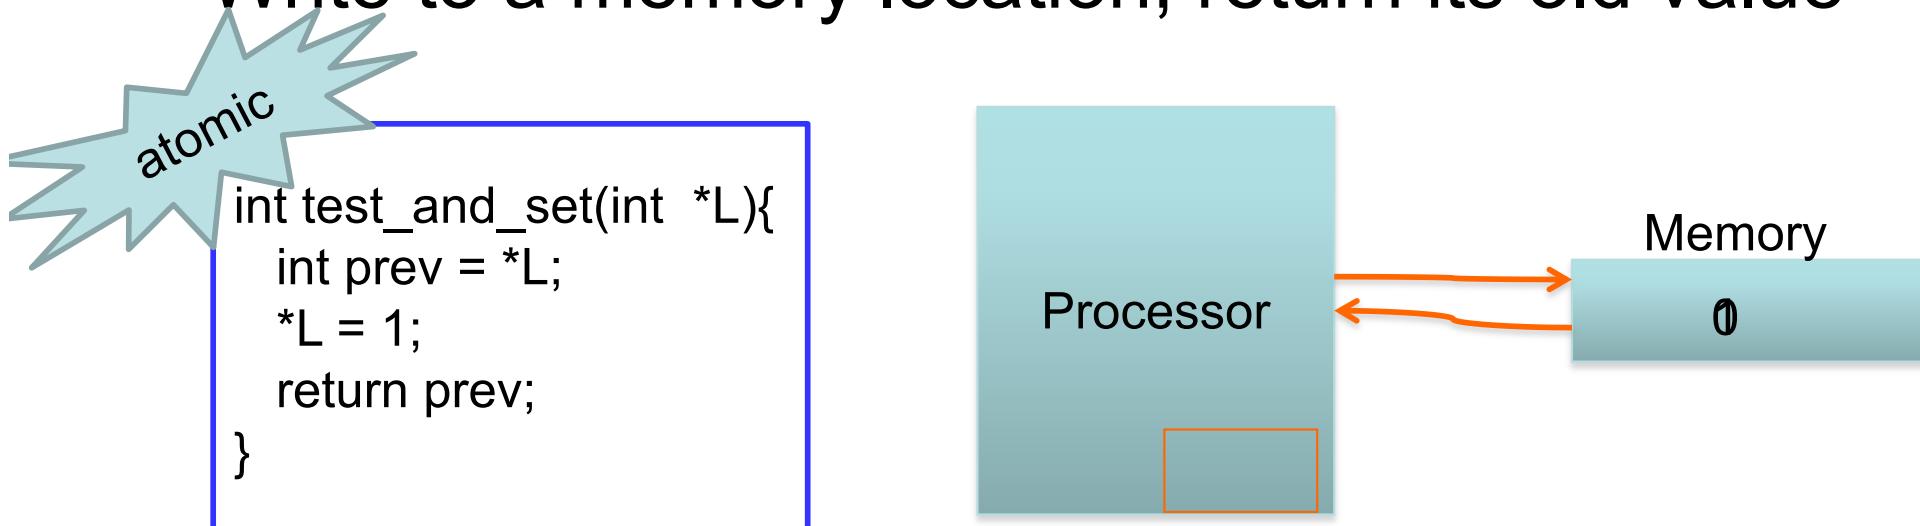
- Write to a memory location, return its old value



equivalent software representation  
(the entire function is executed atomically)

# Hardware Support (Test & Set Instruction)

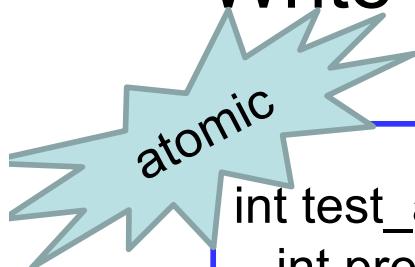
- Write to a memory location, return its old value



equivalent software representation  
(the entire function is executed atomically)

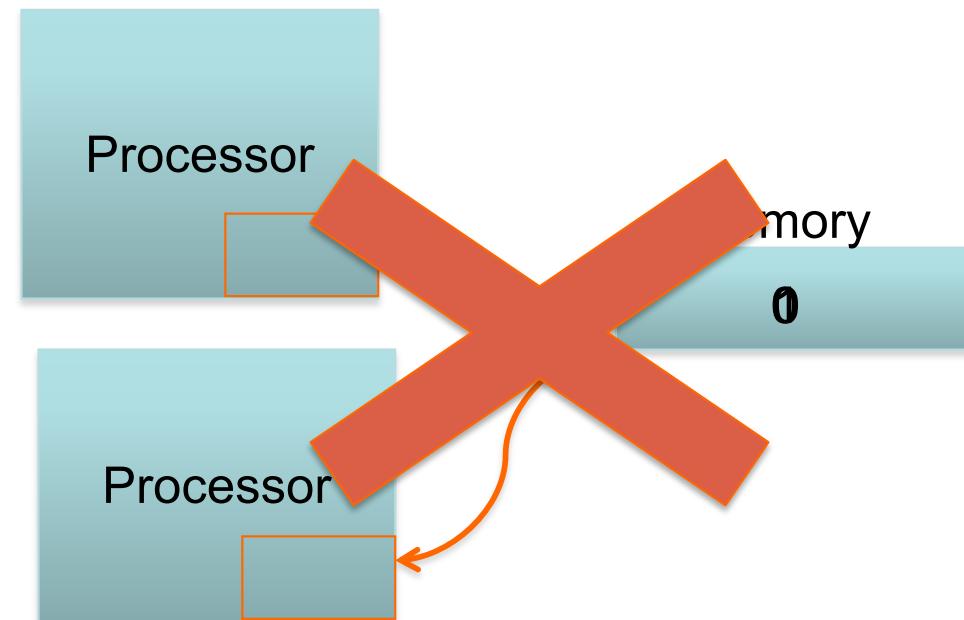
# Hardware Support (Test & Set Instruction)

- Write to a memory location, return its old value



```
int test_and_set(int *L){  
    int prev = *L;  
    *L = 1;  
    return prev;  
}
```

equivalent software representation  
(the entire function is executed  
atomically)



**Why does this work?** If two CPUs execute `test_and_set` at the same time, the hardware ensures that one `test_and_set` does both its steps before the other one starts.

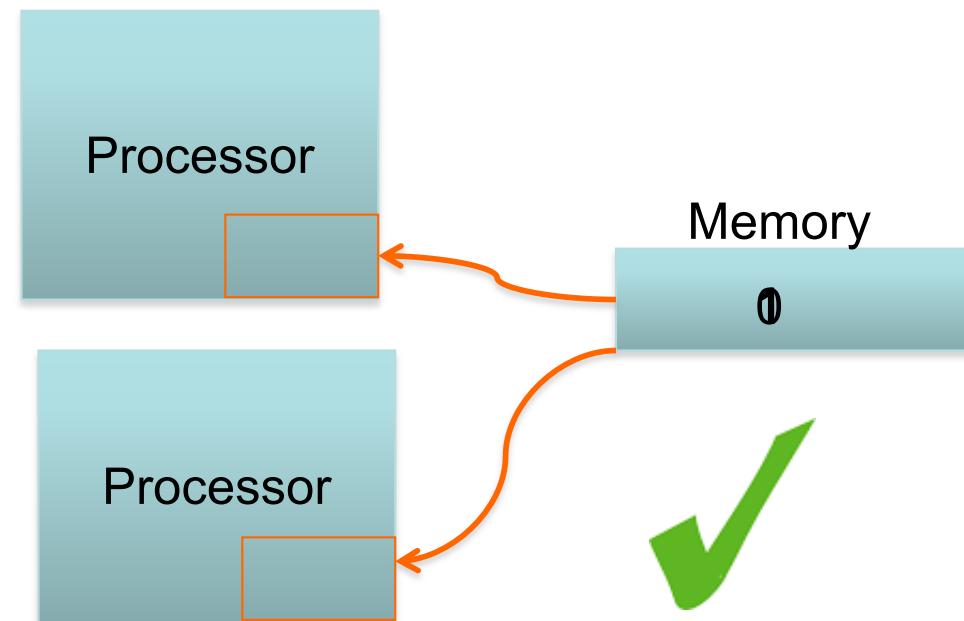
# Hardware Support (Test & Set Instruction)

- Write to a memory location, return its old value



```
int test_and_set(int *L){  
    int prev = *L;  
    *L = 1;  
    return prev;  
}
```

equivalent software representation  
(the entire function is executed  
atomically)



**Why does this work?** If two CPUs execute `test_and_set` at the same time, the hardware ensures that one `test_and_set` does both its steps before the other one starts.

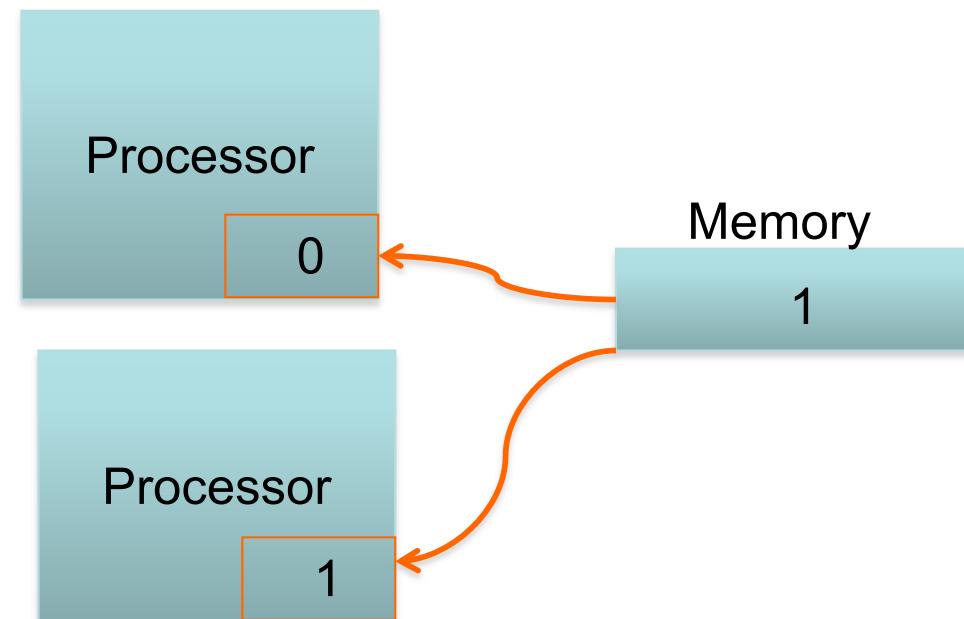
# Hardware Support (Test & Set Instruction)

- Write to a memory location, return its old value



```
int test_and_set(int *L){  
    int prev = *L;  
    *L = 1;  
    return prev;  
}
```

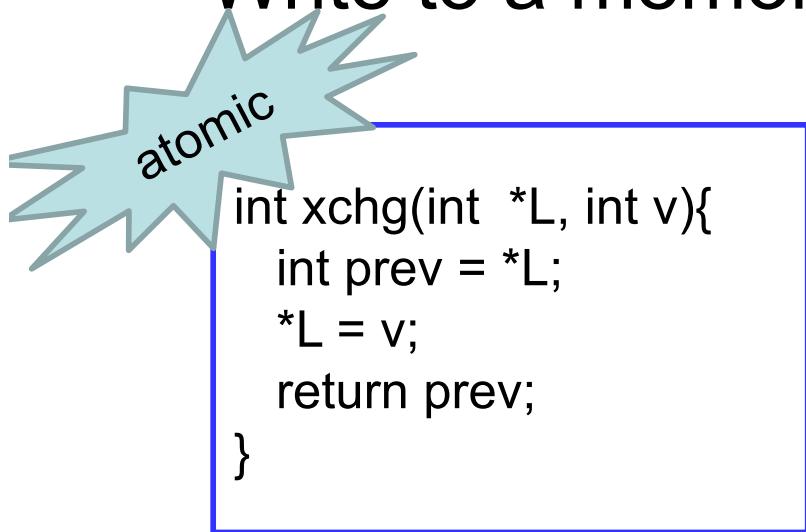
equivalent software representation  
(the entire function is executed  
atomically)



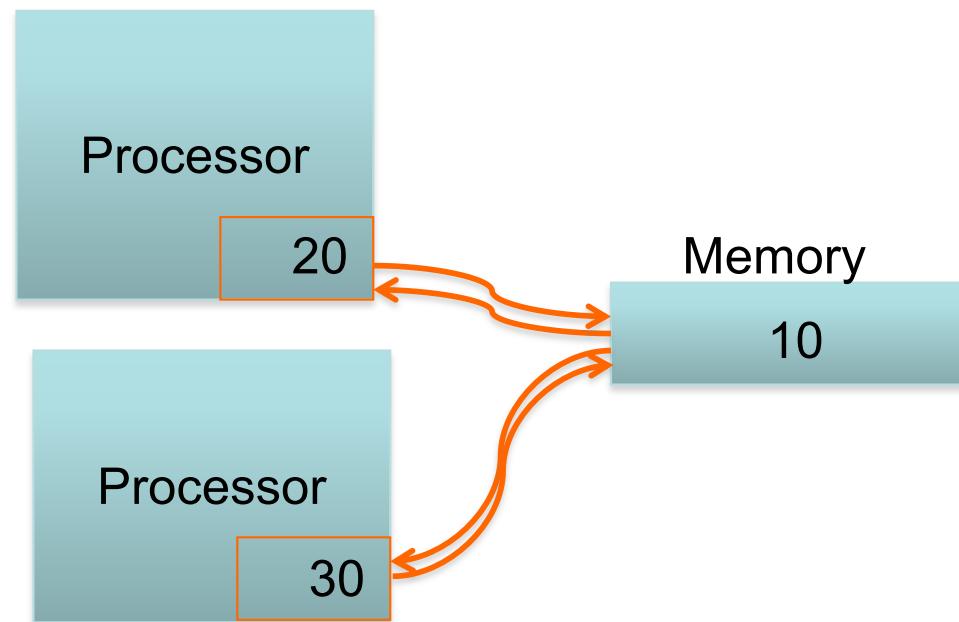
**Why does this work?** If two CPUs execute `test_and_set` at the same time, the hardware ensures that one `test_and_set` does both its steps before the other one starts.

# Intel Hardware Support (xchg Instruction)

- Write to a memory location, return its old value



equivalent software representation  
(the entire function is executed  
atomically)



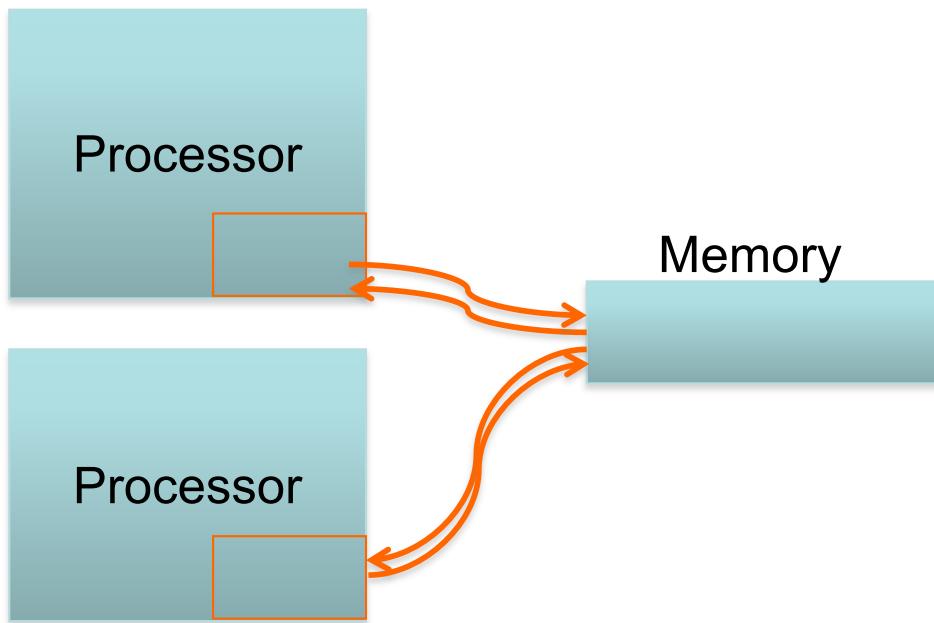
**Why does this work?** If two CPUs execute xchg at the same time, the hardware ensures that one xchg completes, only then the second xchg starts.

# Intel Hardware Support (using xchg instruction)

Note. %eax is returned

typical usage :

xchg reg, mem



```
int xchg(addr, value){  
    %eax = value  
    xchg %eax, (addr)  
}
```

```
void acquire(int *locked){  
    while(1){  
        if(xchg(locked, 1) == 0)  
            break;  
    }  
}
```

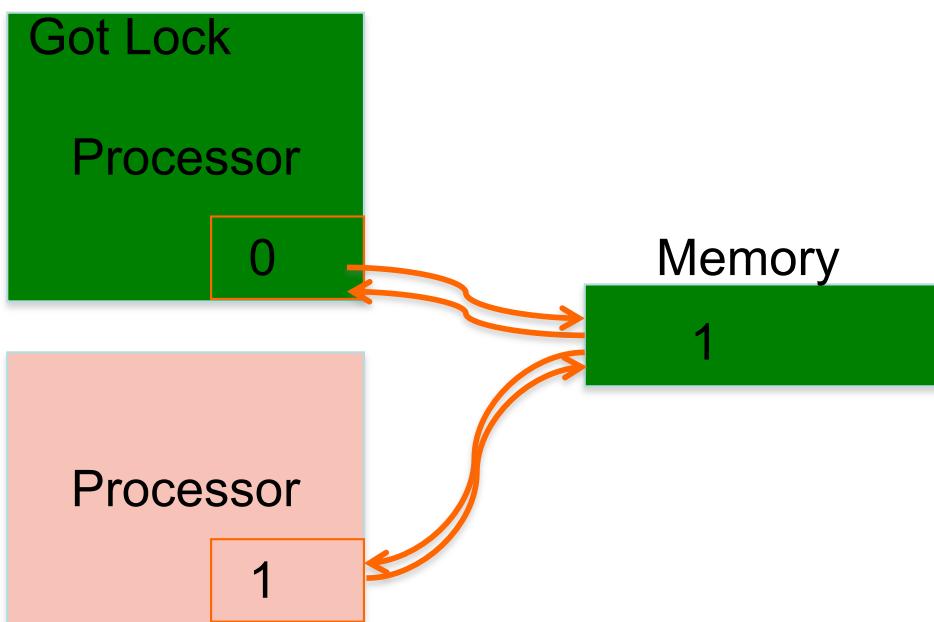
```
void release(int *locked){  
    locked = 0;  
}
```

# Intel Hardware Support (using xchg instruction)

Note. %eax is returned

typical usage :

xchg reg, mem



```
int xchg(addr, value){  
    %eax = value  
    xchg %eax, (addr)  
}
```

```
void acquire(int *locked){  
    while(1){  
        if(xchg(locked, 1) == 0)  
            break;  
    }  
}
```

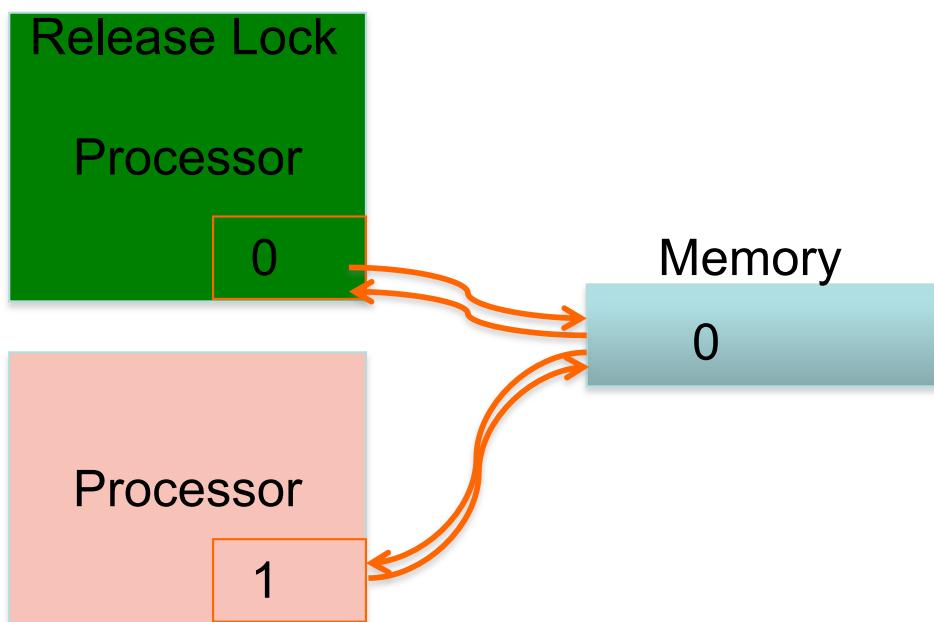
```
void release(int *locked){  
    locked = 0;  
}
```

# Intel Hardware Support (using xchg instruction)

Note. %eax is returned

typical usage :

xchg reg, mem



```
int xchg(addr, value){  
    %eax = value  
    xchg %eax, (addr)  
}
```

```
void acquire(int *locked){  
    while(1){  
        if(xchg(locked, 1) == 0)  
            break;  
    }  
}
```

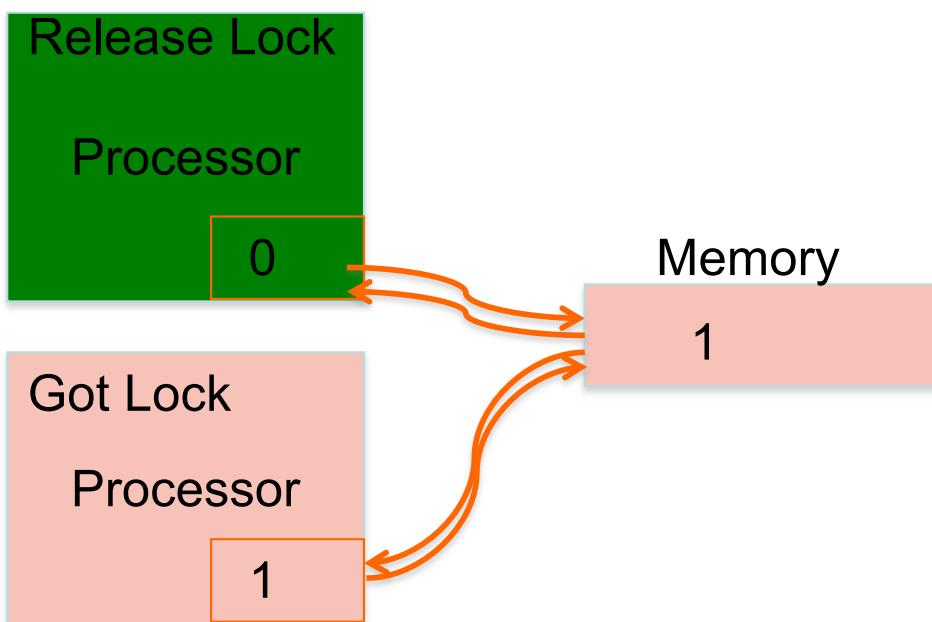
```
void release(int *locked){  
    locked = 0;  
}
```

# Intel Hardware Support (using xchg instruction)

Note. %eax is returned

typical usage :

xchg reg, mem



```
int xchg(addr, value){  
    %eax = value  
    xchg %eax, (addr)  
}
```

```
void acquire(int *locked){  
    while(1){  
        if(xchg(locked, 1) == 0)  
            break;  
    }  
}
```

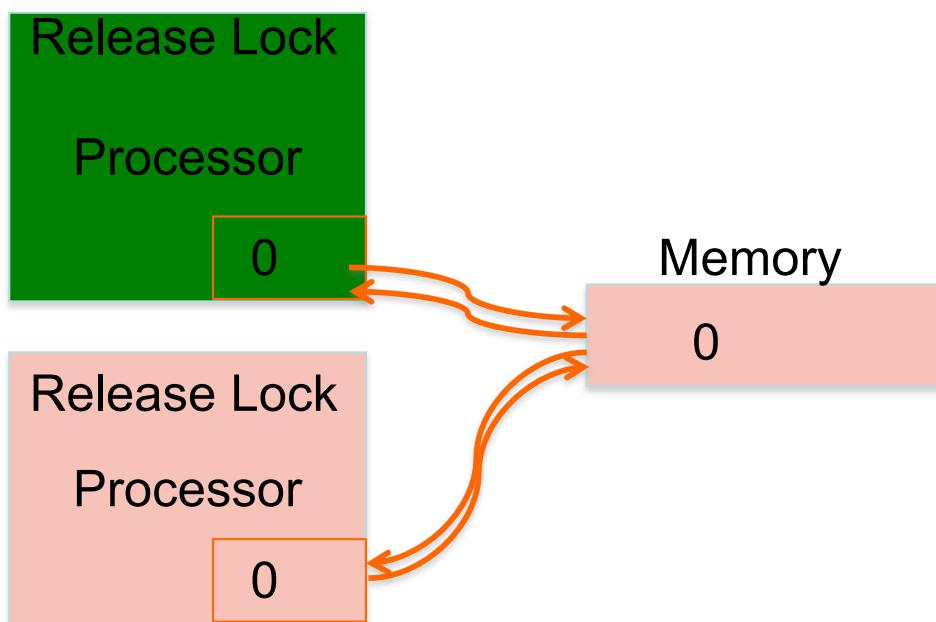
```
void release(int *locked){  
    locked = 0;  
}
```

# Intel Hardware Support (using xchg instruction)

Note. %eax is returned

typical usage :

xchg reg, mem



```
int xchg(addr, value){  
    %eax = value  
    xchg %eax, (addr)  
}
```

```
void acquire(int *locked){  
    while(1){  
        if(xchg(locked, 1) == 0)  
            break;  
    }  
}
```

```
void release(int *locked){  
    locked = 0;  
}
```

# High Level Constructs

- Spinlock
- Mutex
- Semaphore

# Spinlocks Usage

Process 1

```
acquire(&locked)  
critical section  
release(&locked)
```

Process 2

```
acquire(&locked)  
critical section  
release(&locked)
```

- One process will **acquire** the lock
- The other will wait in a loop repeatedly checking if the lock is available
- The lock becomes available when the former process **releases** it

```
int xchg(addr, value){  
    %eax = value  
    xchg %eax, (addr)  
}
```

```
void acquire(int *locked){  
    while(1){  
        if(xchg(locked, 1) == 0)  
            break;  
    }  
}
```

```
void release(int *locked){  
    locked = 0;  
}
```



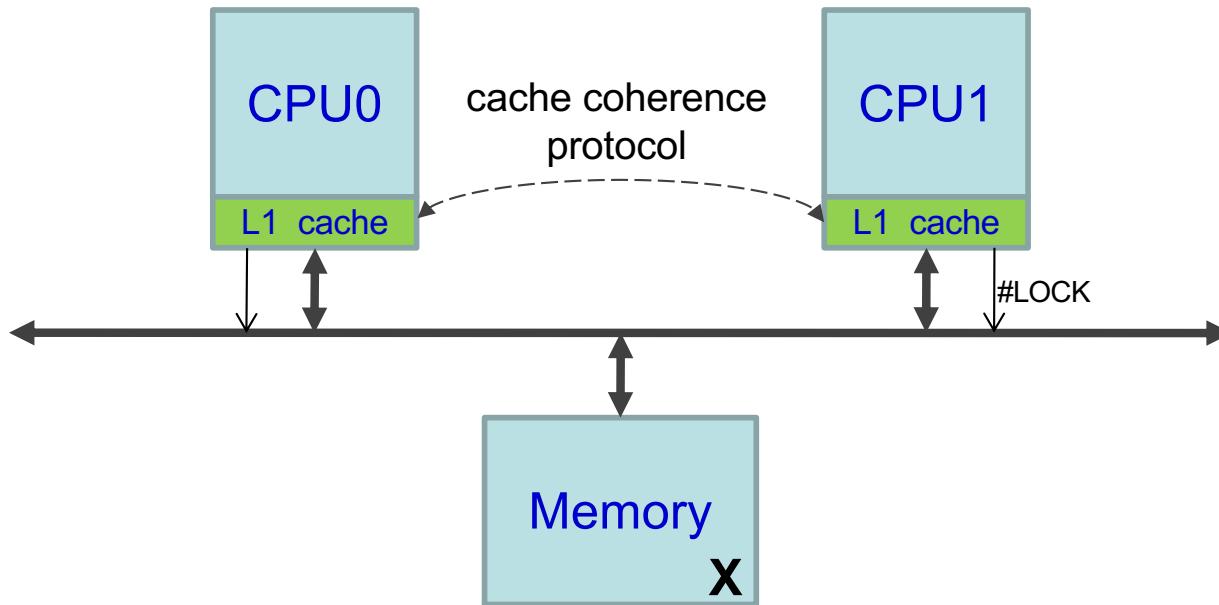
# Issues with Spinlocks

`xchg %eax, X`

- No compiler optimizations should be allowed
  - Should not make X a register variable
    - Write the loop in assembly or use volatile
- Should not reorder **memory** loads and stores
  - Use serialized instructions (which forces instructions not to be reordered)
  - Luckily xchg is already implements serialization

# More issues with Spinlocks

xchg %eax, X



- No caching of (X) possible. All xchg operations are bus transactions.
  - CPU asserts the LOCK, to inform that there is a ‘locked’ memory access
- acquire function in spinlock invokes xchg in a loop...each operation is a bus transaction .... **huge performance hits**

# A better acquire

```
int xchg(addr, value){  
    %eax = value  
    xchg %eax, (addr)  
}
```

```
void acquire(int *locked){  
    reg = 1  
    while(1)  
        if(xchg(locked, reg) == 0)  
            break;  
}
```

Original.

Loop with xchg.

Bus transactions.

Huge overheads

```
void acquire(int *locked) {  
    reg = 1;  
    while (xchg(locked, reg) == 1)  
        while (*locked == 1);  
}
```



Better way

inner loop allows caching of locked. Access cache instead of memory.

# Spinlocks

## (when should it be used?)

- Characteristic : **busy waiting**
  - Useful for short critical sections, where much CPU time is not wasted waiting
    - eg. To increment a counter, access an array element, etc.
  - Not useful, when the period of wait is unpredictable or will take a long time
    - eg. Not good to read page from disk.
    - Use mutex instead (...mutex)

# Spinlock in pthreads

```
#include <pthread.h>
#include <stdio.h>

int global_counter;
pthread_spinlock_t splk;

void *thread_fn(void *arg){
    long id = (long) arg;
    while(1){
        pthread_spin_lock(&splk); → lock
        if (id == 1) global_counter++;
        else global_counter--;
        pthread_spin_unlock(&splk); → unlock
        printf("%d(%d)\n", id, global_counter);
        sleep(1);
    }

    return NULL;
}

int main(){
    pthread_t t1, t2;

    pthread_spin_init(&splk, PTHREAD_PROCESS_PRIVATE); → create spinlock
    pthread_create(&t1, NULL, thread_fn, (void *)1);
    pthread_create(&t2, NULL, thread_fn, (void *)2);
    pthread_join(t1, NULL);
    pthread_join(t2, NULL);
    pthread_spin_destroy(&splk); → destroy spinlock
    printf("Exiting main\n");
    return 0;
}
```

# Mutexes

- Can we do better than busy waiting?
  - If critical section is locked then yield CPU
    - Go to a SLEEP state
  - While unlocking, wake up sleeping process

```
int xchg(addr, value){  
    %eax = value  
    xchg %eax, (addr)  
}  
  
void lock(int *locked){  
    while(1){  
        if(xchg(locked, 1) == 0)  
            break;  
        else  
            sleep();  
    }  
}  
  
void unlock(int *locked){  
    locked = 0;  
    wakeup();  
}
```

# Thundering Herd Problem

- A large number of processes wake up (almost simultaneously) when the event occurs.
  - All waiting processes wake up
  - Leading to several context switches
  - All processes go back to sleep except for one, which gets the critical section
    - Large number of context switches
    - Could lead to starvation

```
int xchg(addr, value){  
    %eax = value  
    xchg %eax, (addr)  
}  
  
void lock(int *locked){  
    while(1){  
        if(xchg(locked, 1) == 0)  
            break;  
        else  
            sleep();  
    }  
}  
  
void unlock(int *locked){  
    locked = 0;  
    wakeup();  
}
```

# Thundering Herd Problem

- The Solution
  - When entering critical section, push into a queue before blocking
  - When exiting critical section, wake up only the first process in the queue

```
int xchg(addr, value){  
    %eax = value  
    xchg %eax, (addr)  
}  
  
void lock(int *locked){  
    while(1){  
        if(xchg(locked, 1) == 0)  
            break;  
        else{  
            // add this process to Queue  
            sleep();  
        }  
    }  
}  
  
void unlock(int *locked){  
    locked = 0;  
    // remove process P from queue  
    wakeup(P)  
}
```

# pthread Mutex

- `pthread_mutex_lock`
- `pthread_mutex_unlock`

# Locks and Priorities

- What happens when a high priority task requests a lock, while a low priority task is in the critical section
  - Priority Inversion
  - Possible solution
    - Priority Inheritance

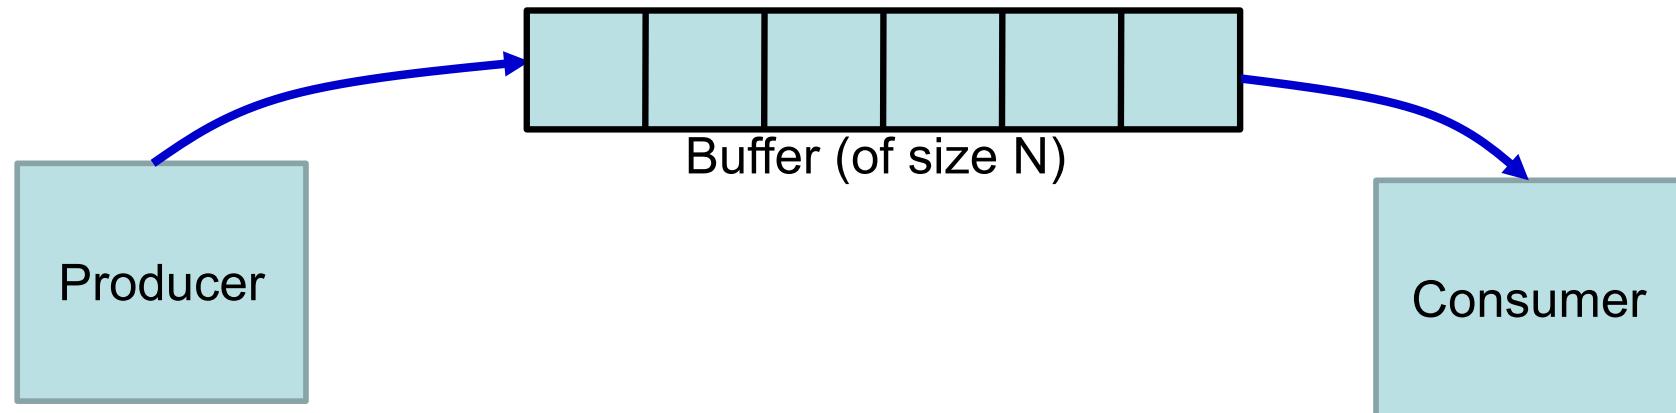
Interesting Read : Mass Pathfinder

[http://research.microsoft.com/en-us/um/people/mbj/mars\\_pathfinder/mars\\_pathfinder.html](http://research.microsoft.com/en-us/um/people/mbj/mars_pathfinder/mars_pathfinder.html)

# Semaphores

# Producer – Consumer Problems

- Also known as *Bounded buffer Problem*
- Producer produces and stores in buffer, Consumer consumes from buffer
- Trouble when
  - Producer produces, but buffer is full
  - Consumer consumes, but buffer is empty



# Producer-Consumer Code

Buffer of size N

```
int count=0;  
Mutex mutex, empty, full;
```

```
1 void producer(){  
2     while(TRUE){  
3         item = produce_item();  
4         if (count == N) sleep(empty);  
5         lock(mutex);  
6         insert_item(item); // into buffer  
7         count++;  
8         unlock(mutex);  
9         if (count == 1) wakeup(full);  
10    }  
}
```

```
1 void consumer(){  
2     while(TRUE){  
3         if (count == 0) sleep(full);  
4         lock(mutex);  
5         item = remove_item(); // from buffer  
6         count--;  
7         unlock(mutex);  
8         if (count == N-1) wakeup(empty);  
9         consume_item(item);  
10    }  
}
```

# Lost Wakeups

- Consider the following context of instructions
- Assume buffer is initially empty

context switch

```
-3 read count value // count < 0  
3 item = produce_item();  
5 lock(mutex);  
6 insert_item(item); // into buffer  
7 count++; // count = 1  
8 unlock(mutex)  
9 test (count == 1) // yes  
9 signal(full);  
3 test (count == 0) // yes  
3 wait();
```

Note, the wakeup is lost.

Consumer waits even though buffer is not empty.

**Eventually producer and consumer will wait infinitely**

consumer

still uses the old value of count (ie 0)

# Semaphores

- Proposed by Dijkstra in 1965
- Functions **down** and **up** must be atomic
- **down** also called **P** (Proberen Dutch for try)
- **up** also called **V** (Verhogen, Dutch form make higher)
- Can have different variants
  - Such as blocking, non-blocking
- If S is initially set to 1,
  - Blocking semaphore similar to a Mutex
  - Non-blocking semaphore similar to a spinlock

```
void down(int *S){  
    while( *S <= 0);  
    *S--;  
}  
  
void up(int *S){  
    *S++;  
}
```

# Producer-Consumer with Semaphores

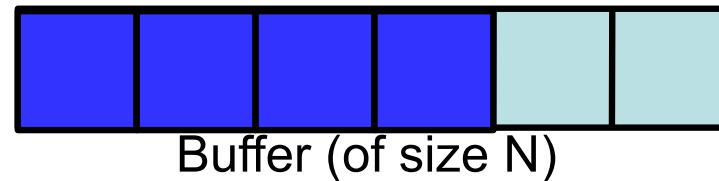
Buffer of size N

int count;

**full = 0, empty = N**

Number of filled blocks in  
the buffer

Number of empty blocks in  
the buffer

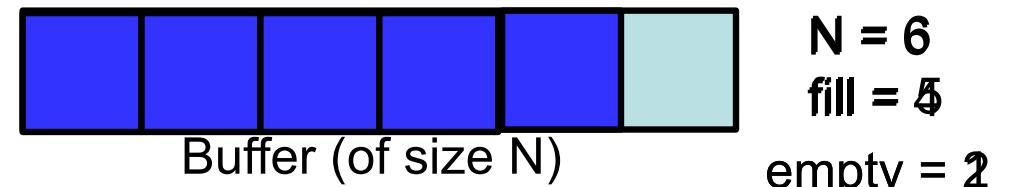


N = 6  
fill = 4  
empty = 2

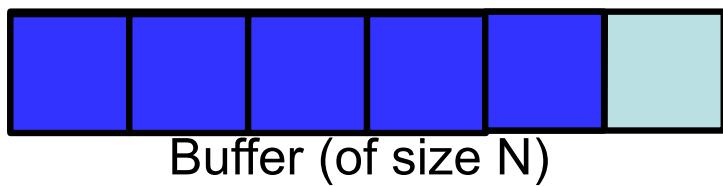
# Producer-Consumer with Semaphores

**full = 0, empty = N**

```
void producer(){
    while(TRUE){
        item = produce_item();
        down(empty);
        insert_item(item); // into buffer
        up(full);
    }
}
```



# Producer-Consumer with Semaphores



$N = 6$   
 $\text{fill} = 5$   
 $\text{empty} = 2$

```
void consumer(){
    while(TRUE){
        → down(full);
        → item = remove_item(); // from buffer
        → up(empty);
        → consume_item(item);
    }
}
```

# Producer-Consumer with Semaphores

## The FULL Buffer

```
void producer(){
    while(TRUE){
        item = produce_item();
        down(empty);

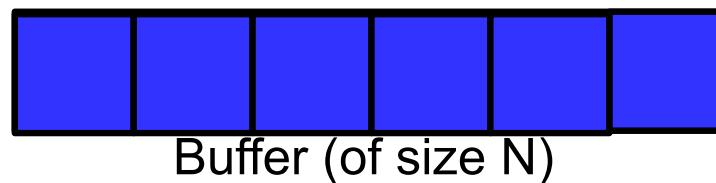
        insert_item(item); // into buffer

        up(full);
    }
}
```

```
void consumer(){
    while(TRUE){
        down(full);

        item = remove_item(); // from buffer

        up(empty);
        consume_item(item);
    }
}
```



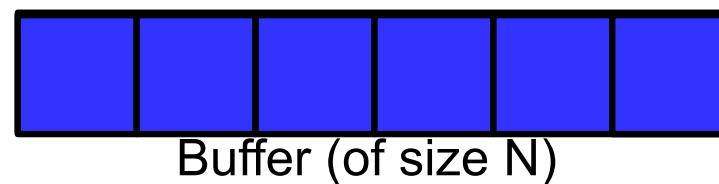
$N = 6$   
 $\text{fill} = 6$   
 $\text{empty} = 0$

# Producer-Consumer with Semaphores

## The FULL Buffer

```
void producer(){
    while(TRUE){
        item = produce_item();
        → down(empty);
        → insert_item(item); // into buffer
        → up(full);
    }
}
```

```
void consumer(){
    while(TRUE){
        down(full);
        item = remove_item(); // from buffer
        up(empty);
        consume_item(item);
    }
}
```



$N = 6$   
 $\text{fill} = 6$   
 $\text{empty} = 0$

# Producer-Consumer with Semaphores

## The Empty Buffer

```
void producer(){
    while(TRUE){
        item = produce_item();
        down(empty);

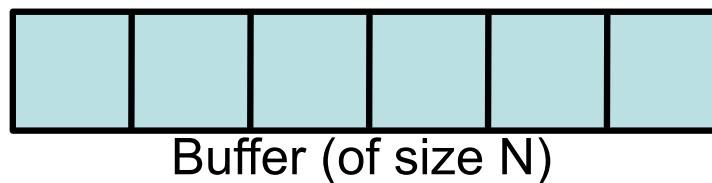
        insert_item(item); // into buffer

        up(full);
    }
}
```

```
void consumer(){
    while(TRUE){
        down(full);

        item = remove_item(); // from buffer

        up(empty);
        consume_item(item);
    }
}
```



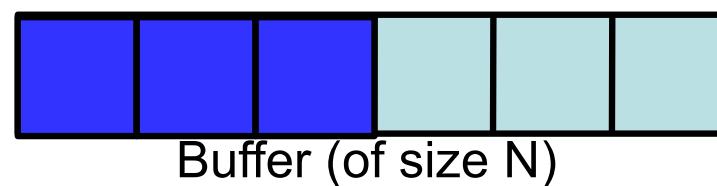
$N = 6$   
 $\text{fill} = 0$   
 $\text{empty} = 6$

# Producer-Consumer with Semaphores

## Serializing Access to the Buffer

```
void producer(){
    while(TRUE){
        item = produce_item();
        down(empty);
        lock(mutex)
        insert_item(item); // into buffer
        unlock(mutex)
        up(full);
    }
}
```

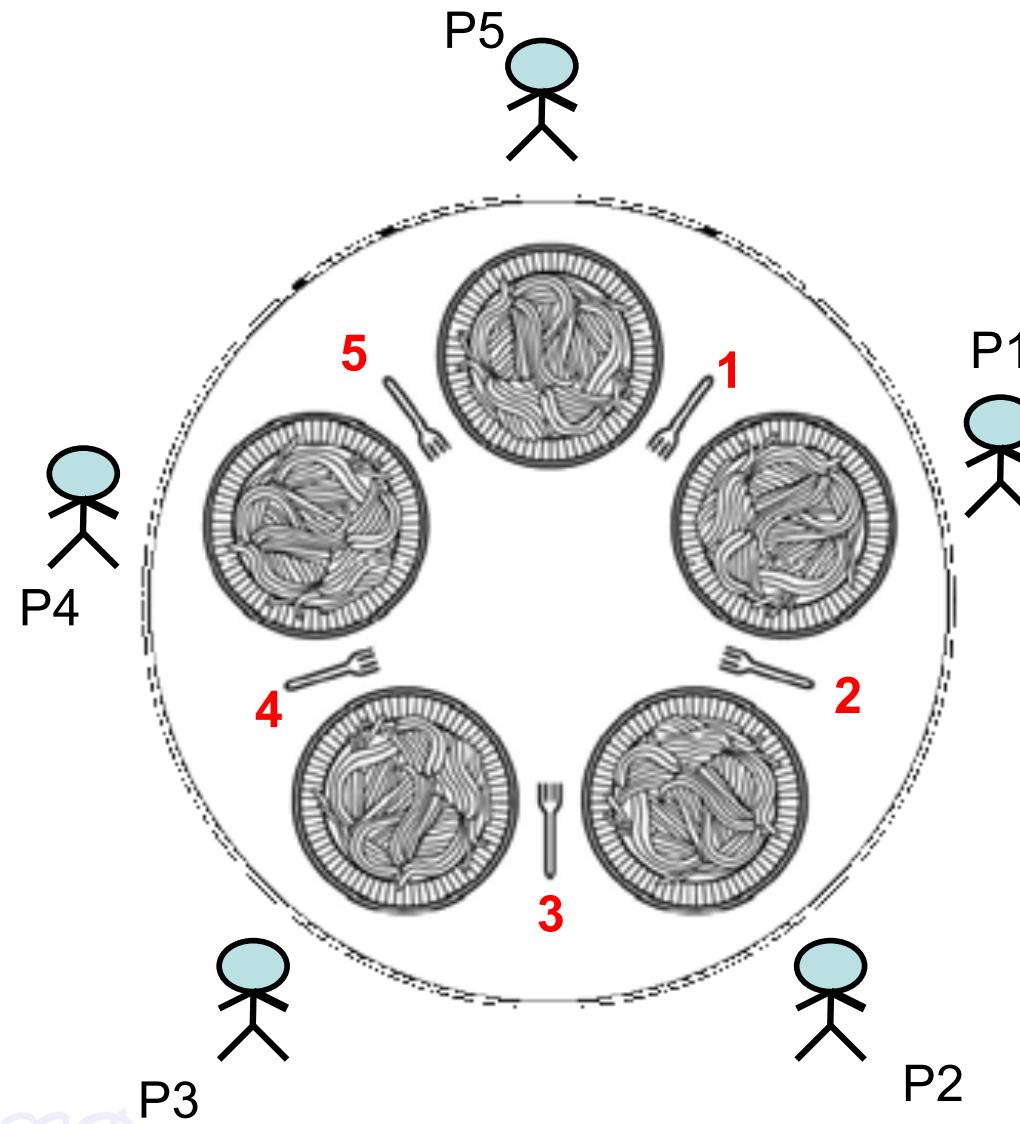
```
void consumer(){
    while(TRUE){
        down(full);
        lock(mutex)
        item = remove_item(); // from buffer
        unlock(mutex)
        up(empty);
        consume_item(item);
    }
}
```



N = 6  
fill = 3  
empty = 3

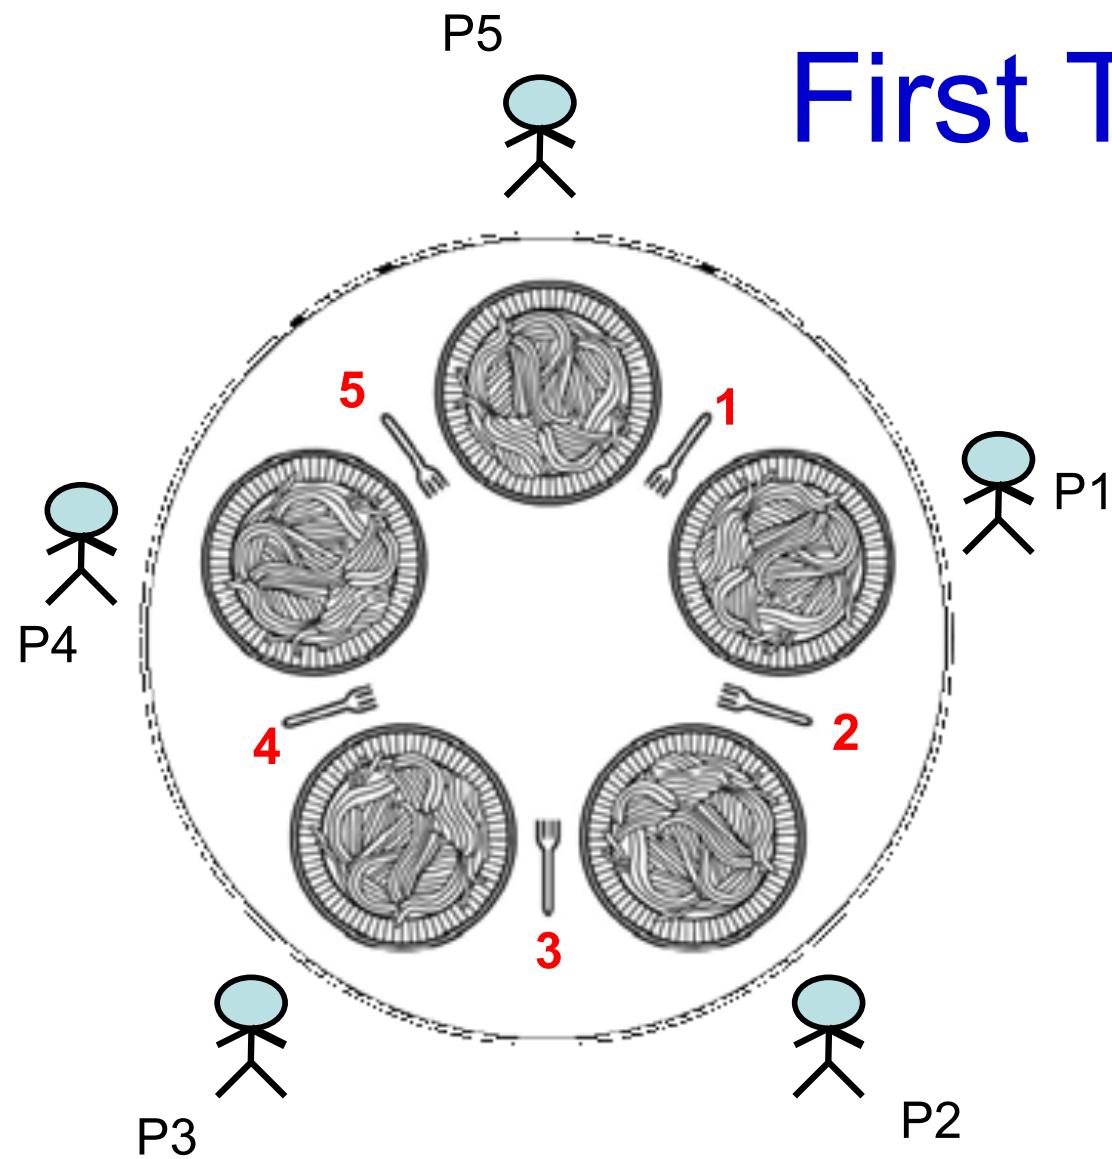
# Dining Philosophers Problem

# Dining Philosophers Problem



- **Philosophers either think or eat**
- To eat, a philosopher needs to hold both forks (the one on his left and the one on his right)
- If the philosopher is not eating, he is thinking.
- **Problem Statement :** Develop an algorithm where no philosopher starves.

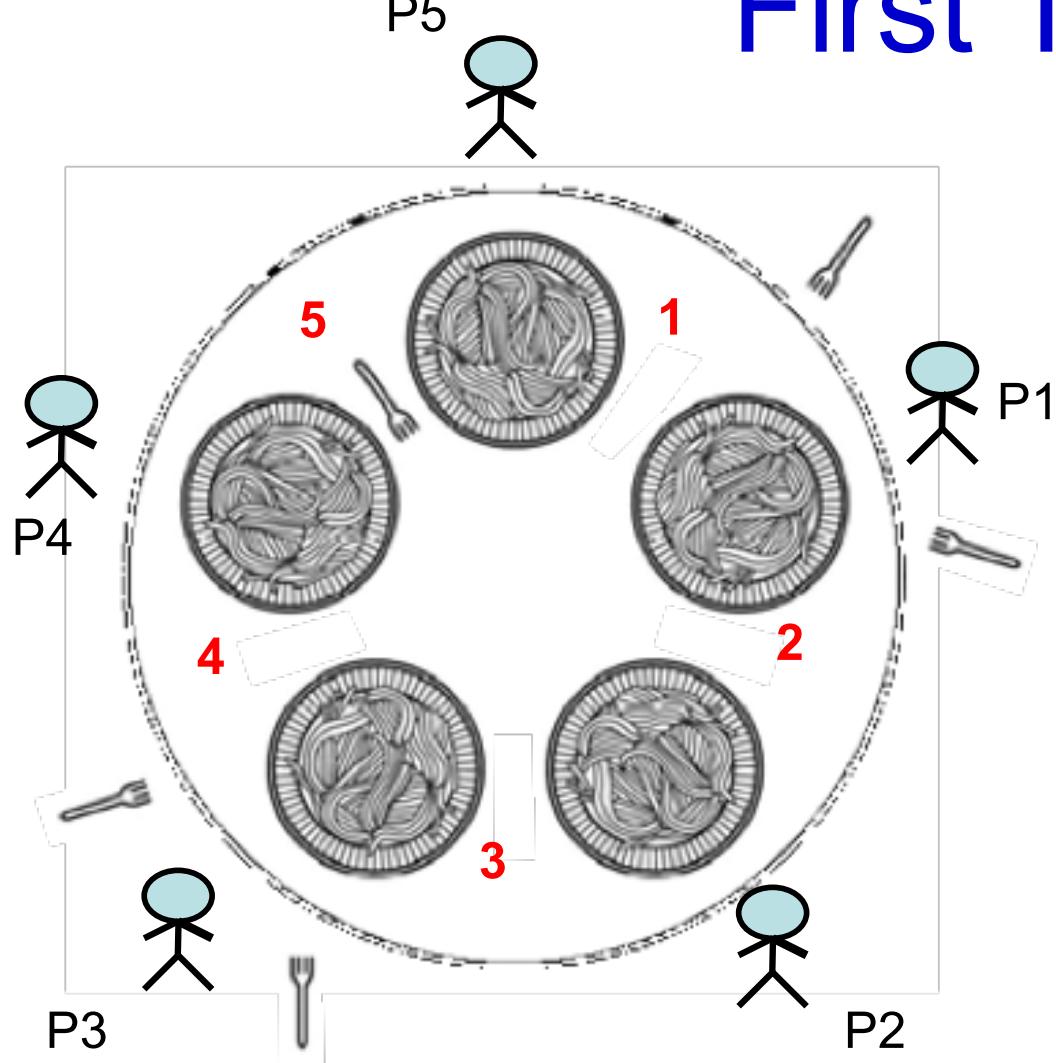
# First Try



```
#define N 5

void philosopher(int i){
    while(TRUE){
        think(); // for some_time
        take_fork(Ri);
        take_fork(Li);
        eat();
        put_fork(Ri);
        put_fork(Li);
    }
}
```

# First Try

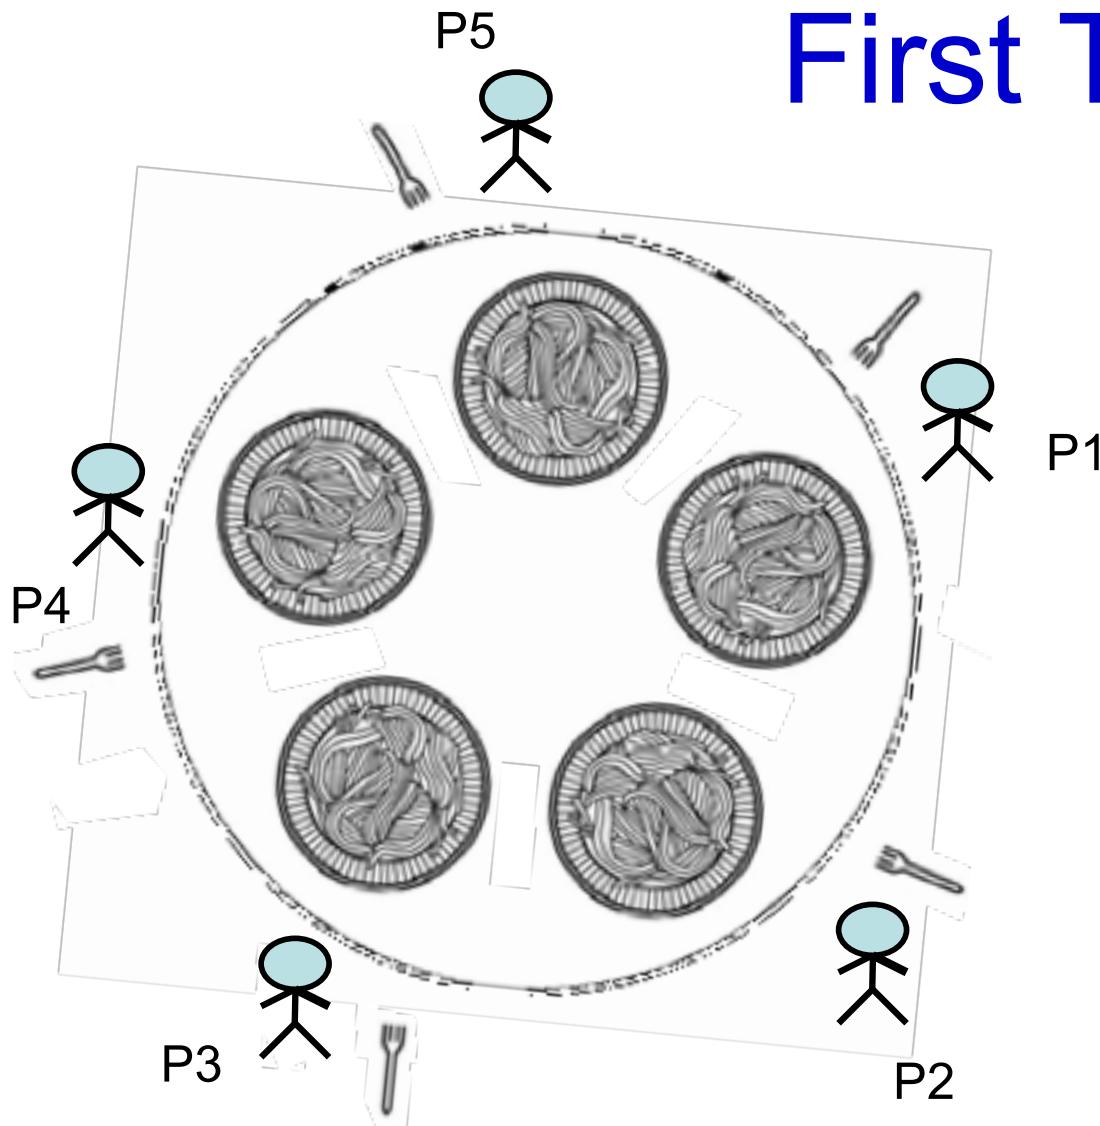


```
#define N 5
```

```
void philosopher(int i){  
    while(TRUE){  
        think(); // for some_time  
        take_fork(Ri);  
        take_fork(Li);  
        eat();  
        put_fork(Ri);  
        put_fork(Li);  
    }  
}
```

What happens if only philosophers P1 and P3 are always given the priority?  
P4, P5, and P2 starves... so scheme needs to be fair

# First Try



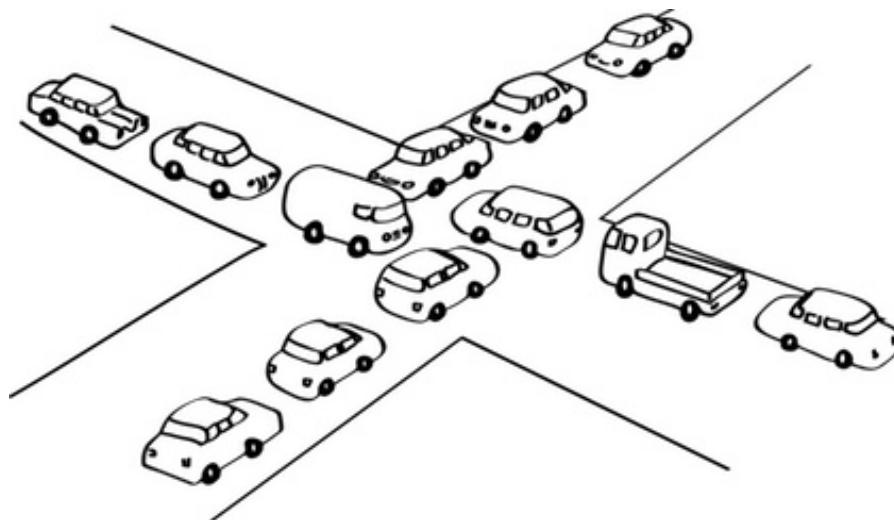
```
#define N 5
```

```
void philosopher(int i){  
    while(TRUE){  
        think(); // for some_time  
        take_fork(Ri);  
        take_fork(Li);  
        eat();  
        put_fork(Ri);  
        put_fork(Li);  
    }  
}
```

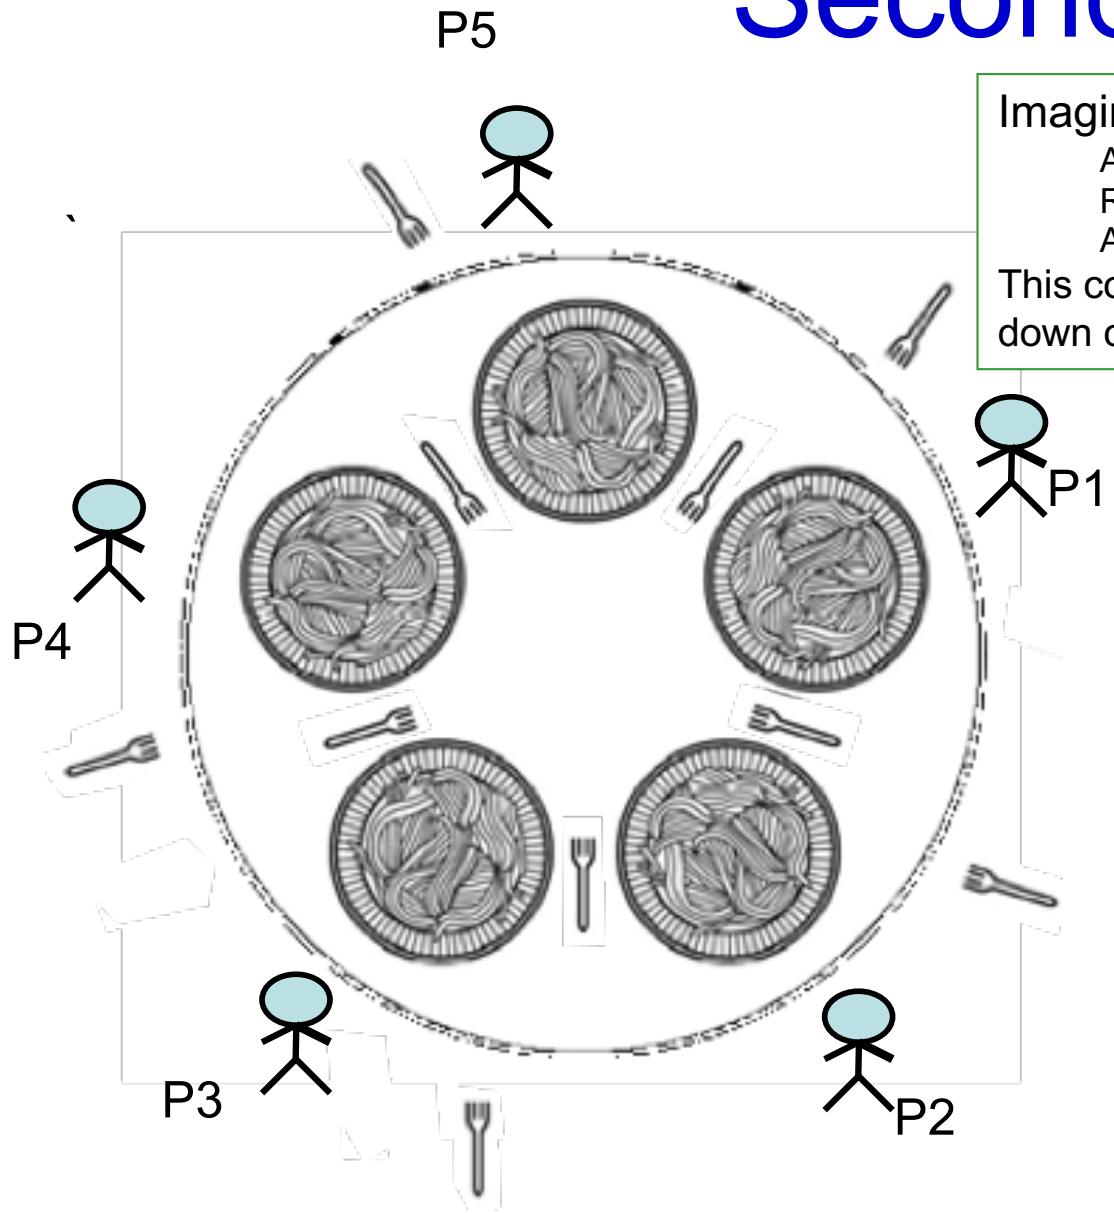
What happens if all philosophers decide to pick up their right forks at the same time?  
Possible starvation due to deadlock

# Deadlocks

- A situation where programs continue to run indefinitely without making any progress
- Each program is waiting for an event that another process can cause



# Second try



Imagine,

All philosophers start at the same time

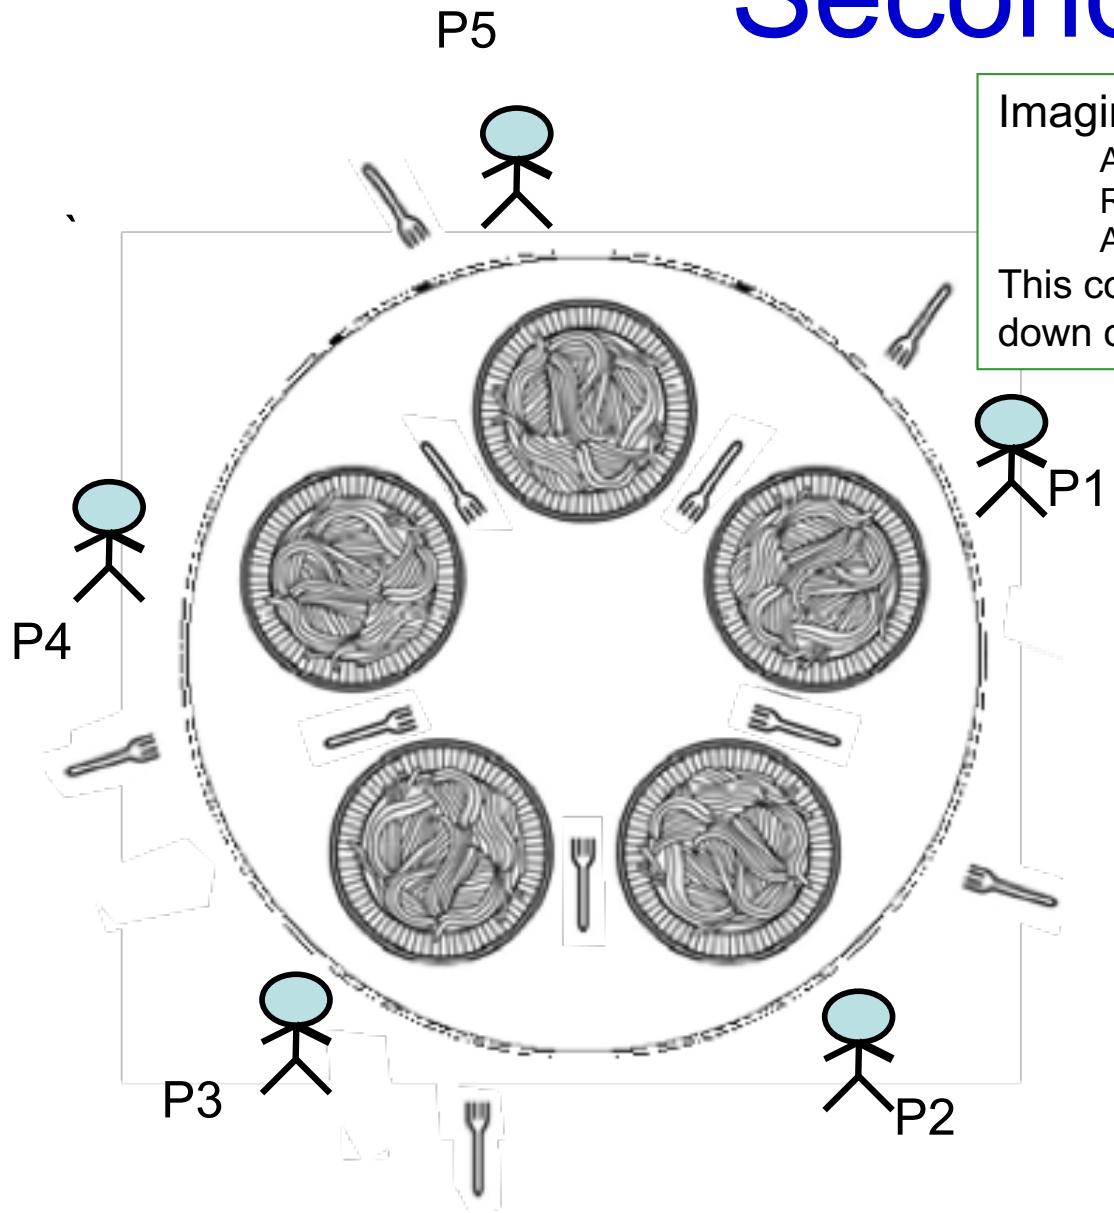
Run simultaneously

And think for the same time

This could lead to philosophers taking fork and putting it down continuously. a deadlock.

```
while(TRUE){  
    think();  
    take_fork(Ri);  
    if (available(Li)){  
        take_fork(Li);  
        eat();  
    }else{  
        put_fork(Ri);  
        sleep(T)  
    }  
}
```

# Second try



Imagine,

All philosophers start at the same time

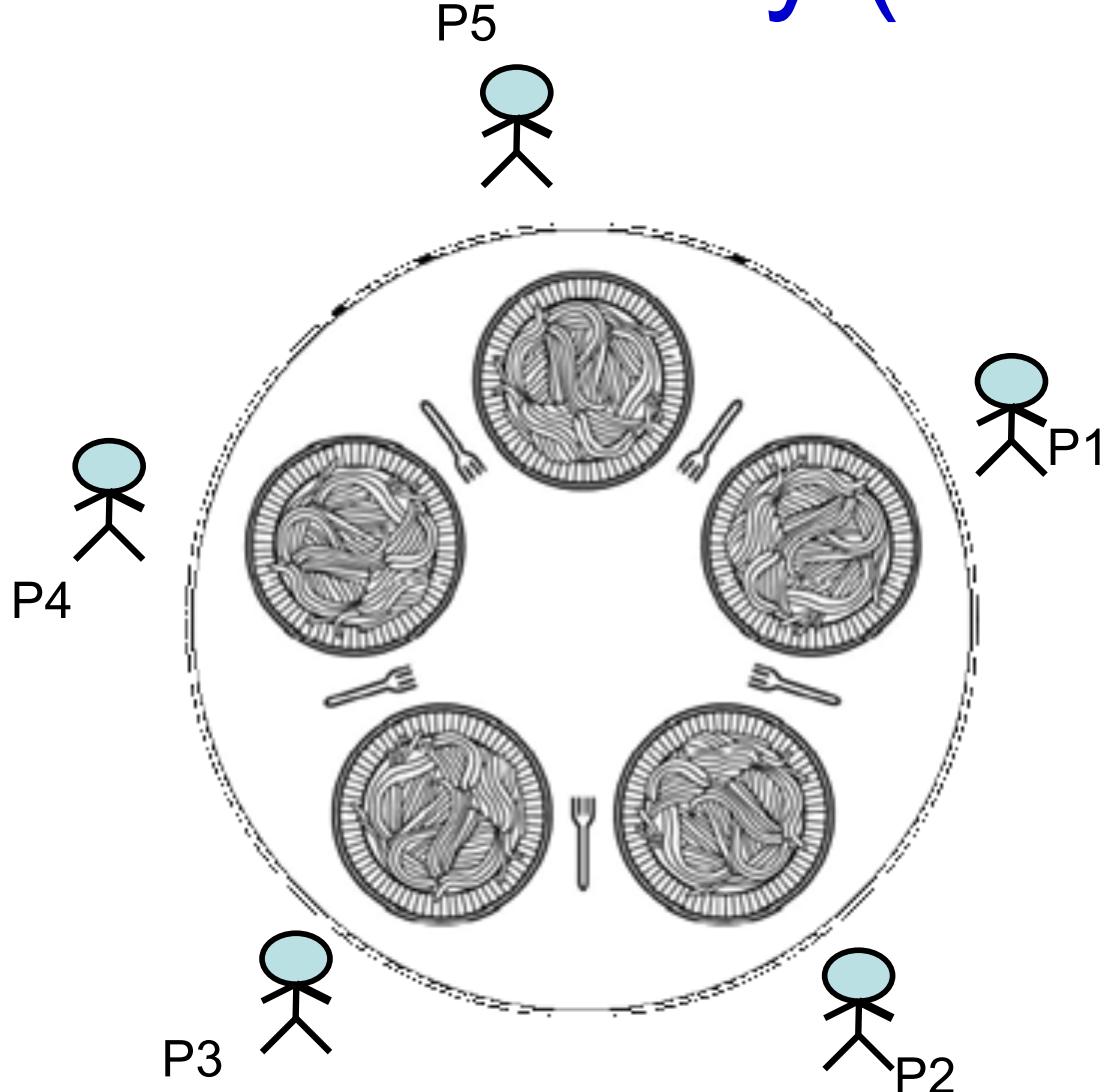
Run simultaneously

And think for the same time

This could lead to philosophers taking fork and putting it down continuously. a deadlock.

```
while(TRUE){  
    think();  
    take_fork(Ri);  
    if (available(Li){  
        take_fork(Li);  
        eat();  
    }else{  
        put_fork(Ri);  
        sleep(T);  
    }  
}
```

# Second try (a better solution)



```
#define N 5

void philosopher(int i){
    while(TRUE){
        think();
        take_fork(Ri);
        if (available(Li)){
            take_fork(Li);
            eat();
        }else{
            put_fork(Ri);
            sleep(random_time);
        }
    }
}
```

# Solution using Mutex

- Protect critical sections with a mutex
- Prevents deadlock
- But has performance issues
  - Only one philosopher can eat at a time

```
#define N 5

void philosopher(int i){
    while(TRUE){
        think(); // for some_time
        wait(mutex);
        take_fork(Ri);
        take_fork(Li);
        eat();
        put_fork(Ri);
        put_fork(Li);
        signal(mutex);
    }
}
```

# Solution with Semaphores

Uses N semaphores ( $s[1], s[2], \dots, s[N]$ ) all initialized to 0, and a mutex

Philosopher has 3 states: HUNGRY, EATING, THINKING

*A philosopher can only move to EATING state if neither neighbor is eating*

```
void philosopher(int i){  
    while(TRUE){  
        think();  
        take_forks(i);  
        eat();  
        put_forks();  
    }  
}
```

```
void take_forks(int i){  
    lock(mutex);  
    state[i] = HUNGRY;  
    test(i);  
    unlock(mutex);  
    down(s[i]);  
}
```

```
void put_forks(int i){  
    lock(mutex);  
    state[i] = THINKING;  
    test(LEFT);  
    test(RIGHT);  
    unlock(mutex);  
}
```

```
void test(int i){  
    if (state[i] = HUNGRY && state[LEFT] != EATING && state[RIGHT] != EATING){  
        state[i] = EATING;  
        up(s[i]);  
    }  
}
```

# Solution to Dining Philosophers

```
void philosopher(int i){  
    while(TRUE){  
        think();  
        take_forks(i);  
        eat();  
        put_forks();  
    }  
}
```

```
void take_forks(int i){  
    lock(mutex);  
    state[i] = HUNGRY;  
    test(i);  
    unlock(mutex);  
    down(s[i]);  
}
```

```
void put_forks(int i){  
    lock(mutex);  
    state[i] = THINKING;  
    test(LEFT);  
    test(RIGHT)  
    unlock(mutex);  
}
```

```
void test(int i){  
    if (state[i] = HUNGRY &&  
        state[LEFT] != EATING &&  
        state[RIGHT] != EATING){  
        state[i] = EATING;  
        up(s[i]);  
    }  
}
```

	P1	P2	P3	P4	P5
state	T	T	T	T	T
semaphore	0	0	0	0	0

# Solution to Dining Philosophers

```
void philosopher(int i){  
    while(TRUE){  
        → think();  
        → take_forks(i);  
        → eat();  
        put_forks();  
    }  
}
```

```
void take_forks(int i){  
    lock(mutex);  
    → state[i] = HUNGRY;  
    → test(i);  
    → unlock(mutex);  
    → down(s[i]);  
}
```

```
void put_forks(int i){  
    lock(mutex);  
    state[i] = THINKING;  
    test(LEFT);  
    test(RIGHT);  
    unlock(mutex);  
}
```

s[i] is 1, so down will not block.  
The value of s[i] decrements by 1.

```
void test(int i){  
    → if (state[i] = HUNGRY &&  
        state[LEFT] != EATING &&  
        state[RIGHT] != EATING){  
        → state[i] = EATING;  
        up(s[i]);  
    }  
}
```

	P1	P2	P3	P4	P5
state	T	T	E	T	T
semaphore	0	0	0	0	0

# Solution to Dining Philosophers

```
void philosopher(int i){  
    while(TRUE){  
        think();  
        → take_forks(i);  
        → eat();  
        put_forks();  
    }  
}
```

```
void take_forks(int i){  
    → lock(mutex);  
    → state[i] = HUNGRY;  
    → test(i);  
    unlock(mutex);  
    down(s[i]);  
}
```

```
void put_forks(int i){  
    lock(mutex);  
    state[i] = THINKING;  
    test(LEFT);  
    test(RIGHT);  
    unlock(mutex);  
}
```

```
void test(int i){  
    → if (state[i] = HUNGRY &&  
        state[LEFT] != EATING &&  
        state[RIGHT] != EATING){  
        state[i] = EATING;  
        up(s[i]);  
    }  
}
```

	P1	P2	P3	P4	P5
state	T	T	E	H	T
semaphore	0	0	0	0	0

# Solution to Dining Philosophers

```
void philosopher(int i){  
    while(TRUE){  
        think();  
        take_forks(i);  
        → eat();  
        → put_forks();  
    }  
}
```

```
void take_forks(int i){  
    lock(mutex);  
    state[i] = HUNGRY;  
    test(i);  
    unlock(mutex);  
    down(s[i]);  
}
```

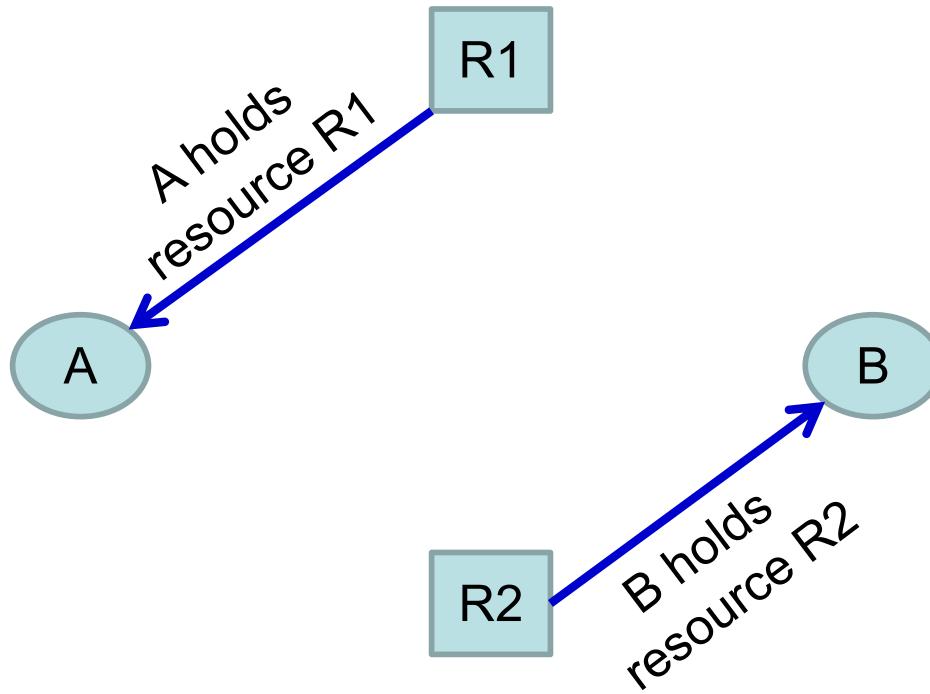
```
void put_forks(int i){  
    lock(mutex);  
    → state[i] = THINKING;  
    → test(LEFT);  
    → test(RIGHT)  
    unlock(mutex);  
}
```

```
void test(int i){  
    → if (state[i] = HUNGRY &&  
         state[LEFT] != EATING &&  
         state[RIGHT] != EATING){  
        → state[i] = EATING;  
        up(s[i]);  
    }  
}
```

	P1	P2	P3	P4	P5
state	T	T	E	■	T
semaphore	0	0	0	0	0

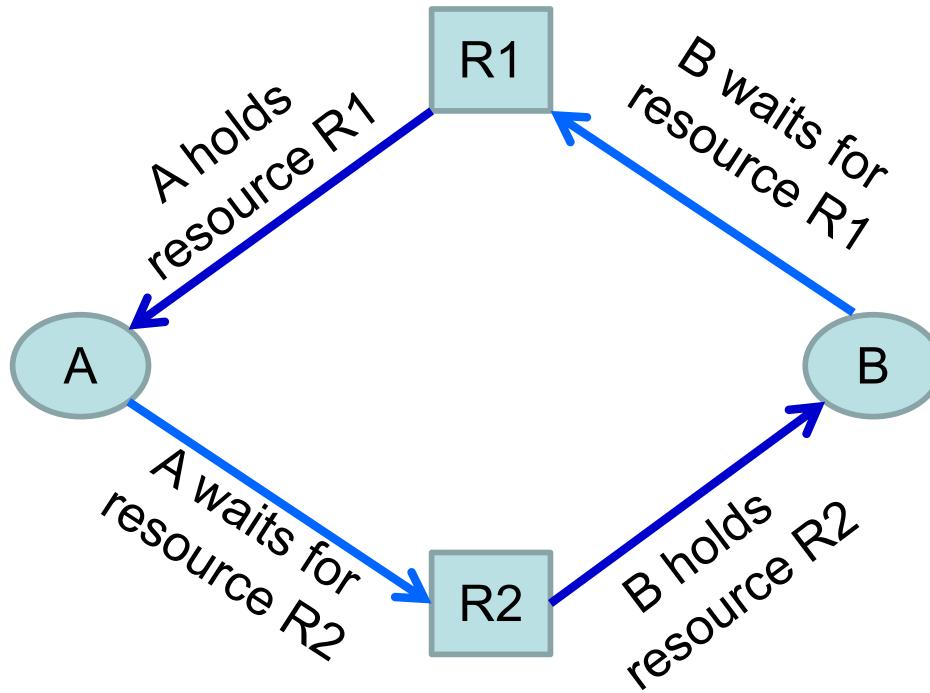
# Deadlocks

# Deadlocks



Consider this situation:

# Deadlocks

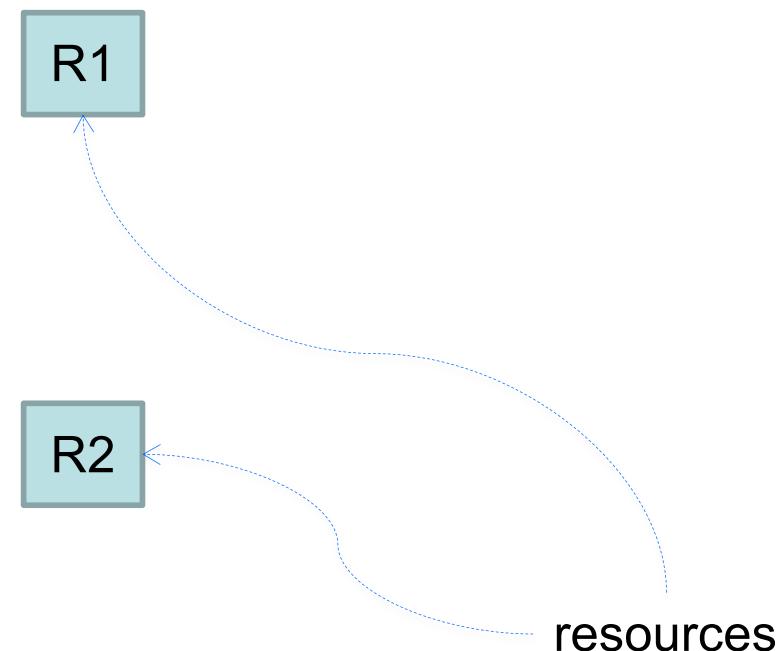


## A Deadlock Arises:

Deadlock : A set of processes is deadlocked if each process in the set is waiting for an event that only another process in the set can cause.

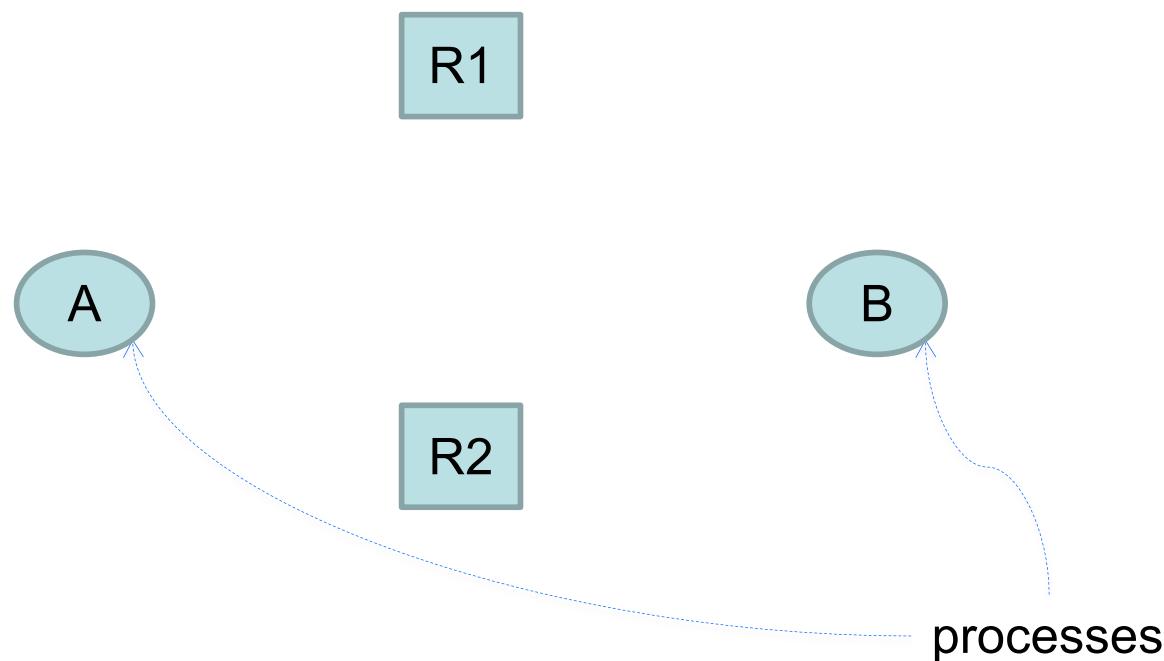
# Resource Allocation Graph

- It is a directed graph
- Used to model resource allocations



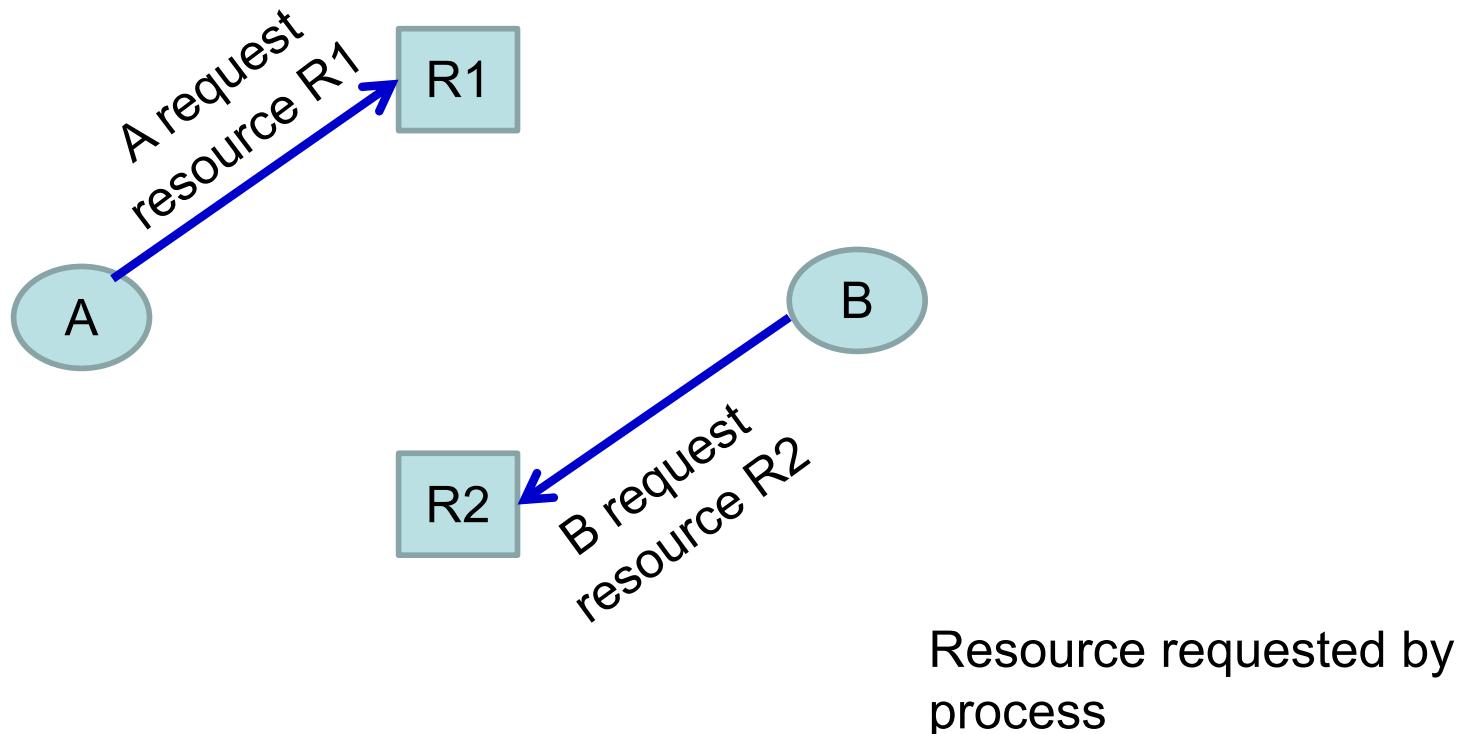
# Resource Allocation Graph

- It is a directed graph
- Used to model resource allocations



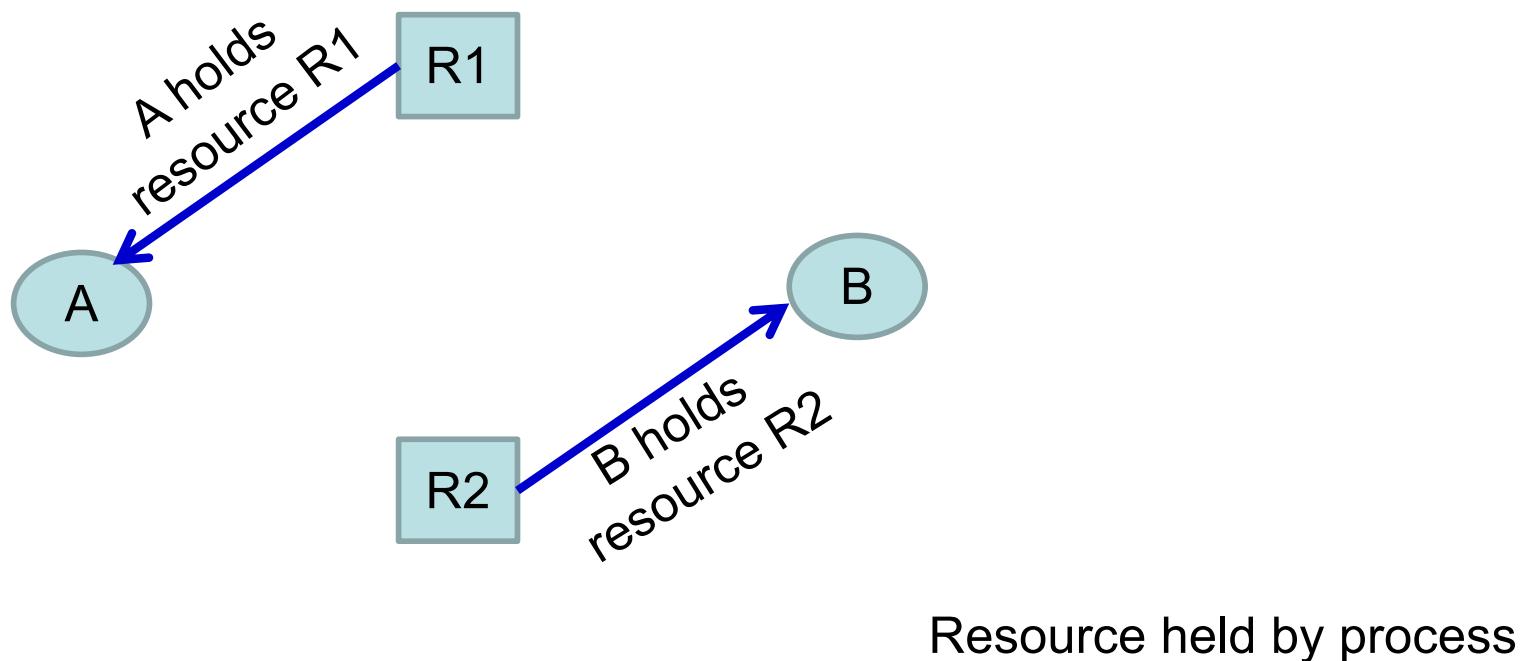
# Resource Allocation Graph

- It is a directed graph
- Used to model resource allocations



# Resource Allocation Graph

- It is a directed graph
- Used to model resource allocations



# Conditions for Resource Deadlocks

## 1. Mutual Exclusion

- Each resource is either available or currently assigned to exactly one process

## 2. Hold and wait

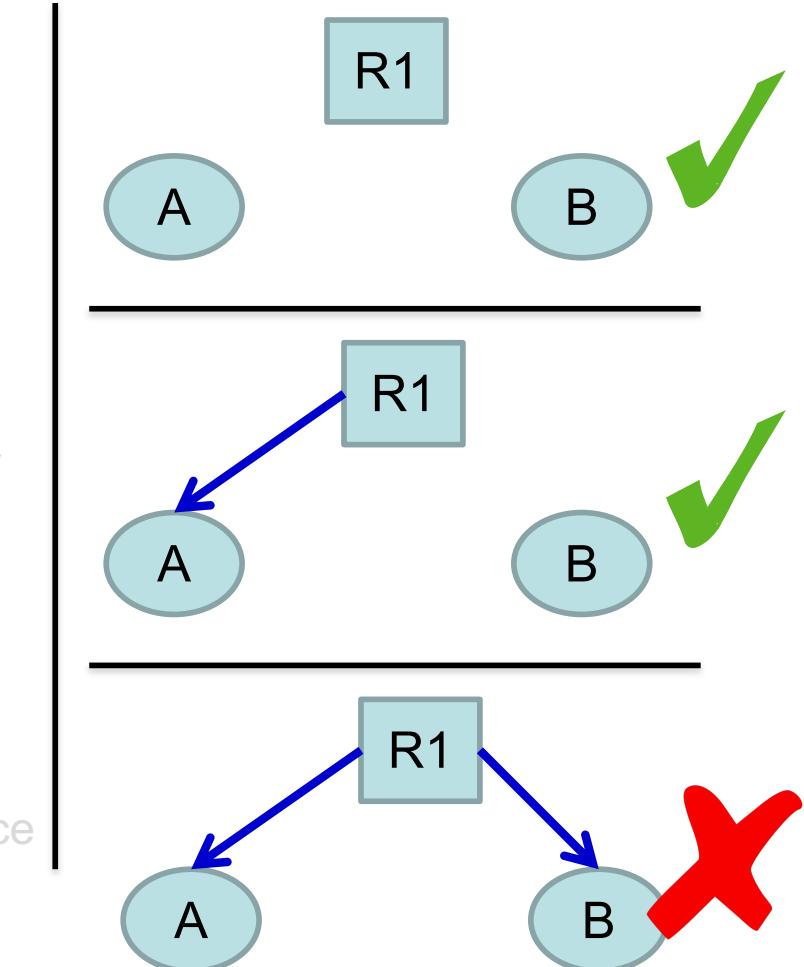
- A process holding a resource requests another resource

## 3. No preemption

- Resources previously granted cannot be forcibly taken away from a process
- They must be explicitly released by the process holding them

## 4. Circular wait

- There must be a circular chain of two or more processes, each of which is waiting for a resource held by the next member of the chain



# Conditions for Resource Deadlocks

## 1. Mutual Exclusion

- Each resource is either available or currently assigned to exactly one process

## 2. Hold and wait

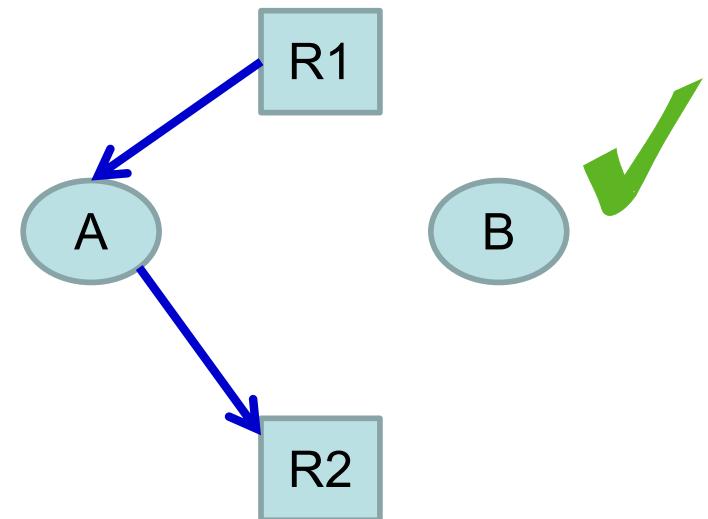
- A process holding a resource requests another resource

## 3. No preemption

- Resources previously granted cannot be forcibly taken away from a process
- They must be explicitly released by the process holding them

## 4. Circular wait

- There must be a circular chain of two or more processes, each of which is waiting for a resource held by the next member of the chain



# Conditions for Resource Deadlocks

## 1. Mutual Exclusion

- Each resource is either available or currently assigned to exactly one process

## 2. Hold and wait

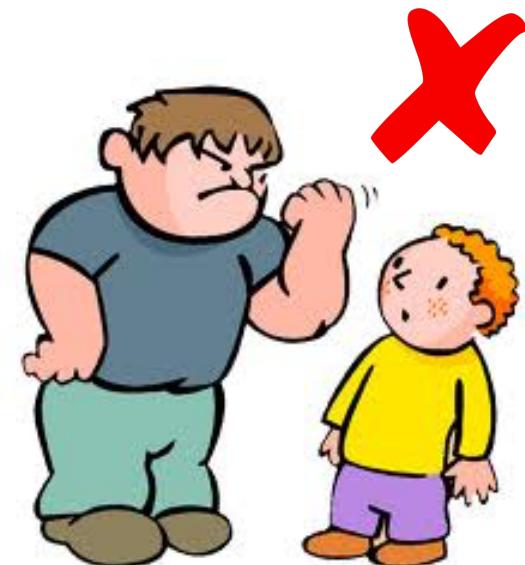
- A process holding a resource requests another resource

## 3. No preemption

- Resources previously granted cannot be forcibly taken away from a process
- They must be explicitly released by the process holding them

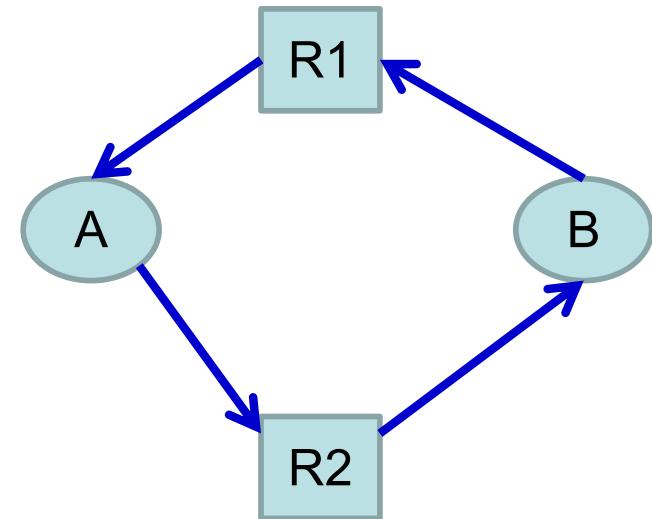
## 4. Circular wait

- There must be a circular chain of two or more processes, each of which is waiting for a resource held by the next member of the chain



# Conditions for Resource Deadlocks

1. Mutual Exclusion
  - Each resource is either available or currently assigned to exactly one process
2. Hold and wait
  - A process holding a resource requests another resource
3. No preemption
  - Resources previously granted cannot be forcibly taken away from a process
  - They must be explicitly released by the process holding them
4. Circular wait
  - There must be a circular chain of two or more processes, each of which is waiting for a resource held by the next member of the chain



# Conditions for Resource Deadlocks

## 1. Mutual Exclusion

- Each resource is either available or currently assigned to exactly one process

## 2. Hold and wait

- A process holding a resource requests another resource

## 3. No preemption

- Resources previously granted cannot be forcibly taken away from a process
- They must be explicitly released by the process holding them

## 4. Circular wait

- There must be a circular chain of two or more processes, each of which is waiting for a resource held by the next member of the chain

All four of these conditions must be present for a resource deadlock to occur!!

# Deadlocks : (A Chanced Event)

- Ordering of resource requests and allocations are probabilistic, thus deadlock occurrence is also probabilistic

A  
Request R  
Request S  
Release R  
Release S  
(a)

B  
Request S  
Request T  
Release S  
Release T  
(b)

C  
Request T  
Request R  
Release T  
Release R  
(c)



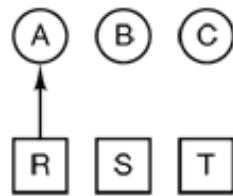
Deadlock occurs

A	B	C
Request R	Request S	Request T
Request S	Request T	Request R
Release R	Release S	Release T
Release S	Release T	Release R
(a)	(b)	(c)

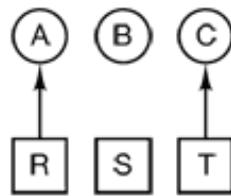
No deadlock occurrence  
(B can be granted S  
after step q)

1. A requests R
2. C requests T
3. A requests S
4. C requests R
5. A releases R
6. A releases S  
no deadlock

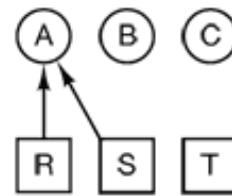
(k)



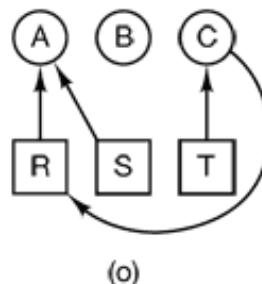
(l)



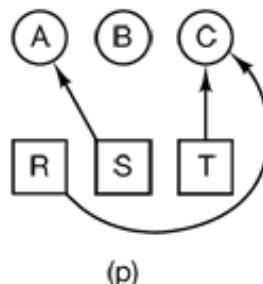
(m)



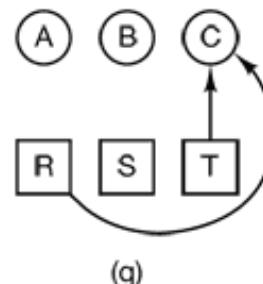
(n)



(o)



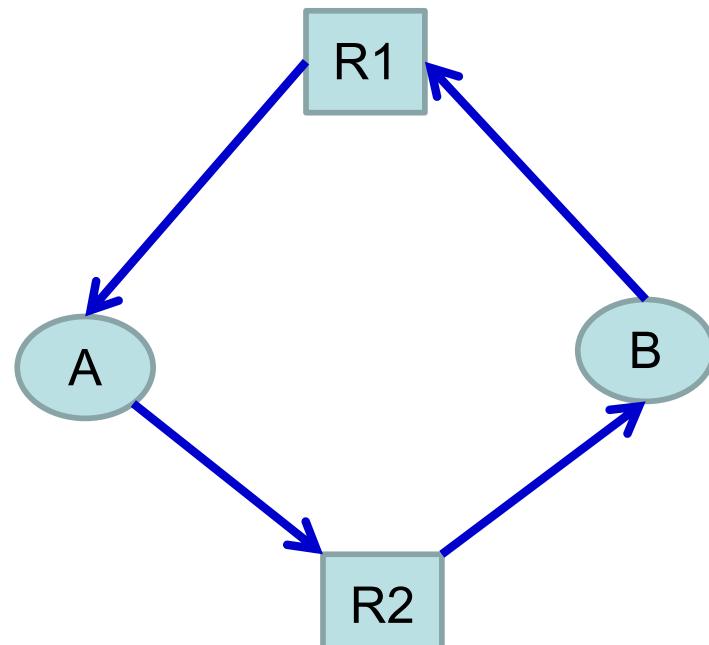
(p)



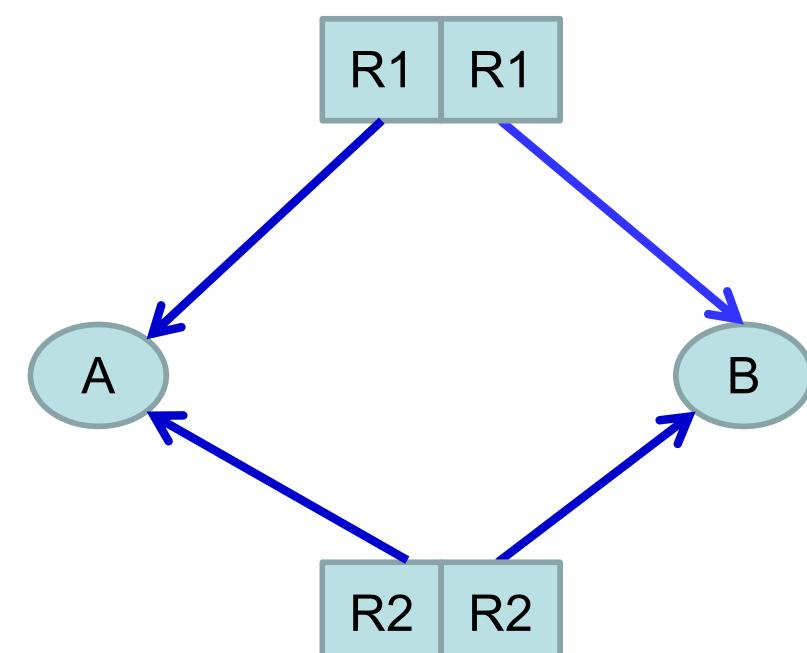
(q)

# Multiple Resources

- Having multiple resources can potentially reduce the chance of having a deadlock



A deadlocked state



No Deadlocks

# Should Deadlocks be handled?

- Preventing / detecting deadlocks could be tedious
- Can we live without detecting / preventing deadlocks?
  - What is the probability of occurrence?
  - What are the consequences of a deadlock? (How critical is a deadlock?)



# Handling Deadlocks

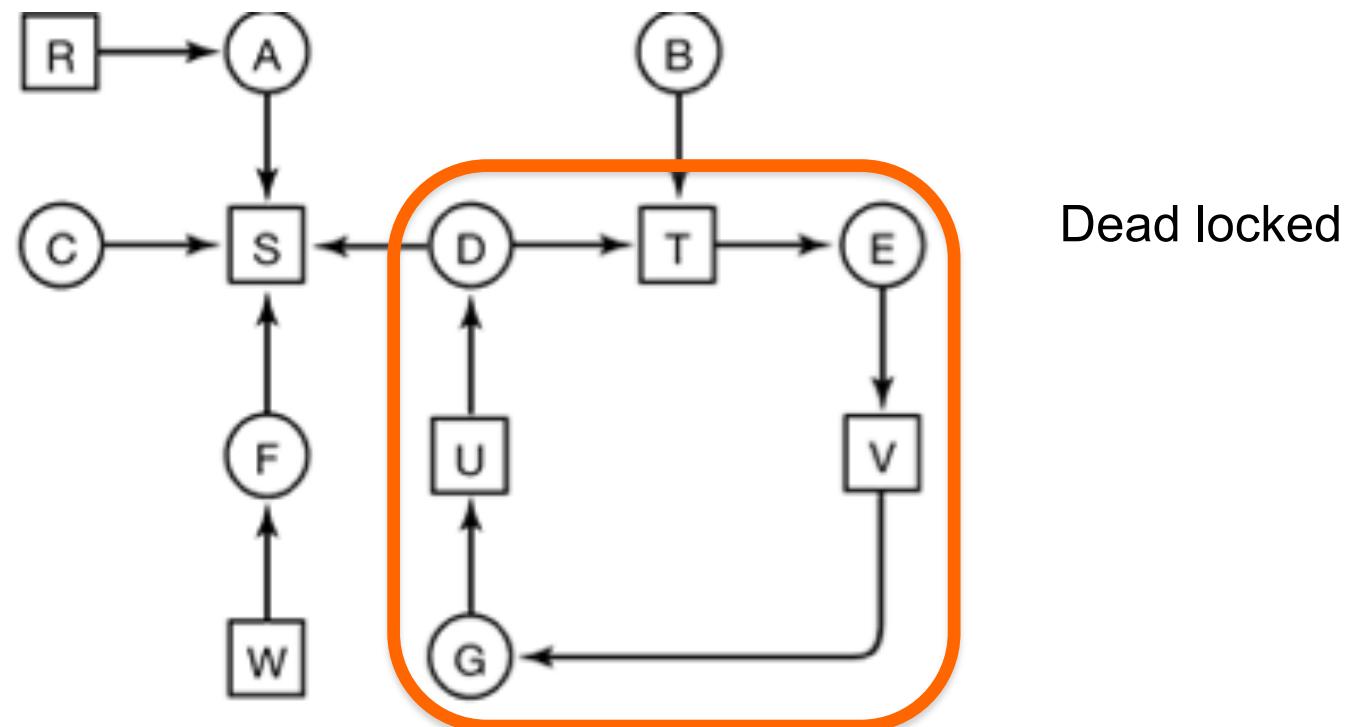
- Detection and Recovery
- Avoidance
- Prevention

# Deadlock Detection

- How can an OS detect when there is a deadlock?
- OS needs to keep track of
  - Current resource allocation
    - Which process has which resource
  - Current request allocation
    - Which process is waiting for which resource
- Use this information to detect deadlocks

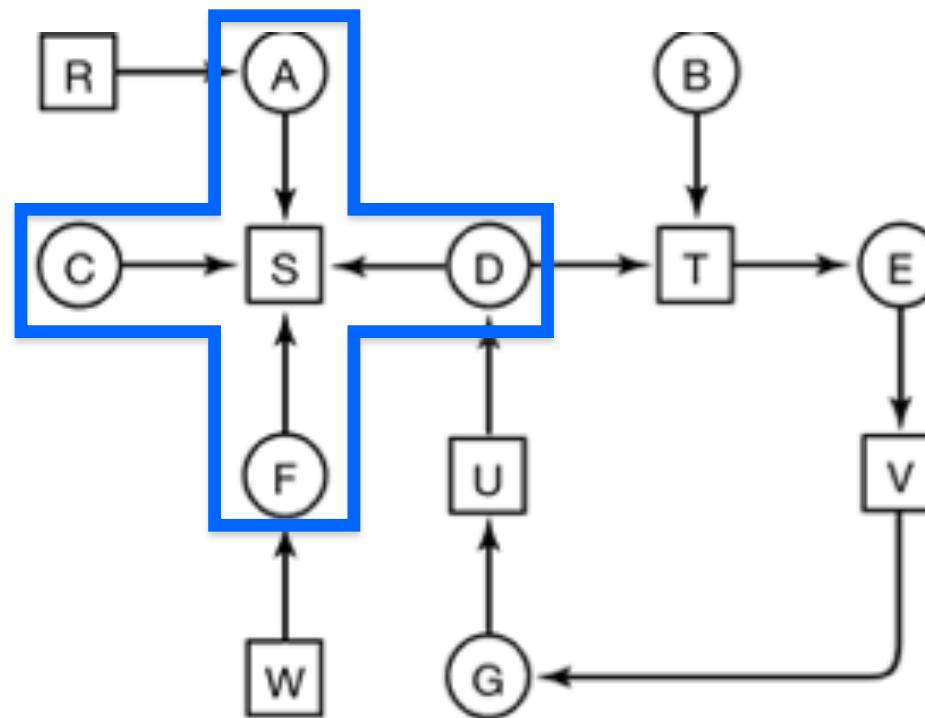
# Deadlock Detection

- Deadlock detection with **one resource of each type**
- Find cycles in resource graph



# Deadlock Detection

- Deadlock detection with **one resource of each type**
- There should be atleast one sequence of resource allocation, to avoid a deadlock



Not a Dead lock as  
an allocation  
sequence possible

**Sample Allocation Sequence**

S allocated to A, after A  
completes S allocated to C  
then, S allocated to F, and  
finally S allocated to D

# Deadlock Detection

- Deadlock detection with multiple resources of each type

$$E = \begin{pmatrix} 4 & 2 & 3 & 1 \end{pmatrix}$$

**Existing Resource Vector**

Tape drives	Plotters	Scanners	CD Roms
-------------	----------	----------	---------

$$A = \begin{pmatrix} 2 & 1 & 0 & 0 \end{pmatrix}$$

**Resources Available**

Tape drives	Plotters	Scanners	CD Roms
-------------	----------	----------	---------

$$\sum_{i=1}^n C_{ij} + A_j = E_j$$

$P_1$       Current allocation matrix

$P_2$        $C = \begin{bmatrix} 0 & 0 & 1 & 0 \\ 2 & 0 & 0 & 1 \\ 0 & 1 & 2 & 0 \end{bmatrix}$

**Current Allocation Matrix**

*Who has what!!*

Request matrix

$R = \begin{bmatrix} 2 & 0 & 0 & 1 \\ 1 & 0 & 1 & 0 \\ 2 & 1 & 0 & 0 \end{bmatrix}$

**Request Matrix**

*Who is waiting for what!!*

Process  $P_i$  holds  $C_i$  resources and requests  $R_i$  resources, where  $i = 1$  to  $3$

Goal is to check if there is any sequence of allocations by which all current requests can be met. If so, there is no deadlock.

# Deadlock Detection

- Deadlock detection with multiple resources of each type

$$E = \begin{pmatrix} 4 & 2 & 3 & 1 \end{pmatrix}$$

Tape drives  
Plotters  
Scanners  
CD Roms

**Existing Resource Vector**

$$A = \begin{pmatrix} 2 & 1 & 0 & 0 \end{pmatrix}$$

Tape drives  
Plotters  
Scanners  
CD Roms

**Resources Available**

$$\sum_{i=1}^n C_{ij} + A_j = E_j$$

$P_1$  Current allocation matrix

$$C = \begin{bmatrix} 0 & 0 & 1 & 0 \\ 2 & 0 & 0 & 1 \\ 0 & 1 & 2 & 0 \end{bmatrix}$$

**Current Allocation Matrix**

*Who has what!!*

Request matrix

$$R = \begin{bmatrix} 2 & 0 & 0 & 1 \\ 1 & 0 & 1 & 0 \\ 2 & 1 & 0 & 0 \end{bmatrix}$$

**Request Matrix**

*Who is waiting for what!!*

Process  $P_i$  holds  $C_i$  resources and requests  $R_i$  resources, where  $i = 1$  to  $3$

Goal is to check if there is any sequence of allocations by which all current requests can be met. If so, there is no deadlock.

# Deadlock Detection

- Deadlock detection with multiple resources of each type

$$\begin{matrix} & \text{Tape drives} \\ E = & \begin{matrix} 4 & 2 & 3 & 1 \end{matrix} \\ \text{Existing Resource Vector} & \begin{matrix} \text{Plotters} \\ \text{Scanners} \\ \text{CD Roms} \end{matrix} \end{matrix}$$

$$\begin{matrix} & \text{Tape drives} \\ A = & \begin{matrix} 2 & 1 & 0 & 0 \end{matrix} \\ \text{Resources Available} & \begin{matrix} \text{Plotters} \\ \text{Scanners} \\ \text{CD Roms} \end{matrix} \end{matrix}$$

$$\sum_{i=1}^n C_{ij} + A_j = E_j$$

$$\begin{matrix} P_1 & \text{Current allocation matrix} \\ P_2 & \\ P_3 & \end{matrix}$$
$$C = \begin{bmatrix} 0 & 0 & 1 & 0 \\ 2 & 0 & 0 & 1 \\ 0 & 1 & 2 & 0 \end{bmatrix}$$

**Current Allocation Matrix**

$$\begin{matrix} & \text{Request matrix} \\ R = & \begin{bmatrix} 2 & 0 & 0 & 1 \\ 1 & 0 & 1 & 0 \\ 2 & 1 & 0 & 0 \end{bmatrix} \end{matrix}$$

**Request Matrix**

$P_1$  cannot be satisfied  
 $P_2$  cannot be satisfied  
 $P_3$  can be satisfied

Process  $P_i$  holds  $C_i$  resources and requests  $R_i$  resources, where  $i = 1$  to 3

# Deadlock Detection

- Deadlock detection with multiple resources of each type

	Tape drives	Plotters	Scanners	CD Roms
$E = (4 \ 2 \ 3 \ 1)$	4	2	3	1
<b>Existing Resource Vector</b>				

	Tape drives	Plotters	Scanners	CD Roms
$A = (2 \ 1 \ 0 \ 0)$	2	1	0	0
<b>Resources Available</b>				

$P_1$  Current allocation matrix

$P_2$

$$P_3 \quad C = \begin{bmatrix} 0 & 0 & 1 & 0 \\ 2 & 0 & 0 & 1 \\ 0 & 1 & 2 & 0 \end{bmatrix}$$

**Current Allocation Matrix**

Request matrix

$$R = \begin{bmatrix} 2 & 0 & 0 & 1 \\ 1 & 0 & 1 & 0 \\ 2 & 1 & 0 & 0 \end{bmatrix}$$

**Request Matrix**

$P_3$  runs and its allocation is  $(2, 2, 2, 0)$

On completion it returns the available resources :  $A = (4 \ 2 \ 2 \ 1)$

Either  $P_1$  or  $P_2$  can now run.

**NO Deadlock!!!**

# Deadlock Detection

- Deadlock detection with multiple resources of each type

$$E = (4 \quad 2 \quad 3 \quad 1)$$

**Existing Resource Vector**

	Tape drives	Plotters	Scanners	CD Roms
E	4	2	3	1

$$A = (2 \quad 1 \quad 0 \quad 0)$$

**Resources Available**

	Tape drives	Plotters	Scanners	CD Roms
A	2	1	0	0

$$\sum_{i=1}^n C_{ij} + A_j = E_j$$

$$P_1 \quad \text{Current allocation matrix}$$
$$P_2 \quad C = \begin{bmatrix} 0 & 0 & 1 & 0 \\ 2 & 0 & 0 & 1 \\ 0 & 1 & 2 & 0 \end{bmatrix}$$

**Current Allocation Matrix**

$$P_3$$

Request matrix

$$R = \begin{bmatrix} 2 & 0 & 0 & 1 \\ 1 & 0 & 1 & 0 \\ 2 & 1 & 1 & 0 \end{bmatrix}$$

**Request Matrix**

P<sub>1</sub> cannot be satisfied

P<sub>2</sub> cannot be satisfied

P<sub>3</sub> cannot be satisfied  
**deadlock**

Process P<sub>i</sub> holds C<sub>i</sub> resources and requests R<sub>i</sub> resources, where i = 1 to 3

**Deadlock detected as none of the requests can be satisfied**

# Deadlock Detection

- Deadlock detection with multiple resources of each type

Tape drives  
Plotters  
Scanners  
CD Roms

$$E = (4 \quad 2 \quad 3 \quad 1)$$

**Existing Resource Vector**

Tape drives  
Plotters  
Scanners  
CD Roms

$$A = (2 \quad 1 \quad 0 \quad 0)$$

**Resources Available**

$$\sum_{i=1}^n C_{ij} + A_j = E_j$$

$P_1$  Current allocation matrix  
 $P_2$   
 $P_3$

$$C = \begin{bmatrix} 0 & 0 & 1 & 0 \\ 2 & 0 & 0 & 1 \\ 0 & 1 & 2 & 0 \end{bmatrix}$$

**Current Allocation Matrix**

Request matrix  
 $R = \begin{bmatrix} 2 & 0 & 0 & 1 \\ 1 & 0 & 1 & 0 \\ 2 & 1 & 0 & 0 \end{bmatrix}$

**Request Matrix**

$P_1$  cannot be satisfied  
 $P_2$  cannot be satisfied  
 $P_3$  can be satisfied  
**No deadlock**

Process  $P_i$  holds  $C_i$  resources and requests  $R_i$  resources, where  $i = 1$  to 3  
Deadlock not present since requests can be satisfied

# Deadlock Recovery

What should the OS do when it detects a deadlock?

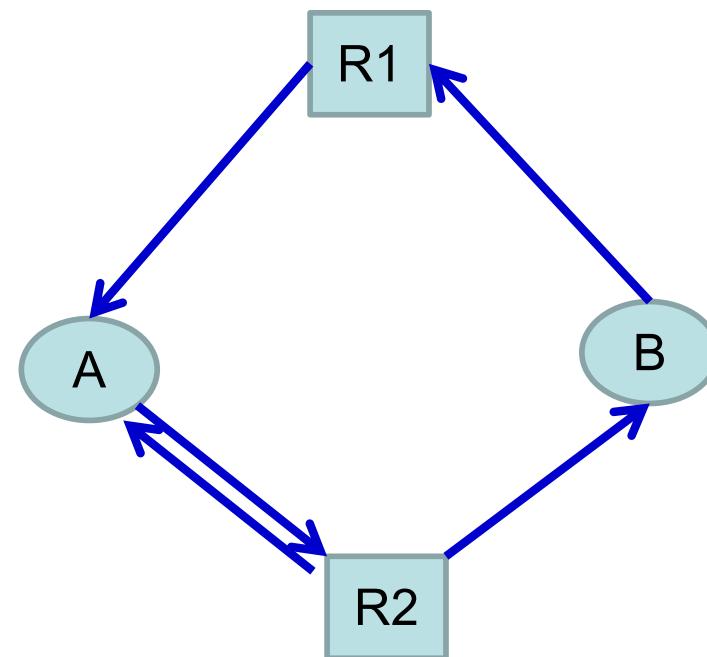
- **Raise an alarm**
  - Tell users and administrator
- Preemption
  - Take away a resource temporarily (frequently not possible)
- Rollback
  - Checkpoint states and then rollback
- Kill process
  - Keep killing processes until deadlock is broken



# Deadlock Recovery

What should the OS do when it detects a deadlock?

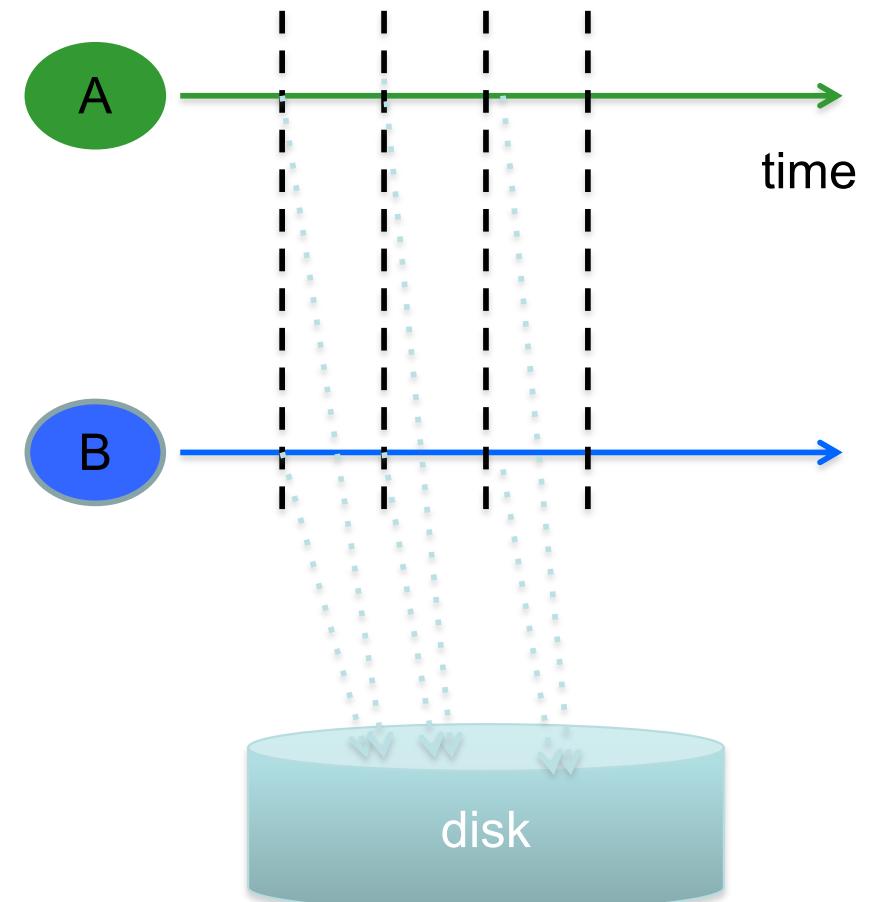
- Raise an alarm
  - Tell users and administrator
- Preemption
  - Take away a resource temporarily (frequently not possible)
- Rollback
  - Checkpoint states and then rollback
- Kill process
  - Keep killing processes until deadlock is broken



# Deadlock Recovery

What should the OS do when it detects a deadlock?

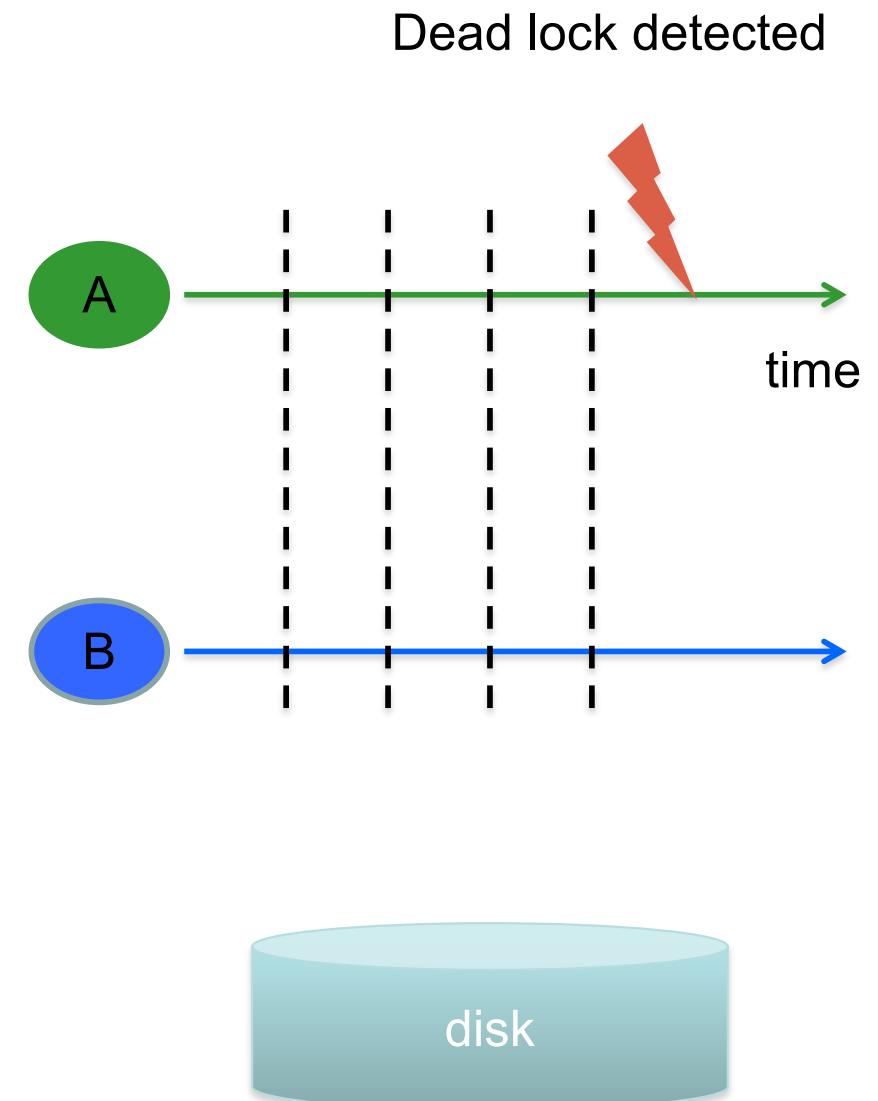
- Raise an alarm
  - Tell users and administrator
- Preemption
  - Take away a resource temporarily (frequently not possible)
- Rollback
  - Checkpoint states and then rollback
- Kill process
  - Keep killing processes until deadlock is broken



# Deadlock Recovery

What should the OS do when it detects a deadlock?

- Raise an alarm
  - Tell users and administrator
- Preemption
  - Take away a resource temporarily (frequently not possible)
- Rollback
  - Checkpoint states and then rollback
- Kill process
  - Keep killing processes until deadlock is broken

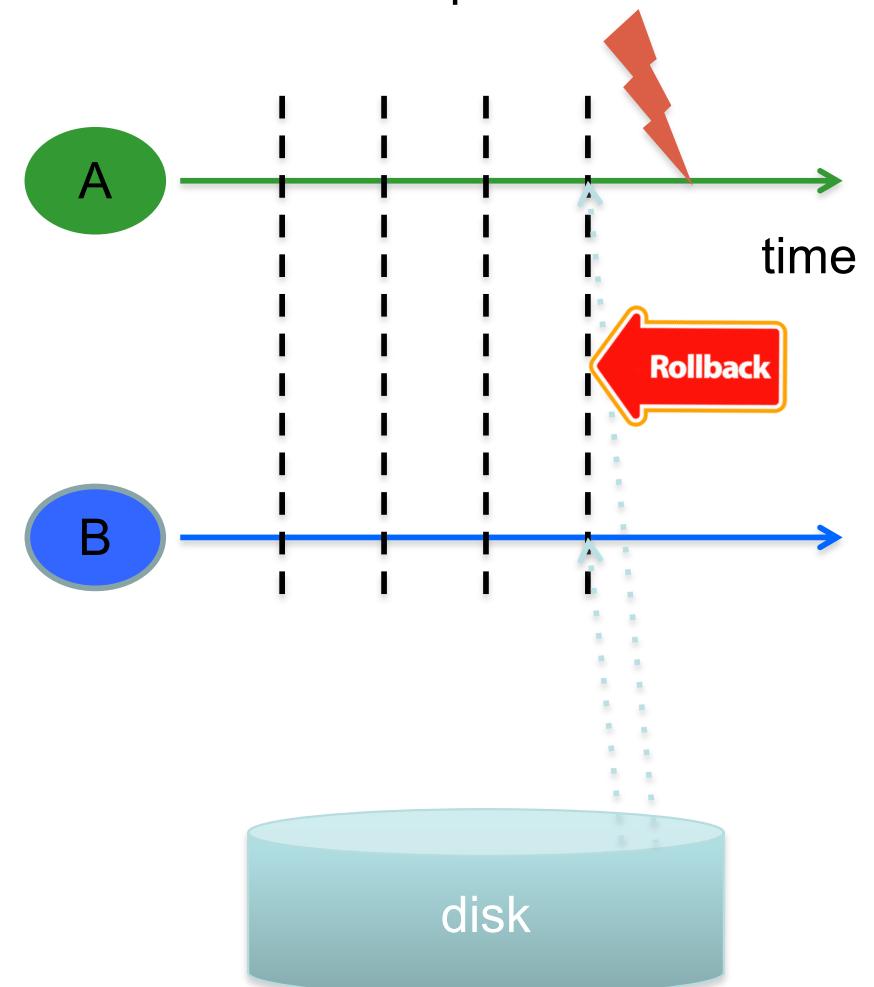


# Deadlock Recovery

What should the OS do when it detects a deadlock?

- Raise an alarm
  - Tell users and administrator
- Preemption
  - Take away a resource temporarily (frequently not possible)
- Rollback
  - Checkpoint states and then rollback
- Kill process
  - Keep killing processes until deadlock is broken

Dead lock detected, then roll back to the last checkpoint



# Deadlock Avoidance

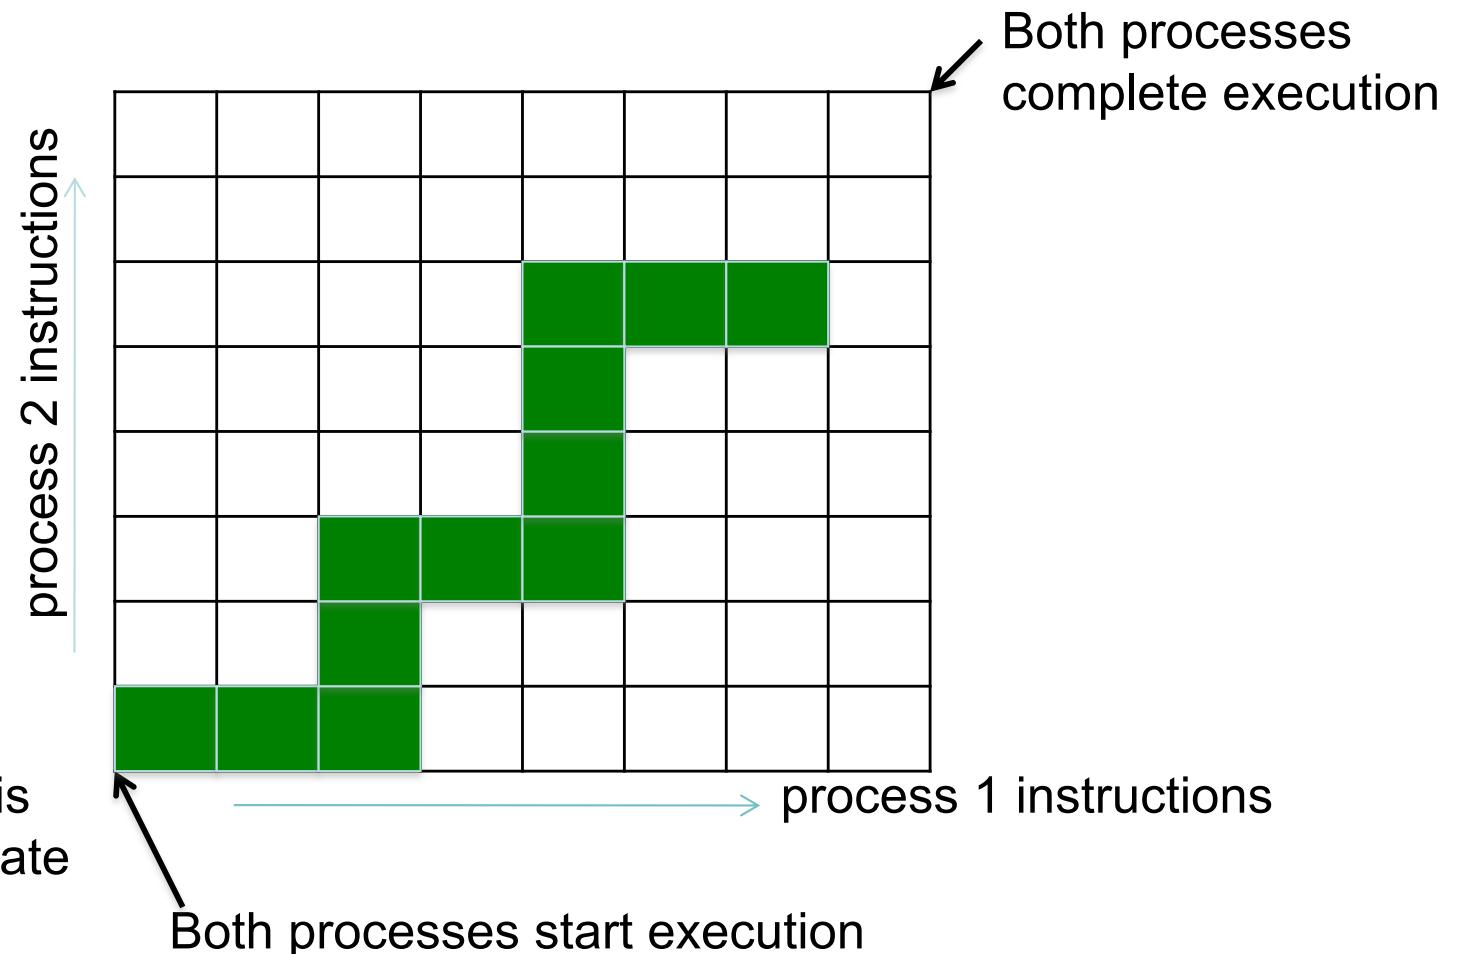
- Basic idea, always make the right choice!



# Deadlock Avoidance

- System decides in advance if allocating a resource to a process will lead to a deadlock

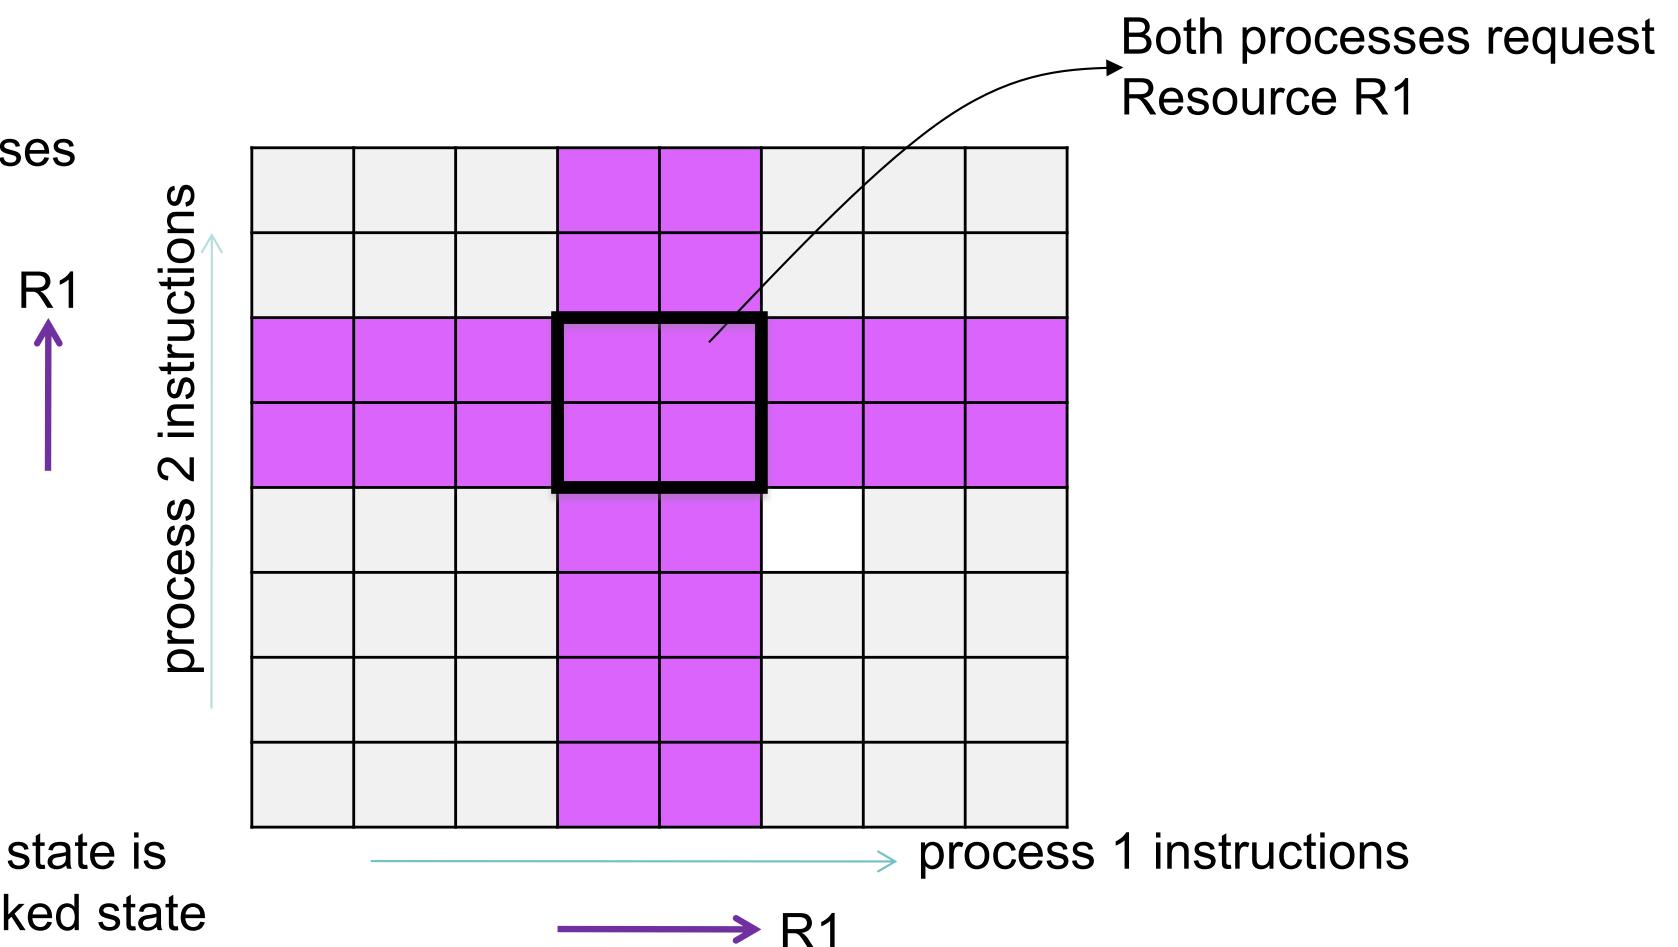
State matrix  
For two processes



# Deadlock Avoidance

- State matrix for two processes

State matrix  
For two processes

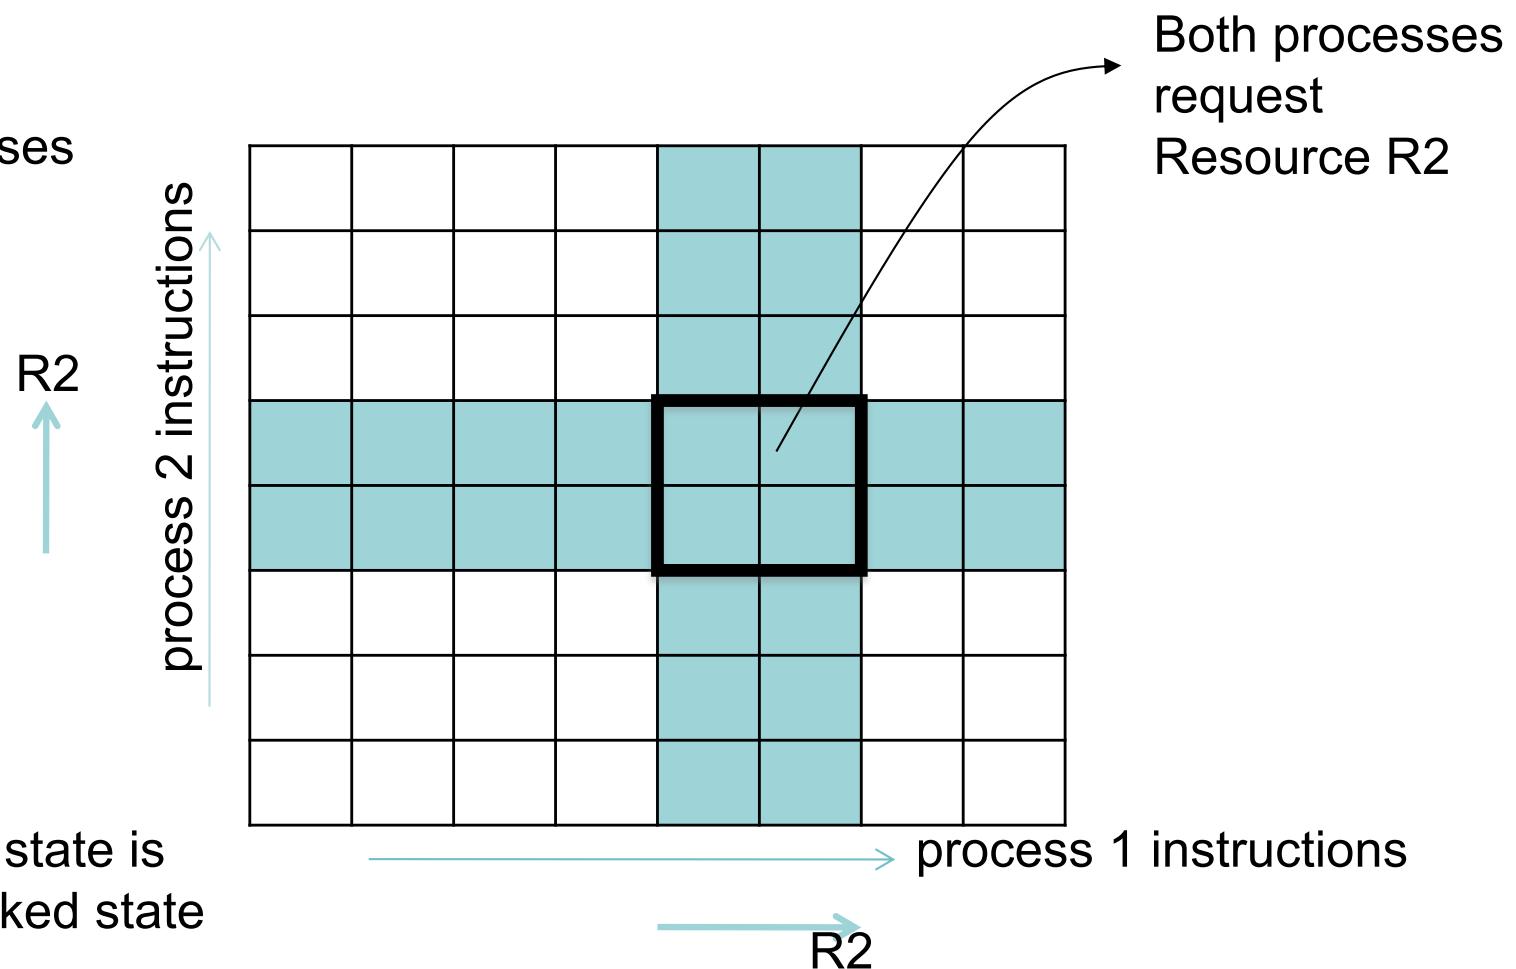


Note: unsafe state is  
not a deadlocked state

# Deadlock Avoidance

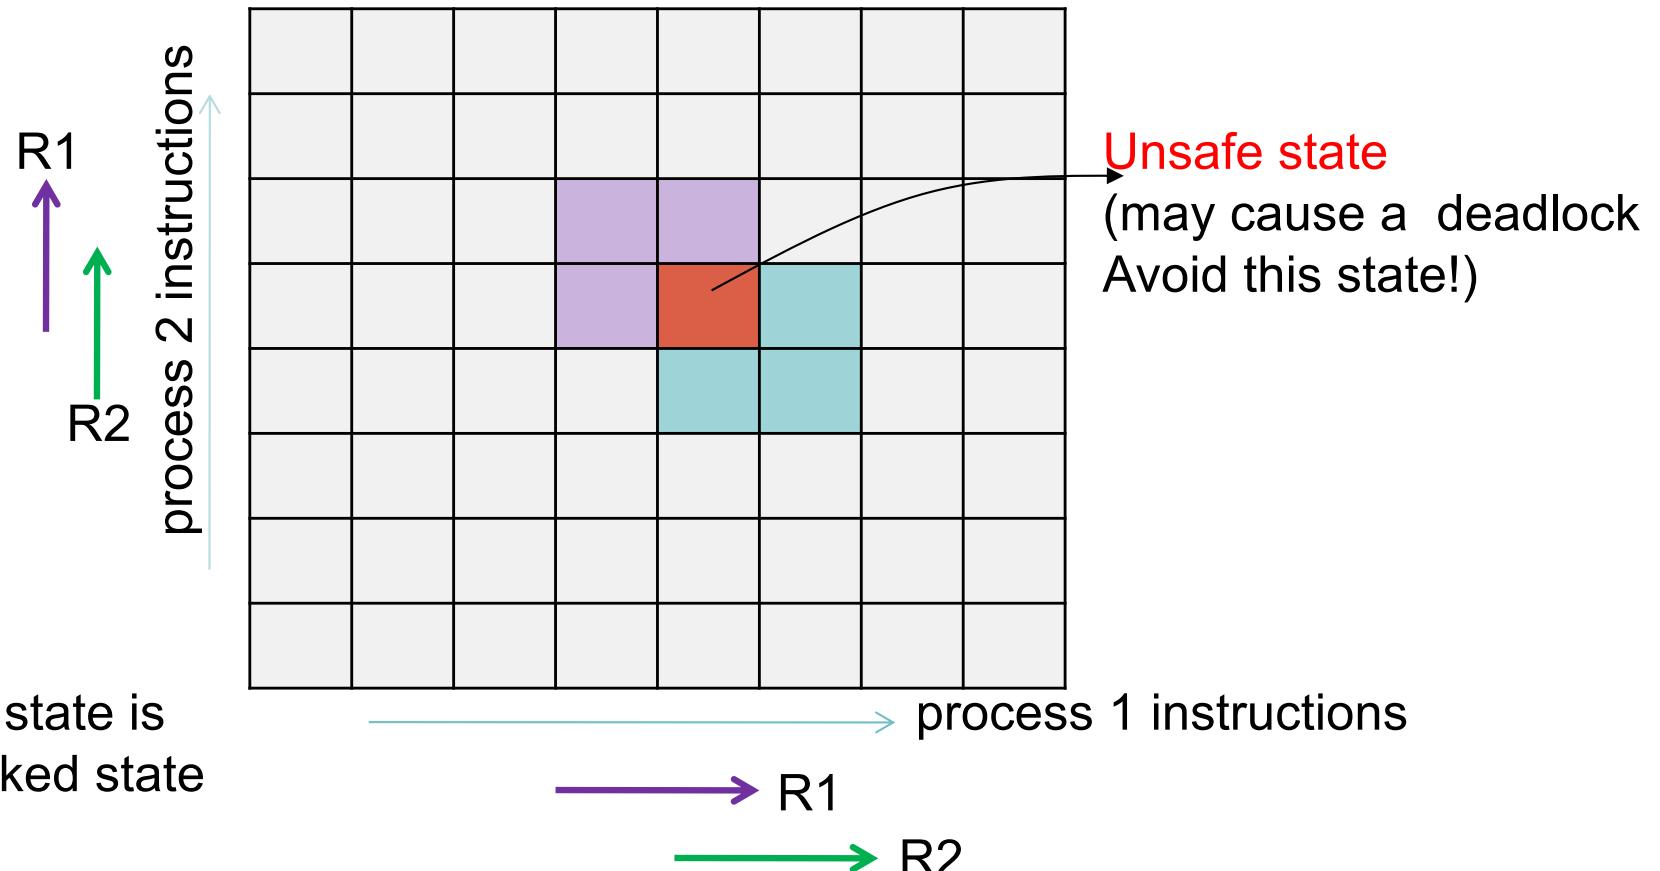
- State matrix for two processes

State matrix  
For two processes



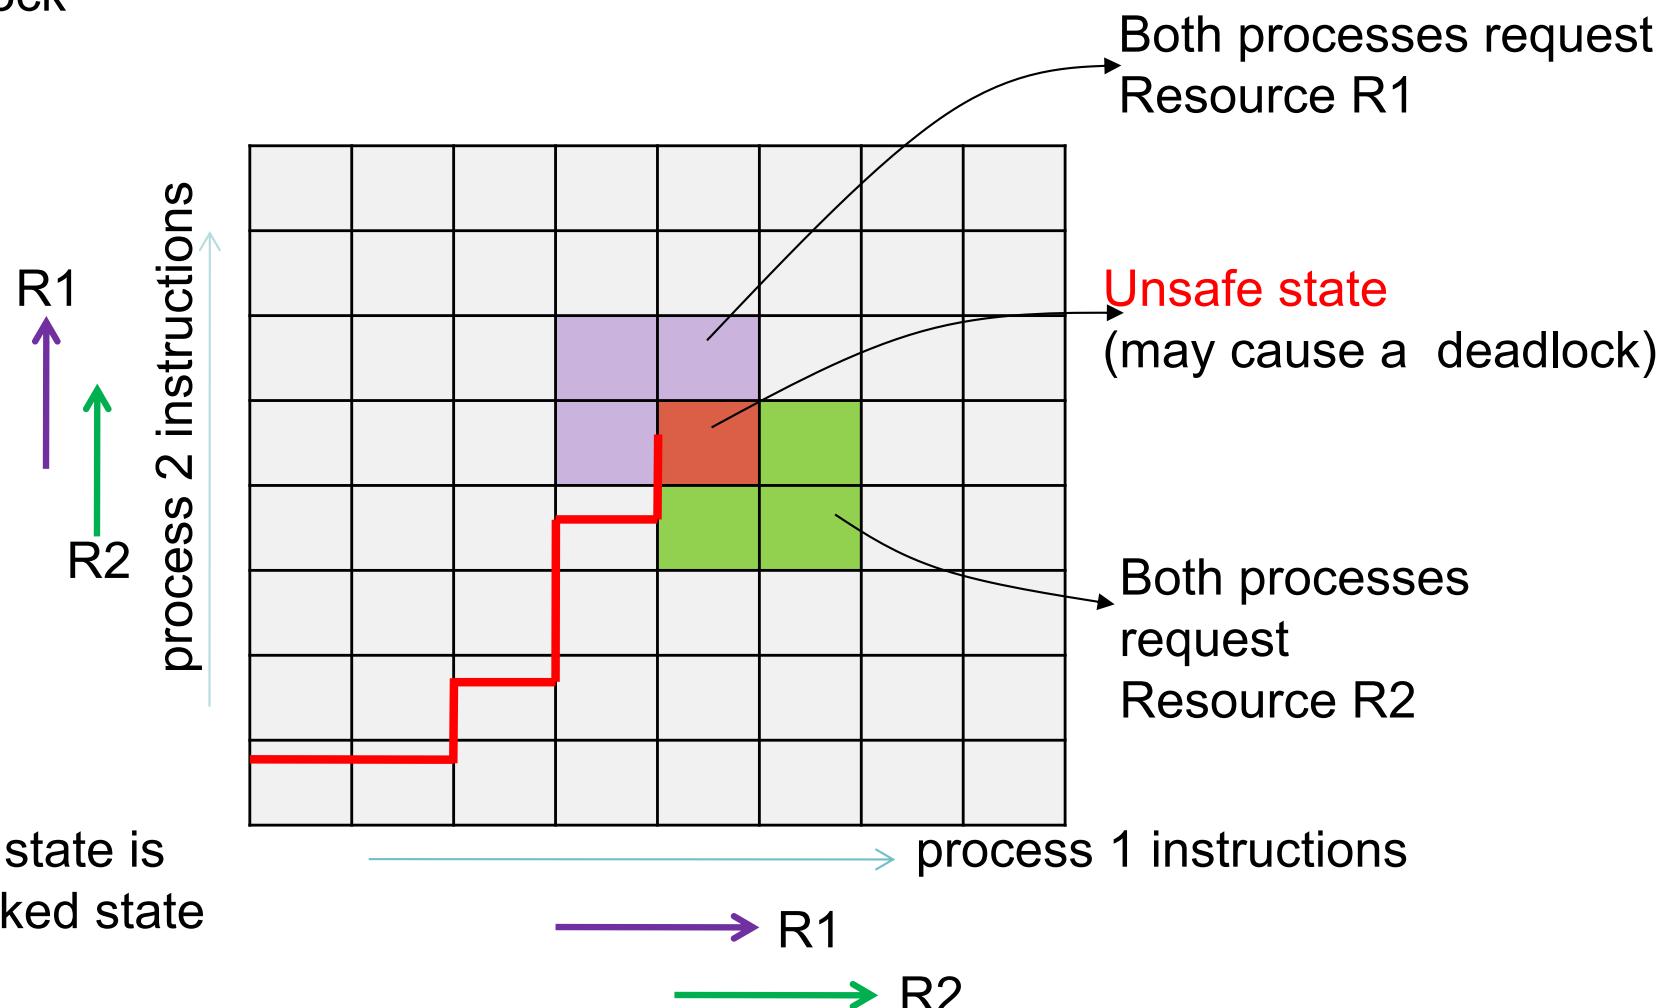
# Deadlock Avoidance

- System decides in advance if allocating a resource to a process will lead to a deadlock



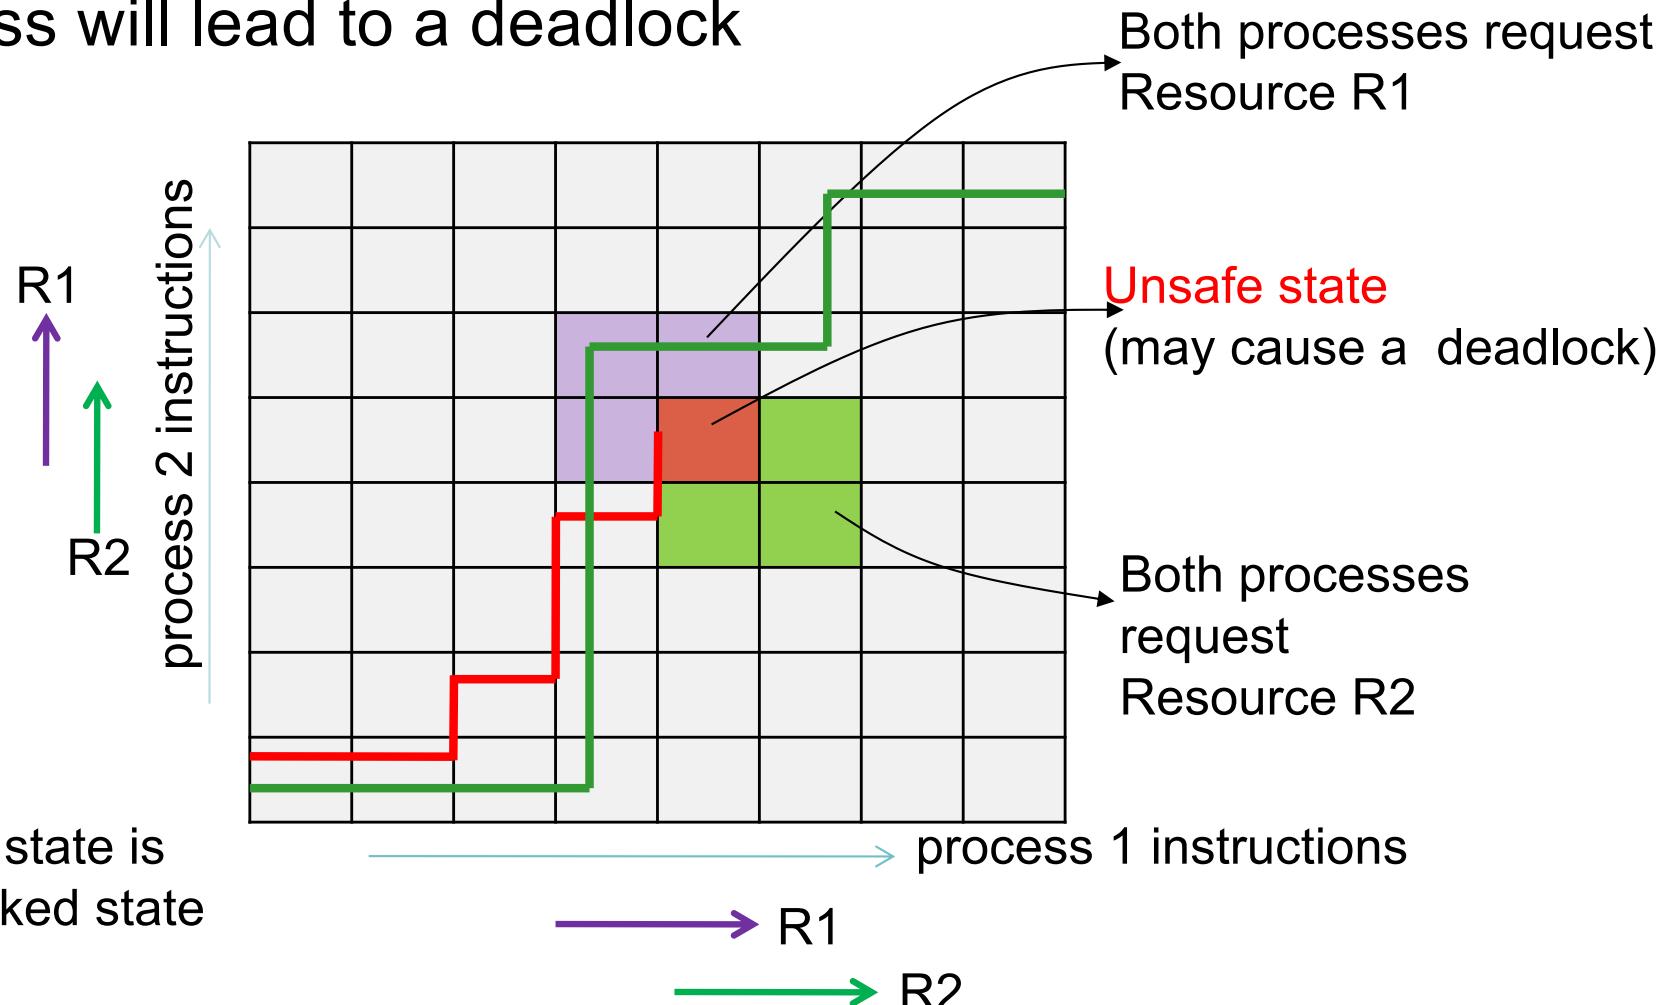
# Deadlock Avoidance

- System decides in advance if allocating a resource to a process will lead to a deadlock



# Deadlock Avoidance

- System decides in advance if allocating a resource to a process will lead to a deadlock



# Deadlock Avoidance

- Is there an algorithm that can always avoid deadlocks by conservatively making the right choice.
- Ensures system never reaches an unsafe state
- Safe state** : A state is said to be safe, if there is some scheduling order in which every process can run to completion even if all of them suddenly requests their maximum number of resources immediately
- An unsafe state **does not have to** lead to a deadlock; *it could lead to a deadlock*

# Example with a Banker

- Consider a banker with 4 clients ( $P_1, P_2, P_3, P_4$ ).
  - Each client has certain credit limits (totaling 20 units)
  - The banker knows that max credits will not be used at once, so he keeps only 10 units

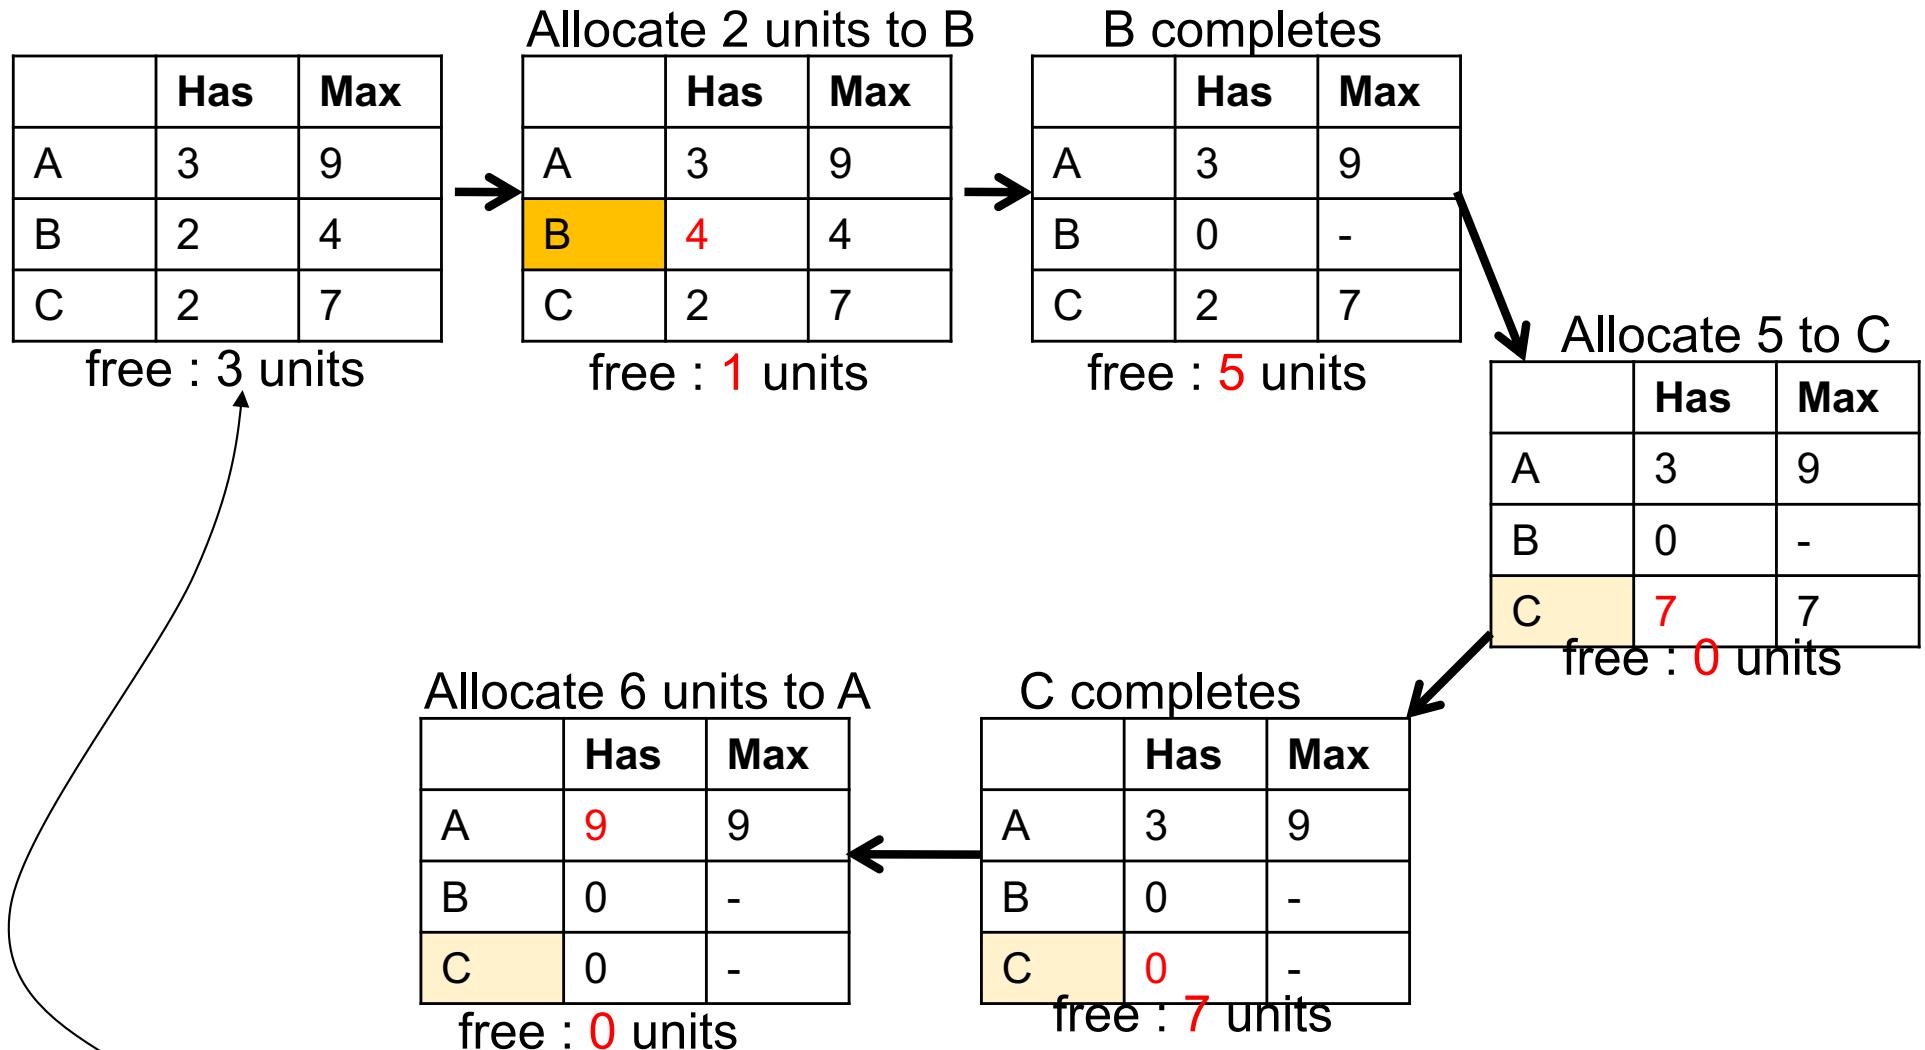
	Has	Max
A	3	9
B	2	4
C	2	7

Total : 10 units  
free : 3 units



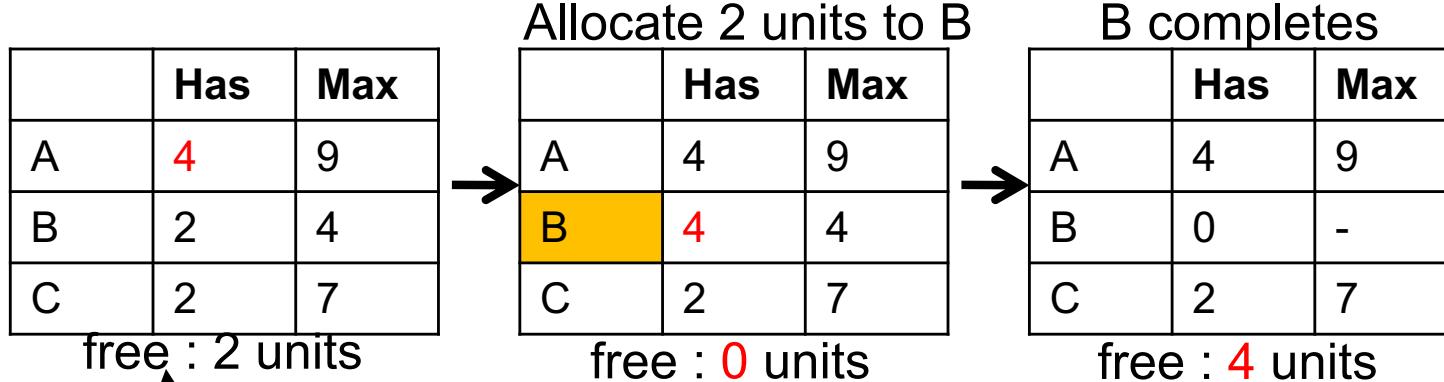
- Clients declare **maximum** credits in advance. The banker can allocate credits provided no unsafe state is reached.

# Safe State



This is a safe state because there is some scheduling order in which every process executes

# Unsafe State



This is an unsafe state because there exists NO scheduling order in which every process executes

# Banker's Algorithm (with a single resource)

When a request occurs

- If(is\_system\_in\_a\_safe\_state)
  - Grant request
- else
  - postpone until later



Please read Banker's Algorithm with multiple resources from Modern Operating Systems, Tanenbaum

# Deadlock Prevention

- Deadlock avoidance not practical, need to know maximum requests of a process
- Deadlock prevention
  - Prevent at-least one of the 4 conditions
    1. Mutual Exclusion
    2. Hold and wait
    3. No preemption
    4. Circular wait

# Prevention

## 1. Preventing Mutual Exclusion

- Not feasible in practice
- But OS can ensure that resources are optimally allocated

## 2. Hold and wait

- One way is to achieve this is to require all processes to request resources before starting execution
  - May not lead to optimal usage
  - May not be feasible to know resource requirements

## 3. No preemption

- Pre-empt the resources, such as by virtualization of resources (eg. Printer spools)

## 4. Circular wait

- One way, process holding a resource cannot hold a resource and request for another one
- Ordering requests in a sequential / hierarchical order.

# Hierarchical Ordering of Resources

- Group resources into levels  
(i.e. prioritize resources numerically)
- A process may only request resources at higher levels than any resource it currently holds
- Resource may be released in any order
- eg.
  - Semaphore  $s_1, s_2, s_3$  (with priorities in increasing order)  
 $\text{down}(S_1); \text{down}(S_2); \text{down}(S_3) ; \rightarrow \text{allowed}$   
 $\text{down}(S_1); \text{down}(S_3); \text{down}(S_2); \rightarrow \text{not allowed}$